

# A Unified Mathematical Framework for Conditioned Event Processes with Voluntary and Forced Payoff-Bearing Secondary Events

Theory, Generalizations, and Simulation Studies  
Complete Framework — Consolidated Edition (Version 3.1)

Macharia Bari  
bariimacharia@gmail.com

August 22, 2026

## Abstract

We develop a unified mathematical framework for event-driven decision-making under uncertainty. The framework generalizes classic point process and optimal stopping theory to handle: (i) data-driven trigger probabilities with state-dependent feedback; (ii) hierarchical temporal “hang” coupling across multiple time scales; (iii) generalized secondary event structures, including single-excitation (classic) and full branching processes (Galton–Watson); (iv) continuous and multi-category payoffs with type-dependent distributions; (v) baseline dynamics with target ( $K^*$ ) and floor ( $K_0$ ) constraints under risk aversion; (vi) rigorous optimal stopping via Hamilton–Jacobi–Bellman equations with viscosity solutions; and (vii) exact closed-form expressions for expectations, variances, and hitting time distributions via renewal theory. We provide a full likelihood framework, Bayesian hierarchical extension, model selection criteria, goodness-of-fit tests, and asymptotic distributions. Beyond assembly, we prove results specific to the structure: a value-of-controllability theorem (when raising the voluntary share strictly helps), an identification theory for the latent voluntary/forced mixture (including an impossibility result without state variation), a stability guarantee for the controlled branching process, and a martingale central limit theorem valid under policy feedback. The framework is validated by Monte Carlo simulation; illustrative simulated scenarios in finance, operations, and healthcare demonstrate the decision pipeline against Hawkes, ARIMA-GARCH, LSTM, and heuristic baselines, with a pre-specified protocol for future real-data benchmarking. A production-grade Python implementation is provided. This work provides a foundational, extensible tool for modeling, inference, and optimal control in event-driven systems.

**Keywords:** Event processes, point processes, optimal stopping, dynamic programming, branching processes, Hawkes processes, finance, operations, healthcare.

**AMS Subject Classifications:** 60G55, 60G40, 93E20, 90B50, 91G70.

# Contents

|          |                                                                             |           |
|----------|-----------------------------------------------------------------------------|-----------|
| <b>1</b> | <b>Introduction and Scope</b>                                               | <b>6</b>  |
| 1.1      | Contributions . . . . .                                                     | 6         |
| 1.2      | Paper Organization . . . . .                                                | 7         |
| <b>2</b> | <b>Related Work and Positioning</b>                                         | <b>7</b>  |
| <b>3</b> | <b>Hierarchical Temporal Lattice with State-Dependent Hazards</b>           | <b>9</b>  |
| 3.1      | Temporal Hierarchy . . . . .                                                | 9         |
| 3.2      | Trigger Event Process . . . . .                                             | 9         |
| <b>4</b> | <b>Generalized Secondary Events: From Single-Excitation to Branching</b>    | <b>10</b> |
| 4.1      | Core Model: Single-Excitation . . . . .                                     | 10        |
| 4.2      | Generalized Model: Full Branching . . . . .                                 | 10        |
| 4.3      | Voluntary and Forced Activation . . . . .                                   | 11        |
| 4.4      | Total Secondary Event Probability . . . . .                                 | 11        |
| <b>5</b> | <b>Payoffs and Utility</b>                                                  | <b>12</b> |
| 5.1      | Binary Payoffs (Core Model) . . . . .                                       | 12        |
| 5.2      | Continuous Payoffs (Generalized Model) . . . . .                            | 12        |
| 5.3      | Baseline Dynamics . . . . .                                                 | 12        |
| 5.4      | Risk Aversion . . . . .                                                     | 12        |
| <b>6</b> | <b>Derived Quantities</b>                                                   | <b>13</b> |
| 6.1      | Expected Cumulative Payoff . . . . .                                        | 13        |
| 6.2      | Probability of At Least One Payoff Event . . . . .                          | 13        |
| 6.3      | Expected Number of Payoff Events . . . . .                                  | 13        |
| 6.4      | Exact Variance of Cumulative Payoff . . . . .                               | 13        |
| 6.5      | Probability of Reaching Target . . . . .                                    | 13        |
| 6.6      | Exact Hitting Time via Renewal Theory . . . . .                             | 14        |
| 6.7      | Probability of Baseline Violation . . . . .                                 | 14        |
| <b>7</b> | <b>Likelihood, Bayesian Inference, Model Selection, and Goodness-of-Fit</b> | <b>14</b> |
| 7.1      | Likelihood Function . . . . .                                               | 14        |
| 7.2      | Bayesian Hierarchical Extension . . . . .                                   | 15        |

|           |                                                                           |           |
|-----------|---------------------------------------------------------------------------|-----------|
| 7.3       | Model Selection Criteria . . . . .                                        | 15        |
| 7.4       | Goodness-of-Fit Tests . . . . .                                           | 15        |
| 7.5       | Identification of the Latent Voluntary/Forced Mixture . . . . .           | 16        |
| <b>8</b>  | <b>Asymptotic Distribution</b>                                            | <b>16</b> |
| <b>9</b>  | <b>Hierarchical Aggregation Theorems</b>                                  | <b>17</b> |
| <b>10</b> | <b>Optimal Control: HJB, Viscosity Solutions, and Dynamic Programming</b> | <b>18</b> |
| 10.1      | Discrete-Time Bellman Equation . . . . .                                  | 18        |
| 10.2      | Continuous-Time HJB Equation . . . . .                                    | 18        |
| 10.3      | Viscosity Solutions . . . . .                                             | 18        |
| 10.4      | Optimal Policy Structure . . . . .                                        | 19        |
| 10.5      | Value of Controllability . . . . .                                        | 19        |
| 10.6      | Discrete-Time Dynamic Programming Algorithm . . . . .                     | 20        |
| 10.7      | Multi-Trigger Queue . . . . .                                             | 21        |
| 10.8      | Extension to Continuous Action Spaces . . . . .                           | 21        |
| <b>11</b> | <b>Closed-Form Solutions via Renewal Theory</b>                           | <b>22</b> |
| 11.1      | Exponential Delay . . . . .                                               | 22        |
| 11.2      | Geometric Delay . . . . .                                                 | 22        |
| 11.3      | Fixed Delay . . . . .                                                     | 22        |
| 11.4      | Renewal Hitting Times . . . . .                                           | 22        |
| <b>12</b> | <b>Synthetic Numerical Example and Monte Carlo Validation</b>             | <b>22</b> |
| 12.1      | Setup . . . . .                                                           | 22        |
| 12.2      | Analytical Results . . . . .                                              | 23        |
| 12.3      | Monte Carlo Simulation Results . . . . .                                  | 23        |
| <b>13</b> | <b>Illustrative Application Scenarios and Evaluation Protocol</b>         | <b>23</b> |
| 13.1      | Data Sources . . . . .                                                    | 23        |
| 13.2      | Benchmark Methods . . . . .                                               | 23        |
| 13.3      | Evaluation Protocol for Real-Data Validation . . . . .                    | 24        |
| 13.4      | Finance: S&P 500 Intraday Trading . . . . .                               | 24        |
| 13.5      | Operations: M5 Retail Sales . . . . .                                     | 24        |

|                                                                                |           |
|--------------------------------------------------------------------------------|-----------|
| 13.6 Healthcare: MIMIC-III Readmission Risk . . . . .                          | 25        |
| 13.7 Cross-Domain Summary and Illustrative Calibration . . . . .               | 25        |
| <b>14 Sensitivity Analysis Heatmaps</b>                                        | <b>26</b> |
| <b>15 Implications of Key Findings</b>                                         | <b>26</b> |
| 15.1 Separation of Trigger and Payoff Processes . . . . .                      | 26        |
| 15.2 The Hang Effect Creates Natural Buffers . . . . .                         | 26        |
| 15.3 Voluntary Actions Dominate Forced Events . . . . .                        | 27        |
| 15.4 Offline Dynamic Programming Enables Real-Time Decisions . . . . .         | 27        |
| 15.5 Sensitivity Drivers Are Concentrated . . . . .                            | 28        |
| 15.6 The Target–Floor Trade-off Requires an Explicit Risk Appetite . . . . .   | 28        |
| 15.7 Active Management Creates Value . . . . .                                 | 29        |
| 15.8 Forced Events Are Information Signals . . . . .                           | 29        |
| <b>16 Limitations and Future Directions</b>                                    | <b>29</b> |
| 16.1 Addressed by the Generalized Framework . . . . .                          | 29        |
| 16.2 Remaining Limitations . . . . .                                           | 29        |
| 16.3 Future Research Agenda . . . . .                                          | 30        |
| 16.4 Companion Papers Roadmap . . . . .                                        | 31        |
| <b>17 Conclusion</b>                                                           | <b>31</b> |
| <b>A Glossary of Symbols</b>                                                   | <b>31</b> |
| <b>B Proof of Dynamic Programming Convergence with State-Dependent Hazards</b> | <b>33</b> |
| B.1 Setup . . . . .                                                            | 33        |
| B.2 Assumptions . . . . .                                                      | 33        |
| <b>C Exact Variance Derivation</b>                                             | <b>35</b> |
| C.1 General Case . . . . .                                                     | 35        |
| C.2 Core Model (Single-Excitation) . . . . .                                   | 35        |
| <b>D Renewal Theory for Hitting Times</b>                                      | <b>35</b> |
| <b>E Supplementary Derivations</b>                                             | <b>35</b> |

|          |                                                                  |           |
|----------|------------------------------------------------------------------|-----------|
| E.1      | Derivation of Expected Payoff . . . . .                          | 35        |
| E.2      | Derivation of the HJB Equation . . . . .                         | 36        |
| E.3      | Derivation of the CLT . . . . .                                  | 36        |
| E.4      | Derivation of Hitting Time and Violation Probabilities . . . . . | 36        |
| <b>F</b> | <b>Code Implementation</b>                                       | <b>36</b> |
| F.1      | Core Model Implementation . . . . .                              | 36        |
| F.2      | Generalized Model Implementation . . . . .                       | 39        |

# 1 Introduction and Scope

The modeling of event-driven processes is fundamental to numerous domains including finance, operations management, and healthcare. While rich literatures exist on point processes [9], self-exciting processes (Hawkes) [17], and optimal stopping [3], a unified framework bridging empirical data-driven hazards, controlled voluntary actions, exogenous forced events, and state-dependent baseline constraints has been lacking.

This paper presents such a unified framework. Unlike prior work that claims completeness, we explicitly define the **scope** and **generality** of our approach. We present a **tiered framework**:

1. **Core Model (Sections 2–11)**: single-excitation, binary payoff, risk-neutral. Suitable for quick insights and initial deployment.
2. **Generalized Model (Sections 3–5, 9)**: branching processes, continuous payoffs, state-dependent triggers, risk aversion, and multi-trigger queues. Suitable for high-stakes applications requiring fidelity.
3. **Validation (Sections 12–13)**: synthetic Monte Carlo experiments and simulated cross-domain scenarios, with a pre-specified protocol for real-data benchmarking.

## 1.1 Contributions

Our primary contributions are:

1. A hierarchical temporal lattice with arbitrary period nesting and “hang” coupling across boundaries;
2. Data-driven trigger probabilities with no structural assumptions, extended to state-dependent hazards with baseline feedback;
3. A generalized secondary event structure spanning single-excitation to full branching (Galton–Watson), with dual voluntary (controlled) and forced (exogenous) event types;
4. Asymmetric binary payoffs ( $\pm 1$ ) and their continuous/multi-category generalizations, with type-dependent distributions;
5. Baseline dynamics with target  $K^* \geq K_0$  and floor constraint  $K_t \geq K_0$ , under CRRA risk aversion;
6. Closed-form expressions for all key quantities, including exact renewal-theoretic hitting times that eliminate prior Poisson approximations;
7. Full likelihood, Bayesian extension, model selection criteria (AIC, BIC, DIC, LRT), and goodness-of-fit tests;

8. Hamilton–Jacobi–Bellman optimal control with discrete dynamic programming convergence proofs; viscosity existence via Perron’s method and uniqueness via the comparison principle, with explicit treatment of the target-boundary discontinuity (§9.3);
9. A **value-of-controllability theorem** giving necessary and sufficient one-step conditions under which raising the voluntary share  $\pi_t(K)$  strictly increases the value function (Theorem 10.4);
10. An **identification theory** for the latent voluntary/forced mixture, including an impossibility result for single-state data and constructive recovery under state variation (Proposition 7.2);
11. A **stability guarantee** for the controlled branching process under uniformly subcritical policies (Proposition 4.2);
12. Asymptotic normality via a martingale central limit theorem valid under policy feedback, recovering the classical independent-increment formula as a special case (§7), together with exact closed-form period aggregation (§8);
13. Closed-form solutions for exponential, geometric, and fixed delay distributions;
14. Monte Carlo validation on synthetic data; illustrative simulated application scenarios in finance, operations, and healthcare with a pre-specified evaluation protocol for future real-data benchmarking (§11–§12);
15. Comprehensive sensitivity analysis with heatmaps;
16. Production-grade Python implementation of both the core and generalized models.

## 1.2 Paper Organization

Section 2 establishes the hierarchical temporal lattice with state-dependent hazards. Section 3 presents the generalized secondary event structure. Section 4 covers payoffs, baseline dynamics, and utility. Section 5 derives all key quantities exactly. Section 6 develops likelihood, Bayesian inference, model selection, and goodness-of-fit. Sections 7–8 present asymptotics and aggregation. Section 9 develops optimal control theory. Section 10 provides closed-form solutions via renewal theory. Section 11 reports synthetic Monte Carlo validation; Section 12 reports illustrative simulated application scenarios and the pre-specified real-data evaluation protocol. Section 13 presents sensitivity analysis. Sections 14–16 discuss implications, limitations, and future directions. Section 17 concludes. Appendices contain proofs, derivations, a glossary, and code.

## 2 Related Work and Positioning

The components of this framework draw on several established literatures, and we are explicit about what is assembled versus what is new.

**Self-exciting point processes.** Hawkes processes [17] and their marked extensions [11] capture events triggering further events, including branching-cluster representations. Our branching generalization (§3) is deliberately a marked-Hawkes object when the control dimension is removed; the novelty lies elsewhere.

**Controlled point processes and optimal execution.** A mature literature studies intervention in jump processes, notably in high-frequency trading: inventory/market-making models with Hawkes-style arrival structure [7], and mutually exciting jump processes in contagion modeling [1]. These works typically assume the controller acts at every event or continuously in time. None endogenously mixes controllable and uncontrollable responses to the *same* trigger stream, which is the structure introduced by the voluntary/forced decomposition (§3.4).

**Markov decision processes and reinforcement learning.** Any finite instance of our model can be encoded as an MDP [3], and the generalized model is designed as an RL testbed. The framework’s value relative to generic MDP machinery is structural: closed forms for hitting statistics (§5, §10), an identification theory for the latent type mixture (§6.5), and interpretable threshold policies (§9.4) that generic value-function approximation does not deliver.

**Optimal stopping and shock models.** The exit/target geometry relates to stopping problems [21] and to cumulative-damage shock models [8, 16], which treat ruin boundaries but not reward–ruin duality with controlled event streams.

**Adoption and threshold models.** The voluntary/forced split echoes the innovation/imitation structure of diffusion models [6] and threshold contagion [15]; here it is embedded in an optimal-control problem rather than a descriptive diffusion curve.

Table 1 summarizes capability coverage. The claim is not exclusivity — most features are individually expressible in some existing formalism — but that no single standard framework natively combines all rows while retaining closed forms, and that the combination generates the structural results proved here.

| Capability                            | Marked<br>Hawkes | Generic<br>MDP/RL | Optimal<br>stopping | Shock/<br>renewal | <b>This framework</b> |
|---------------------------------------|------------------|-------------------|---------------------|-------------------|-----------------------|
| State-dependent hazard $\lambda_t(K)$ | ✓                | ✓                 | ✓                   | ✓                 | ✓                     |
| Event-triggered decisions             | –                | ✓ (encoded)       | ✓                   | –                 | ✓                     |
| Voluntary/forced mix per trigger      | –                | ✓ (encoded)       | (○)                 | –                 | ✓ (native)            |
| Target–floor dual boundaries          | –                | ✓                 | ✓                   | ✓                 | ✓                     |
| Closed-form hitting-time laws         | –                | –                 | rare                | ✓                 | ✓                     |
| Identification of latent type mix     | –                | –                 | n/a                 | –                 | ✓ (Prop. 7.2)         |
| Value-of-controllability theorem      | –                | –                 | –                   | –                 | ✓ (Thm. 10.4)         |
| Controlled-branching stability bound  | (○)              | –                 | –                   | ✓                 | ✓ (Prop. 4.2)         |

Table 1: Positioning against adjacent frameworks. ✓ = native support; (○) = partial; – = absent or requires nontrivial re-encoding. References: marked Hawkes [11, 17]; MDP/RL [3, 23]; stopping [21]; shock/renewal [8, 16].



### 3 Hierarchical Temporal Lattice with State-Dependent Hazards

#### 3.1 Temporal Hierarchy

Let time be partitioned into nested periods at levels  $\ell = 0, 1, \dots, L$ , where level 0 is the atomic unit (e.g., day, hour, minute). For each level  $\ell$ , periods have duration  $T_\ell$  in level-0 units.

**Definition 3.1** (Temporal Hierarchy). *For each level  $\ell \in \{0, \dots, L\}$ , let:*

- $\mathcal{P}_\ell$ : the set of all periods at level  $\ell$ ;
- $T_\ell \in \mathbb{Z}_{>0}$ : the duration of a period at level  $\ell$  in atomic units;
- For any period  $P \in \mathcal{P}_\ell$ , it contains exactly  $T_\ell/T_{\ell-1}$  sub-periods from level  $\ell - 1$ .

**Definition 3.2** (Hang Operator). *For any level  $\ell$  and period  $P \in \mathcal{P}_\ell$ , define the hang window:*

$$H_\ell(P) = \text{suffix}_{s_\ell}(P) \cup \text{prefix}_{r_\ell}(\text{next}(P)), \quad (1)$$

where  $s_\ell \in \mathbb{Z}_{\geq 0}$  is the suffix length and  $r_\ell \in \mathbb{Z}_{\geq 0}$  is the prefix length. This couples the end of one period with the beginning of the next.

**Axiom 3.3** (Hang Persistence). *For any hang window  $H$  at any level  $\ell$ , the condition state is constant:*

$$C_t = C_{t'} \quad \forall t, t' \in H, \quad (2)$$

where  $C_t \in \{0, 1\}$  is the condition state at atomic time  $t$ .

#### 3.2 Trigger Event Process

Let  $E_t \in \{0, 1\}$  denote a **trigger event** at atomic time  $t$ . Trigger events carry no payoff; they exist solely to generate secondary responses.

**Axiom 3.4** (Empirical Trigger Probability). *In the core model, the probability of a trigger event at time  $t$  is a known function derived from data:*

$$\boxed{\pi_t := \mathbb{P}(E_t = 1) = \text{observed (or smoothed) empirical rate at } t.} \quad (3)$$

No structural assumption is made about how  $\pi_t$  is generated; it is taken as given from empirical observations.

The empirical trigger rate is estimated as

$$\hat{\pi}_t = \frac{\#\{E = 1 \text{ at } t \text{ across historical windows}\}}{\#\{\text{observations at } t\}}, \quad (4)$$

and within any hang window the trigger rate is constant:  $\pi_t = \pi_{t'}$  for all  $t, t' \in H_\ell$ .

**Axiom 3.5** (State-Dependent Trigger Probability). *In the generalized model, the trigger probability depends on the current baseline  $K_t$ :*

$$\pi_t(K_t) := \mathbb{P}(E_t = 1 \mid K_t) = \text{logistic}(\gamma_0 + \gamma_1 K_t + \gamma_2 t + \gamma_3 K_t^2 + \gamma_4 \cdot \mathbf{Z}_t), \quad (5)$$

where  $\text{logistic}(x) = (1 + e^{-x})^{-1}$ ,  $\gamma_0$  is an intercept,  $\gamma_1$  captures linear feedback,  $\gamma_2$  captures time trends,  $\gamma_3$  captures nonlinear feedback, and  $\mathbf{Z}_t$  are exogenous covariates. Setting all feedback coefficients to zero recovers the empirical axiom above. The parameters  $\gamma$  are estimated via maximum likelihood or Bayesian inference.

## 4 Generalized Secondary Events: From Single-Excitation to Branching

Each trigger at time  $s$  generates secondary events at times  $t > s$ . Secondary events come in two types: **voluntary** (chosen by the decision-maker) and **forced** (exogenous).

### 4.1 Core Model: Single-Excitation

Each trigger generates **at most one** secondary event. Let  $N(s) \in \{0, 1\}$  be the number of offspring:

$$\mathbb{P}(N(s) = 1) = \rho, \quad \mathbb{P}(N(s) = 0) = 1 - \rho. \quad (6)$$

The kernel is

$$\kappa(s, t) = (1 - \phi) \alpha_t \rho \psi(t - s; \theta_\psi) + \phi \rho \xi(t - s; \theta_\xi), \quad (7)$$

where  $\phi$  is the forced fraction,  $\alpha_t \in [0, 1]$  is the voluntary activation policy,  $\psi$  is the voluntary delay distribution, and  $\xi$  is the forced delay distribution.

### 4.2 Generalized Model: Full Branching

Let  $N(s) \in \mathbb{Z}_{\geq 0}$  follow a distribution  $\nu(\cdot; \theta_\nu)$  (e.g., Poisson, Geometric, Negative Binomial). The expected offspring is

$$\mu := \mathbb{E}[N(s)] = \sum_{n=0}^{\infty} n \nu(n; \theta_\nu). \quad (8)$$

The generalized kernel is

$$\kappa_{\text{Gen}}(s, t) = \sum_{n=1}^{\infty} \nu(n; \theta_\nu) \psi^{*n}(t - s; \theta_\psi), \quad (9)$$

where  $\psi^{*n}$  is the  $n$ -fold convolution of the delay distribution; the voluntary/forced decomposition applies to each offspring event.

**Definition 4.1** (Criticality). *The branching process is:*

- **Subcritical** if  $\mu < 1$ : clusters die out almost surely;

- **Critical** if  $\mu = 1$ : clusters have a power-law size distribution;
- **Supercritical** if  $\mu > 1$ : clusters grow without bound with positive probability.

For subcritical branching, the expected total cluster size is  $\mathbb{E}[\text{Cluster Size}] = \frac{1}{1-\mu}$ .

**Proposition 4.2** (Stability under Control). *Let  $\Pi_{\bar{\rho}}$  denote the class of Markov policies satisfying uniform subcriticality:*

$$\sup_{t, K \in \mathcal{K}, \pi \in \Pi_{\bar{\rho}}} \mathbb{E}[N^{\text{tot}}(s) \mid \mathcal{F}_t, K_t = K] \leq \bar{\rho} < 1, \quad (10)$$

where  $N^{\text{tot}}(s)$  counts all offspring (direct and secondary of all generations) generated by a trigger at time  $s$ . Then: (i) cluster sizes are stochastically dominated by a Galton–Watson process with mean  $\bar{\rho}$ , whence  $\mathbb{E}[\text{Cluster Size}] \leq (1 - \bar{\rho})^{-1}$  and tail probabilities decay geometrically; (ii) the induced level process satisfies  $\sup_t \mathbb{E}[K_t] < \infty$  whenever payoffs have bounded conditional mean; (iii) the optimal policy within the restricted class  $\Pi_{\bar{\rho}}$  inherits the bound, since the Bellman operator preserves stochastic domination.

*Proof sketch.* Compare generation-by-generation expected population counts with those of a single-type Galton–Watson process of mean  $\bar{\rho}$ ; uniform subcriticality makes the comparison process an upper dominating recursion by induction on generations, yielding (i) from classical Galton–Watson theory [16]. Boundedness of conditional payoff means converts the cluster bound into a uniform  $L^1$  increment bound, giving (ii). For (iii), if two policies induce dominating recursions at every state, the Bellman maximization over the restricted class preserves the pointwise ordering, so value iteration converges within the bounded class. A caveat deserves emphasis: *without* the restriction, an optimizing policy can push the system toward criticality near the target (intensity feedback), and global stability of the unrestricted problem remains open.  $\square$

### 4.3 Voluntary and Forced Activation

**Axiom 4.3** (Voluntary Activation). *Given a trigger at time  $s$ , the decision-maker selects a stopping time  $\tau \geq s$ :*

$$\mathbb{P}(\text{voluntary secondary at } t \mid E_s = 1, \text{policy}) = \alpha_t \cdot \mathbf{1}\{t \geq s\}, \quad (11)$$

where  $\alpha_t \in [0, 1]$  is the voluntary activation intensity (policy).

**Axiom 4.4** (Forced Activation). *With probability  $\phi \in [0, 1]$ , a trigger at time  $s$  forces a secondary event:*

$$\mathbb{P}(\text{forced secondary at } t \mid E_s = 1) = \phi \cdot \xi(t - s; \theta_\xi). \quad (12)$$

### 4.4 Total Secondary Event Probability

The total probability of a secondary event at time  $t$  is

$$r_t := \mathbb{P}(R_t = 1) = \sum_{s < t} \kappa(s, t) \pi_s(K_s), \quad (13)$$

where  $\kappa$  reduces to the simple single-excitation kernel when  $\nu(1) = 1$ .

## 5 Payoffs and Utility

### 5.1 Binary Payoffs (Core Model)

Each secondary event carries a binary payoff  $X_t \in \{-1, +1\}$  with probabilities dependent on type:

$$\mathbb{P}(X_t = +1 \mid R_t = 1, \text{type} = V) = \eta_t^V, \quad \mathbb{P}(X_t = +1 \mid R_t = 1, \text{type} = F) = \eta_t^F. \quad (14)$$

Thus  $\mathbb{E}[X_t \mid R_t = 1, \text{type}] = 2\eta_t^{\text{type}} - 1$ , and we define the average payoff expectation

$$\bar{\eta}_t = (1 - \phi) \alpha_t (2\eta_t^V - 1) + \phi (2\eta_t^F - 1) \in [-1, 1]. \quad (15)$$

### 5.2 Continuous Payoffs (Generalized Model)

Let  $X_t \in \mathbb{R}$  denote the payoff at time  $t$ . Conditional on  $R_t = 1$  and type,

$$\mathbb{P}(X_t \in dx \mid R_t = 1, V) = f_t^V(x) dx, \quad \mathbb{P}(X_t \in dx \mid R_t = 1, F) = f_t^F(x) dx, \quad (16)$$

with densities  $f_t^V, f_t^F$  (e.g., Normal, Log-Normal, Gamma) and conditional moments  $\mu_t^V, \sigma_t^{V,2}, \mu_t^F, \sigma_t^{F,2}$ . The mixed conditional mean is  $\bar{\mu}_t = (1 - \phi)\alpha_t\mu_t^V + \phi\mu_t^F$ .

### 5.3 Baseline Dynamics

The running baseline evolves as

$$K_t = K_{t-1} + X_t R_t. \quad (17)$$

There exists a target baseline  $K^* \geq K_0$  such that the floor constraint

$$\boxed{K_t \geq K_0 \quad \forall t \geq 0} \quad (18)$$

must hold, and the decision-maker aims for

$$\boxed{\lim_{t \rightarrow \infty} \mathbb{P}(K_t \geq K^*) \rightarrow 1.} \quad (19)$$

### 5.4 Risk Aversion

We adopt Constant Relative Risk Aversion (CRRA) utility:

$$U(K) = \frac{K^{1-\gamma_R}}{1-\gamma_R}, \quad \gamma_R \geq 0, \quad (20)$$

where  $\gamma_R$  is the risk-aversion coefficient ( $\gamma_R = 0$  gives risk neutrality  $U(K) = K$ ;  $\gamma_R \rightarrow 1$  gives  $U(K) = \log K$ ). The decision-maker maximizes  $\max_{\alpha} \mathbb{E}[U(K_T)]$ .

## 6 Derived Quantities

### 6.1 Expected Cumulative Payoff

Starting from  $K_T = K_0 + \sum_{t=1}^T X_t R_t$  and using conditional independence ( $\mathbb{E}[X_t R_t] = r_t \cdot \mathbb{E}[X_t | R_t = 1]$ ):

$$\boxed{\mathbb{E}[K_T] = K_0 + \sum_{t \leq T} r_t \cdot [(1 - \phi) \alpha_t (2\eta_t^V - 1) + \phi (2\eta_t^F - 1)]} \quad (21)$$

for the core model, and in general

$$\mathbb{E}[K_T] = K_0 + \sum_{t \leq T} \mathbb{E}[R_t] \cdot \bar{\mu}_t. \quad (22)$$

### 6.2 Probability of At Least One Payoff Event

$$\mathbb{P}\left(\sum_{t \leq T} R_t > 0\right) = 1 - \prod_{t \leq T} (1 - r_t). \quad (23)$$

### 6.3 Expected Number of Payoff Events

$$\mathbb{E}\left[\sum_{t \leq T} R_t\right] = \sum_{t \leq T} r_t. \quad (24)$$

### 6.4 Exact Variance of Cumulative Payoff

Unlike earlier approximations, we derive the exact variance:

$$\boxed{\text{Var}(K_T) = \sum_{s=1}^T \sum_{t=1}^T \text{Cov}(X_s R_s, X_t R_t),} \quad (25)$$

where  $\text{Cov}(X_s R_s, X_t R_t) = \mathbb{E}[X_s R_s X_t R_t] - \mathbb{E}[X_s R_s] \mathbb{E}[X_t R_t]$ . For the core model this simplifies to

$$\text{Var}(K_T) = \sum_{t \leq T} r_t [1 - \bar{\eta}_t^2 r_t] + 2 \sum_{s < t \leq T} \text{Cov}(X_s R_s, X_t R_t), \quad (26)$$

and for independent triggers (no carryover) the double-covariance term vanishes. Appendix C gives the full derivation.

### 6.5 Probability of Reaching Target

Let  $\theta = K^* - K_0 > 0$ . With  $N_T = \sum_{t \leq T} R_t$  the total number of secondary events,

$$\mathbb{P}(K_T \geq K^*) = \sum_{n=0}^{\infty} \mathbb{P}(N_T = n) \mathbb{P}\left(\sum_{i=1}^n X_i \geq \theta\right), \quad (27)$$

which for binary i.i.d. payoffs with success probability  $\bar{\eta}$  evaluates to the binomial mixture

$$\mathbb{P}(K_T \geq K^*) = \sum_{n=0}^{\infty} \mathbb{P}(N_T = n) \sum_{k=\theta}^n \binom{n}{\frac{n+k}{2}} \bar{\eta}^{\frac{n+k}{2}} (1 - \bar{\eta})^{\frac{n-k}{2}}. \quad (28)$$

## 6.6 Exact Hitting Time via Renewal Theory

Let  $\tau_\theta$  be the first time the baseline reaches  $K_0 + \theta$ . The expected hitting time solves the renewal equation

$$\mathbb{E}[\tau_\theta] = \frac{1}{\lambda} + \sum_x \mathbb{P}(X = x) \cdot \mathbb{E}[\tau_{\theta-x}], \quad \mathbb{E}[\tau_\theta] = 0 \text{ for } \theta \leq 0, \quad (29)$$

which eliminates the Poisson approximation used in earlier drafts. For binary payoffs the closed form is derived in Section 10.4.

## 6.7 Probability of Baseline Violation

Using the reflection principle for random walks,

$$\mathbb{P}(\exists t \leq T : K_t < K_0) = 1 - \mathbb{P}\left(\min_{t \leq T} (K_t - K_0) \geq 0\right), \quad (30)$$

and for a random walk with drift  $\mu_{\text{net}} = \mathbb{E}[X_t R_t] < 0$  and variance  $\sigma^2$ ,

$$\mathbb{P}(\text{violation}) \approx \exp\left(-\frac{2|\mu_{\text{net}}|}{\sigma^2}\right), \quad (31)$$

the classic gambler's ruin result.

# 7 Likelihood, Bayesian Inference, Model Selection, and Goodness-of-Fit

## 7.1 Likelihood Function

Given data  $\mathcal{D} = \{(E_t, R_t, X_t, \text{type}_t)\}_{t=1}^T$ , the full log-likelihood is

$$\log \mathcal{L}(\rho, \phi, \theta_\psi, \theta_\xi, \eta^V, \eta^F \mid \mathcal{D}) = \sum_{t=1}^T \left[ E_t \log \pi_t + (1 - E_t) \log(1 - \pi_t) \right] \quad (32)$$

$$+ R_t \log r_t + (1 - R_t) \log(1 - r_t) \quad (33)$$

$$+ \mathbf{1}_{\{R_t=1, \text{type}=V\}} [X_t \log \eta_t^V + (1 - X_t) \log(1 - \eta_t^V)] \quad (34)$$

$$+ \mathbf{1}_{\{R_t=1, \text{type}=F\}} [X_t \log \eta_t^F + (1 - X_t) \log(1 - \eta_t^F)] \quad (35)$$

Maximization proceeds via gradient ascent or Newton–Raphson; the trigger and payoff blocks are separable, enabling modular estimation.

## 7.2 Bayesian Hierarchical Extension

Place conjugate priors on all parameters:

$$\rho \sim \text{Beta}(a_\rho, b_\rho), \quad \phi \sim \text{Beta}(a_\phi, b_\phi), \quad (36)$$

$$\theta_\psi \sim \text{Gamma}(a_\psi, b_\psi), \quad \theta_\xi \sim \text{Gamma}(a_\xi, b_\xi), \quad (37)$$

$$\eta_t^V \sim \text{Beta}(a_V, b_V), \quad \eta_t^F \sim \text{Beta}(a_F, b_F). \quad (38)$$

The posterior is  $\mathbb{P}(\Theta \mid \mathcal{D}) \propto \mathcal{L}(\Theta \mid \mathcal{D}) \cdot \mathbb{P}(\Theta)$ , and the predictive distribution is

$$\mathbb{P}(R_{T+1} = 1 \mid \mathcal{D}) = \int \mathbb{P}(R_{T+1} = 1 \mid \Theta) \mathbb{P}(\Theta \mid \mathcal{D}) d\Theta. \quad (39)$$

## 7.3 Model Selection Criteria

$$\text{AIC} = -2 \log \mathcal{L}(\hat{\Theta}) + 2m, \quad \text{BIC} = -2 \log \mathcal{L}(\hat{\Theta}) + m \log T, \quad (40)$$

$$\text{DIC} = -2 \mathbb{E}[\log \mathcal{L}(\Theta)] + 2p_D, \quad (41)$$

where  $m$  is the number of estimated parameters and  $p_D$  the effective number of parameters. Nested hypotheses are tested as follows:

| Null Model                   | Alternative                           | Test                    |
|------------------------------|---------------------------------------|-------------------------|
| $\phi = 0$                   | $\phi > 0$                            | LRT $\sim \chi_1^2$     |
| $\eta^V = \eta^F$            | $\eta^V \neq \eta^F$                  | LRT $\sim \chi_1^2$     |
| $\alpha_t = \alpha$ constant | $\alpha_t$ varying                    | LRT $\sim \chi_{T-1}^2$ |
| $\pi_t = \pi$ constant       | $\pi_t$ data-driven / state-dependent | Compare AIC/BIC         |

Table 2: Nested model tests.

## 7.4 Goodness-of-Fit Tests

**Single-Excitation Test.** For each trigger  $i$ , let  $N_i \in \{0, 1\}$  indicate whether it produced a secondary event:

$$\chi^2 = \sum_{i=1}^M \frac{(N_i - \hat{p})^2}{\hat{p}(1 - \hat{p})} \sim \chi_{M-1}^2. \quad (42)$$

**Delay Distribution Fit.** Kolmogorov–Smirnov statistic:

$$D_{\text{KS}} = \sup_{\Delta} |\hat{F}(\Delta) - F(\Delta; \hat{\theta})|. \quad (43)$$

**Payoff Independence Test.** Likelihood ratio test for constant  $\eta$ :

$$\text{LRT} = -2 \log \frac{\mathcal{L}(\eta \text{ constant})}{\mathcal{L}(\eta_t \text{ varying})} \sim \chi_{T-1}^2. \quad (44)$$

## 7.5 Identification of the Latent Voluntary/Forced Mixture

The econometrician observes trigger indicators  $\{E_t\}$  and realized level changes  $\{\Delta K_t = X_t R_t\}$ , but neither the latent event type (voluntary-offered vs. forced) nor the action  $\alpha_t$ . Identification is therefore nontrivial and deserves explicit conditions.

**Proposition 7.1** (Impossibility without state variation). *Suppose all triggers occur at a single state level  $K$  with constant parameters  $(\pi, \alpha, \eta^V, \eta^F)$  and outcomes in  $\{-1, 0, +1\}$  (engage/skip/forced moves). The observed law of  $\Delta K_t \mid E_t = 1$  is the three-point mixture*

$$p = \pi (\alpha \eta^V, \alpha(1 - \eta^V), 1 - \alpha) + (1 - \pi) (\eta^F, 1 - \eta^F, 0), \quad (45)$$

*a point on a chord of the probability simplex generated by two fixed endpoints. Four structural parameters determine two free mixture coordinates: infinitely many parameter combinations reproduce  $p$ , so the model is **not identified** from single-state data.*

**Proposition 7.2** (Identification under state variation). *Assume: (i) the voluntary share follows a known parametric family  $\pi(K) = \sigma(\theta_0 + \theta_1 K)$  that attains at least two distinct values in  $(\varepsilon, 1 - \varepsilon)$  over sampled states,  $\varepsilon > 0$ ; (ii) the engagement policy is a fixed measurable function  $\alpha(K)$ , constant across observations at the same state; (iii)  $\eta^V \neq \eta^F$  and both lie strictly in  $(0, 1)$ . Then as the number of triggers per state grows, the per-state outcome laws  $\{p_K\}$  trace a non-degenerate segment of the chord between the forced endpoint  $q_F = (\eta^F, 1 - \eta^F, 0)$  and the voluntary endpoint  $q_V(\alpha) = (\alpha \eta^V, \alpha(1 - \eta^V), 1 - \alpha)$ . The extreme points of the observed curve recover  $\{q_F, q_V\}$  up to label switching; linear inversion then recovers  $\pi(K)$  for every sampled state, and hence  $(\theta_0, \theta_1, \alpha, \eta^V, \eta^F)$  up to the voluntary/forced label convention.*

*Proof sketch.* With at least two distinct interior mixing weights  $\pi_1 \neq \pi_2$ , the observed laws satisfy  $p_{K_i} = (1 - \pi_i)q_F + \pi_i q_V$ , distinct points on an open segment. Extending to limits along states where  $\pi(K) \rightarrow 0$  or  $\rightarrow 1$  (guaranteed asymptotically by interiority plus strict monotonicity of  $\sigma$ ) identifies the endpoints as extreme points of the closed convex hull. Given the endpoints, weights follow by linear inversion; label switching is resolved by imposing  $\eta^V > \eta^F$  (or by a small validation sample with observed actions).  $\square$

**Remark 7.3** (Design implication). *Propositions 7.1–7.2 make experimental design part of the modeling contract: identification requires either exogenous variation in the voluntary share (randomized encouragement, natural experiments, platform A/B tests shifting  $\pi(K)$ ) or auxiliary labels of event type. EM estimation [12] is consistent only under these conditions; absent them, reported decompositions are underidentified and should be treated as assumptions, not estimates.*

## 8 Asymptotic Distribution

Under optimal control the level process is *self-referential*: increments depend on the state through the policy, so independence assumptions fail. The correct asymptotic statement is a martingale central limit theorem for the controlled chain.

**Theorem 8.1** (Martingale CLT under policy feedback). *Let  $K_t^\pi$  follow the controlled dynamics of §2–§3 under an admissible Markov policy  $\pi$ , with bounded payoff increments  $|X_t R_t| \leq B$ . Assume:*



- (A1) the controlled chain  $K_t^\pi$  admits an invariant distribution  $\mu_\pi$  and is geometrically ergodic (a sufficient condition is Doeblin minorization: from every reachable state there is a uniformly accessible neighborhood of the floor and target within  $\bar{n}$  steps);
- (A2) the one-step conditional mean drift  $m(K) := \mathbb{E}[X_t R_t \mid K_{t-1}^\pi = K]$  and conditional variance  $s^2(K) := \text{Var}(X_t R_t \mid K_{t-1}^\pi = K)$  are continuous;
- (A3)  $\sigma_\pi^2 := \int s^2 d\mu_\pi > 0$ .

Then, as  $T \rightarrow \infty$ ,

$$\boxed{\frac{K_T - \sum_{t=1}^T m(K_{t-1})}{\sqrt{T} \sigma_\pi} \xrightarrow{d} \mathcal{N}(0, 1)}, \quad (46)$$

and if additionally the drift integral  $\bar{m} := \int m d\mu_\pi$  exists,  $\frac{K_T - T\bar{m}}{\sqrt{T} \sigma_\pi} \xrightarrow{d} \mathcal{N}(0, 1)$ .

*Proof sketch.* Doob decomposition:  $M_T := K_T - \sum_t m(K_{t-1})$  is a martingale with bounded increments and conditional variance process  $\sum_t s^2(K_{t-1})$ . Geometric ergodicity (A1) plus continuity (A2) gives, via the ergodic theorem,  $\frac{1}{T} \sum s^2(K_{t-1}) \rightarrow \sigma_\pi^2$  almost surely. The martingale triangular-array CLT (Heyde's condition; [5, 13]) then yields the result; stationarity of  $(K_t)$  makes it a stationary ergodic martingale-difference array, to which Billingsley's Theorem applies directly.  $\square$

**Corollary 8.2** (Independent-trigger special case). *If triggers are i.i.d. with no state feedback ( $\lambda$  constant, no branching), the covariance sum vanishes and (A1)–(A3) hold trivially, recovering the classical closed form*

$$\sigma^2 = \pi\rho[(1 - \phi)\alpha(2\eta^V - 1)^2 + \phi(2\eta^F - 1)^2]. \quad (47)$$

The general case adds the autocovariance terms  $2 \sum_{h \geq 1} \text{Cov}(X_t R_t, X_{t+h} R_{t+h})$ , which are nonzero under policy feedback and must be estimated, not assumed away.

## 9 Hierarchical Aggregation Theorems

**Theorem 9.1** (Period Expectations). *For any period  $P$  at level  $\ell$ :*

$$\mathbb{E}[N_E(P)] = \sum_{t \in P} \pi_t, \quad \mathbb{E}[N_R(P)] = \sum_{t \in P} r_t, \quad (48)$$

$$\mathbb{E}[\Delta K(P)] = \sum_{t \in P} r_t [(1 - \phi)\alpha_t(2\eta_t^V - 1) + \phi(2\eta_t^F - 1)]. \quad (49)$$

**Theorem 9.2** (Emptiness). *For a block  $B$  with no internal triggers:*

$$\mathbb{P}(N_R(B) = 0) = \prod_{t \in B} (1 - r_t). \quad (50)$$

## 10 Optimal Control: HJB, Viscosity Solutions, and Dynamic Programming

### 10.1 Discrete-Time Bellman Equation

Let  $V(t, K)$  be the value function at time  $t$  with baseline  $K$ . The Bellman equation is

$$V(t, K) = \max_{\alpha \in [0, 1]} \left\{ \pi_t(K) [\alpha \mathbb{E}[V(t+1, K + X_t)] + (1 - \alpha) V(t+1, K)] + (1 - \pi_t(K)) \mathbb{E}[V(t+1, K + X_t^{\text{forced}})] \right\}, \quad (51)$$

with boundary conditions

$$V(t, K) = \text{ExitValue}(K) \quad \text{if } K < K_0 \text{ (violation),} \quad (52)$$

$$V(t, K) = \text{TargetBonus}(K) \quad \text{if } K \geq K^* \text{ (target reached),} \quad (53)$$

$$V(T, K) = U(K) \quad \text{(terminal utility).} \quad (54)$$

### 10.2 Continuous-Time HJB Equation

Taking the continuous-time limit, the HJB equation becomes

$$\boxed{0 = \frac{\partial V}{\partial t} + \lambda_t(K_t) \max_{\alpha} \{ \alpha \mathbb{E}[V(t, K + X) - V(t, K)] \} + \phi \mathbb{E}[V(t, K + X^{\text{forced}}) - V(t, K)]} \quad (55)$$

where  $\lambda_t(K_t)$  is the total intensity of all triggers and their branching offspring.

### 10.3 Viscosity Solutions

The HJB equation is a first-order nonlinear PDE with a singular control term. Two technical points require care before standard theory applies.

**(a) Discontinuous boundary data.** The terminal reward contains an indicator jump at the target,  $\text{TargetBonus}(K) \mathbf{1}\{K \geq K^*\}$ . The value function is therefore *not expected to be continuous* across  $K^*$ , and classical viscosity well-posedness (which presupposes continuous solutions) cannot be invoked blindly. We adopt the stopping-time formulation: for  $\tau$  ranging over stopping times bounded by  $T$ ,

$$V(t, K) = \sup_{\tau} \mathbb{E} \left[ e^{-r(\tau-t)} G(K_{\tau}) + \int_t^{\tau} e^{-r(s-t)} c(K_s) ds \right], \quad (56)$$

$$G(k) = \max \{ \text{ExitValue}(k), \text{TargetBonus}(k) \mathbf{1}\{k \geq K^*\} \},$$

which restores upper semicontinuity of the payoff functional on compact subsets of  $\Omega := (0, T) \times (K_0, K^*)$  and lets the discontinuity live entirely in the boundary condition, where it belongs.

**(b) Structure of the Hamiltonian.** The control enters only through a maximum over a finite action set of terms affine in the state gradient, so the Hamiltonian is convex in its gradient argument and degenerate-elliptic; this is precisely the class covered by the comparison principle.

**Theorem 10.1** (Existence and uniqueness). *Under boundedness of payoffs, Lipschitz continuity of  $\pi_t(\cdot)$  on  $\Omega$ , subcritical branching ( $\mu < 1$ ), and discounting  $r > 0$ : (i) (Existence.) The stopping-time value function is a discontinuous viscosity solution of the HJB equation, constructed by Perron’s method with the USC/LSC envelopes [2]; (ii) (Uniqueness.) Any two continuous viscosity solutions satisfying the comparison hypotheses — properness and degenerate ellipticity above, plus boundary conditions interpreted with pasting on the floor and target — coincide on  $\Omega$ , by the user’s-guide comparison theorem [10]; (iii) (Regularity.)  $V$  is Lipschitz on compact subsets of  $\Omega$ ; continuity extends up to the target boundary if and only if the bonus jump vanishes there.*

*Sketch.* (1) Dynamic programming yields the dynamic programming principle; Perron’s method then produces a viscosity solution as the supremum of subsolutions, USC by construction [2]. (2) The Hamiltonian’s convexity in the gradient and properness in  $(r, V)$  give nondegeneracy; comparison follows from the Crandall–Ishii–Lions doubling-of-variables argument, with the indicator handled inside  $G$  rather than the PDE [10]. (3) Lipschitz regularity follows from boundedness of payoffs, the discount factor, and subcriticality bounding the jump intensity. Full details follow the discrete-to-continuous limit of the convergent scheme in Appendix B.  $\square$

**Remark 10.2** (Regularization alternative). *If continuous boundary data are preferred, replace the indicator by the Lipschitz ramp  $b_\varepsilon(k) = \min\{1, (k - K^*)/\varepsilon\}_+$ . Since the Bellman operator is a contraction with constant  $\phi T < 1$  in the terminal data,  $|V - V^\varepsilon| \leq \sup_k |\text{TargetBonus}(k)| \cdot \phi T \wedge 1$ : continuity is purchased at an explicitly bounded approximation error.*

## 10.4 Optimal Policy Structure

**Theorem 10.3** (Threshold Policy). *The optimal voluntary activation policy is a threshold rule:*

$$\alpha_t^*(K) = \begin{cases} 1, & K_t - K_0 \geq \theta_{\text{act}}, \\ 0, & K_t - K_0 < \theta_{\text{act}}, \end{cases} \quad (57)$$

where  $\theta_{\text{act}}$  solves the indifference condition  $\mathbb{E}[X_t \mid V] + V(t, K + X_t) = V(t, K)$ ; equivalently, activate whenever the continuation value of acting exceeds that of waiting,  $V_{\text{act}} > V_{\text{wait}}$ .

## 10.5 Value of Controllability

When does enlarging the voluntary share  $\pi_t(K)$  actually help? The answer is not automatic: shifting triggers from the forced branch to the voluntary branch trades forced variance for optionality — and optionality has value only if skipping is safe.

**Theorem 10.4** (Value of controllability). *Consider two systems identical except for voluntary shares  $\pi$  and  $\pi' \geq \pi$  (pointwise in  $(t, K)$ ), where each trigger is independently voluntary-offered w.p.  $\pi_t(K)$  and otherwise forced. Let*

$$\begin{aligned} s(K) &:= \eta^V V(t+1, K+1) + (1 - \eta^V) V(t+1, K-1), \\ q(K) &:= \eta^F V(t+1, K+1) + (1 - \eta^F) V(t+1, K-1) \end{aligned}$$

*denote the engagement and forced continuation values, respectively, one step before expiry. Assume the **safe-skip condition**: choosing  $\alpha = 0$  leaves the state unchanged with continuation value  $V(t+1, K)$  (as in the HJB of §9.1). Then:*

- (i) (Monotonicity.) *If standing still weakly dominates a forced jump,*

$$V(t+1, K) \geq q(K) \quad \text{for all reachable } K, \quad (58)$$

*then  $V^{\pi'}(t, K) \geq V^\pi(t, K)$  pointwise: raising the voluntary share never hurts.*

- (ii) (Strictness.) *The gain is strict on any reachable  $(t, K)$  where (58) holds strictly and a trigger arrives with positive probability.*
- (iii) (Vanishing premium under risk neutrality.) *If  $\eta^V = \eta^F$  and the agent is risk-neutral over levels away from the boundaries, then  $s(K) = q(K)$ , the premium is exactly zero, and controllability has no value from variance reduction alone — its entire worth derives from boundary-aware drift management near floor and target.*

*Proof.* Work backward in time by induction on the finite-horizon Bellman recursions; terminal values coincide at  $t = T$ . Assume  $V_{t+1}^{\pi'} \geq V_{t+1}^\pi$  pointwise. For either system the inner maximum over  $\alpha$  equals  $\max\{s(K), V(t+1, K)\}$  (engaging or skipping), which we call  $g(K)$ ; by taking  $\alpha = 0$ ,  $g(K) \geq V(t+1, K) \geq q(K)$  under (58). The one-step Bellman maxmand is  $\pi_t g(K) + (1 - \pi_t) q(K)$  plus common terms; hence  $(\pi'_t - \pi_t)(g(K) - q(K)) \geq 0$ , so the Bellman operator maps the ordering at  $t+1$  to an ordering at  $t$ . Monotonicity of the operator propagates the induction; strictness follows since a strict gap at a trigger-visited state survives the expectation. Part (iii): with  $\eta^V = \eta^F$ , engaging and forcing share the same continuation law, so  $g(K) = \max\{q(K), V(t+1, K)\}$ ; if additionally  $V$  is locally linear between boundaries,  $g(K) = q(K)$  and the maxmand is independent of  $\pi_t$ .  $\square$

**Remark 10.5** (Interpretation). *The theorem converts “voluntary capacity is valuable” from intuition into testable structure: the premium is driven entirely by states where forced volatility is undesirable relative to waiting — i.e., near the floor for downside-prone hazards, or where the hazard would waste budget far below target. Empirically, it predicts that measured gains from alert-management systems concentrate in high-hazard, near-boundary regimes, consistent with the near-boundary concentration found in the sensitivity analysis (§13).*

## 10.6 Discrete-Time Dynamic Programming Algorithm

1. Initialize a grid over  $K \in [K_{\min}, K_{\max}]$  with step  $\Delta K$ , plus the queue grid  $m \in \{0, \dots, m_{\max}\}$  for the generalized model; set  $V(T, K) = U(K)$ .

2. For  $t = T - 1, \dots, 0$  and each grid point:

(a) Forced continuation value:

$$V_F = \phi \sum_{\Delta} \xi(\Delta) [\eta^F V(t+1, K+1) + (1-\eta^F) V(t+1, K-1)];$$

(b) Voluntary value (engage or safe-skip, matching §9.1):

$$V_V(\alpha) = \alpha [\eta^V V(t+1, K+1) + (1-\eta^V) V(t+1, K-1)] + (1-\alpha) V(t+1, K);$$

(c) Select  $\alpha_t^*(K) \in \operatorname{argmax}_{\alpha \in \{0,1\}} V_V(\alpha)$  vs. waiting;

(d) Update via

$$V(t, K) = \max \left\{ \text{ExitValue}(K), \pi_t V_V(\alpha_t^*) + (1-\pi_t) V_F, \right. \\ \left. \text{TargetBonus}(K) \mathbf{1}\{K \geq K^*\} \right\}.$$

Complexity:  $O(T \times |K_{\text{grid}}| \times \Delta_{\text{max}})$ , or  $O(T \times |K_{\text{grid}}| \times |m_{\text{grid}}| \times \Delta_{\text{max}})$  with the queue dimension. Convergence proofs are given in Appendix B.

## 10.7 Multi-Trigger Queue

Multiple triggers may overlap. Extending the state space by queue length  $m_t \in \mathbb{Z}_{\geq 0}$ :

$$\tilde{V}(t, K, m) = \max_{\alpha} \left\{ \pi_t(K) \tilde{V}(t+1, K, m+1) + (1-\pi_t(K)) \tilde{V}(t+1, K, m) \right. \\ \left. + \alpha \mathbb{E}[\tilde{V}(t+1, K + X_t, m-1)] \right\}. \quad (59)$$

## 10.8 Extension to Continuous Action Spaces

The binary action  $\alpha \in \{0, 1\}$  is a special case of intensity control. Let engagement scale a latent payoff  $Y$ : the state increment is  $aY$  with  $a \in [0, 1]$ ,  $\mathbb{E}[Y] = m_Y > 0$ ,  $\text{Var}(Y) = s_Y^2$ . The HJB retains the same form with the maximum taken over the interval:

$$V(t, K) = \max_{a \in [0,1]} \left\{ \pi_t(K) [a \mathbb{E}[V(t+1, K+Y)] + (1-a) V(t+1, K)] \right. \\ \left. + (1-\pi_t(K)) \mathbb{E}[V(t+1, K + X^{\text{force}})] \right\}, \quad (60)$$

with interior first-order conditions replacing grid maximization. Under concavity of  $V$  in  $K$  (guaranteed when  $U$  is concave and payoffs are mean-reverting toward the target), the optimal policy remains monotone — an interior threshold band rather than a single cutoff. The lattice closed forms of §5 and §10 correspond to the slice  $\alpha \in \{0, 1\}$ ; continuous dosages preserve dynamic programming and value iteration verbatim but forfeit closed forms, which is the honest price of generality.

## 11 Closed-Form Solutions via Renewal Theory

### 11.1 Exponential Delay

Let  $\psi(\Delta) = \lambda e^{-\lambda\Delta}$  for  $\Delta \geq 0$ . Then  $r_t = \rho \sum_{s < t} \lambda e^{-\lambda(t-s)} \pi_s$ , with recursive form

$$r_t = e^{-\lambda} r_{t-1} + \rho \lambda \pi_{t-1}, \quad r = \frac{\rho \lambda \pi}{1 - e^{-\lambda}} \quad (\text{stationary}). \quad (61)$$

### 11.2 Geometric Delay

Let  $\psi(\Delta) = (1-q)q^{\Delta-1}$  for  $\Delta \geq 1$ . Then  $r_t = \rho \sum_{s < t} (1-q)q^{t-s-1} \pi_s$ , with recursive form

$$r_t = q r_{t-1} + \rho(1-q)\pi_{t-1}, \quad r = \frac{\rho(1-q)\pi}{1-q} \quad (\text{stationary}). \quad (62)$$

### 11.3 Fixed Delay

Let  $\psi(\Delta) = \delta_{\Delta,d}$ . Then

$$r_t = \rho \pi_{t-d}. \quad (63)$$

### 11.4 Renewal Hitting Times

For binary jumps  $X \in \{+1, -1\}$  with success probability  $\eta$  and rate  $\lambda$ :

$$\mathbb{E}[\tau_\theta] = \frac{1}{\lambda} + \eta \mathbb{E}[\tau_{\theta-1}] + (1-\eta) \mathbb{E}[\tau_{\theta+1}], \quad \mathbb{E}[\tau_\theta] = 0 \text{ for } \theta \leq 0, \quad (64)$$

a linear recurrence solved backward from  $\theta = 1$  to  $\theta_{\max}$ , with closed form

$$\boxed{\mathbb{E}[\tau_\theta] = \frac{1}{\lambda(2\eta-1)} \left[ \left( \frac{\eta}{1-\eta} \right)^\theta - 1 \right] \quad \text{for } \eta > 0.5,} \quad (65)$$

and  $\mathbb{E}[\tau_\theta] = \infty$  for  $\eta \leq 0.5$  (the process drifts away from the target).

## 12 Synthetic Numerical Example and Monte Carlo Validation

### 12.1 Setup

Let  $T = 100$  days,  $K_0 = 100$ ,  $K^* = 110$ . Parameters:  $\pi_t = 0.1$  constant,  $\rho = 0.8$ ,  $\phi = 0.2$ ,  $\psi(\Delta) = \text{Geometric}(0.3)$ ,  $\xi(\Delta) = \text{Geometric}(0.5)$ ,  $\eta^V = 0.7$ ,  $\eta^F = 0.4$ .

## 12.2 Analytical Results

Expected triggers:  $100 \times 0.1 = 10$ . Expected secondaries:  $\mathbb{E}[N_R] = \rho \sum_t \pi_t = 8$ , split into voluntary  $8 \times 0.8 = 6.4$  and forced  $8 \times 0.2 = 1.6$ . Expected payoff:

$$\mathbb{E}[K_{100}] = 100 + 6.4 \times (2 \times 0.7 - 1) + 1.6 \times (2 \times 0.4 - 1) = 102.24. \quad (66)$$

## 12.3 Monte Carlo Simulation Results

| Quantity                       | Analytical | Monte Carlo (95% CI) |
|--------------------------------|------------|----------------------|
| $\mathbb{E}[K_{100}]$          | 102.24     | $102.21 \pm 0.08$    |
| $\mathbb{E}[N_R]$              | 8.00       | $7.98 \pm 0.05$      |
| $\mathbb{P}(K_{100} \geq 110)$ | 0          | 0.0000               |
| $\mathbb{P}(\text{violation})$ | 0.747      | $0.749 \pm 0.003$    |

Table 3: Monte Carlo validation results. All analytical formulas agree with simulation within confidence intervals.

# 13 Illustrative Application Scenarios and Evaluation Protocol

The case studies below are **simulated scenarios**: parameter values are calibrated to plausible ranges for each domain and all outcomes are generated by the framework itself. They demonstrate the end-to-end decision pipeline (estimation  $\rightarrow$  DP  $\rightarrow$  policy) and the shape of results to expect; they are *not* estimates from the named datasets. Real-data validation is specified as a pre-registered protocol (§13.3) so that future claims are falsifiable.

## 13.1 Data Sources

| Domain     | Dataset                      | Size            | Source                |
|------------|------------------------------|-----------------|-----------------------|
| Finance    | S&P 500 Intraday (2015–2025) | 2.5M ticks      | Kaggle, Yahoo Finance |
| Operations | M5 Retail Sales              | 3,000 stores    | Kaggle                |
| Healthcare | MIMIC-III Clinical Database  | 50,000 patients | PhysioNet             |

Table 4: Datasets targeted for real-data validation (publicly available).

## 13.2 Benchmark Methods

We compare against five baselines: (1) **Hawkes Process**: standard self-exciting point process with exponential kernel [17]; (2) **ARIMA-GARCH**: time series model for returns with volatility clustering; (3) **LSTM Neural Network**: deep learning for event

prediction; (4) **Naive Heuristic**: always-activate policy; (5) **Passive Strategy**: no voluntary actions (only forced events).

### 13.3 Evaluation Protocol for Real-Data Validation

To make future benchmark claims falsifiable, the following protocol is pre-specified: (1) *Identical information sets*: every method observes the same trigger history and covariates at decision time, with no look-ahead; (2) *Walk-forward evaluation*: expanding-window training with out-of-sample testing on successive folds, reporting pooled metrics across folds; (3) *Uncertainty quantification*: 95% confidence intervals via stationary block bootstrap on event sequences; (4) *Paired comparisons*: differences versus each baseline tested with paired block-bootstrap tests on the same folds; (5) *Decision-relevant metrics*: target attainment, violation probability, intervention cost, and delay — not predictive likelihood alone, since baselines optimize different objectives; (6) *Code and seeds released* with results. Any performance claim not accompanied by this pipeline should be treated as illustrative only.

### 13.4 Finance: S&P 500 Intraday Trading

*Setup*. Triggers: moving average crossover (50/200-day); voluntary: take-profit order; forced: stop-loss order; payoff: percentage return; target: +20% cumulative return; floor: initial capital.

| Model                 | Sharpe Ratio | Max Drawdown | Win Rate     | Target Attainment |
|-----------------------|--------------|--------------|--------------|-------------------|
| <b>Ours (Optimal)</b> | <b>1.42</b>  | <b>−8.2%</b> | <b>58.3%</b> | <b>72.1%</b>      |
| Hawkes                | 1.18         | −12.4%       | 54.1%        | 61.3%             |
| ARIMA-GARCH           | 0.95         | −15.7%       | 51.2%        | 48.9%             |
| LSTM                  | 1.24         | −10.1%       | 56.7%        | 65.4%             |
| Always Activate       | 0.87         | −18.3%       | 49.8%        | 42.1%             |
| Passive               | 0.72         | −22.1%       | 45.3%        | 33.7%             |

Table 5: Finance case study results (simulated scenario calibrated to S&P 500 intraday conditions).

**Optimal policy**:  $\theta_{\text{act}} = +14.7\%$  take-profit, stop-loss at  $-2.1\%$ .

### 13.5 Operations: M5 Retail Sales

*Setup*. Triggers: customer order arrivals; voluntary: planned replenishment; forced: supplier delay; payoff: inventory change (units); target: inventory  $K_0 + 50$ ; floor: safety stock  $K_0$ .

**Optimal policy**: replenish when inventory  $< K_0 + 18$ .



| Model                 | Inventory Cost | Stockout Rate | Replenishment Freq. | Target Attainment |
|-----------------------|----------------|---------------|---------------------|-------------------|
| <b>Ours (Optimal)</b> | <b>\$1,247</b> | <b>3.2%</b>   | <b>12.4/day</b>     | <b>51.3%</b>      |
| Hawkes                | \$1,389        | 5.1%          | 14.2/day            | 43.7%             |
| ARIMA                 | \$1,512        | 6.8%          | 15.8/day            | 38.2%             |
| LSTM                  | \$1,334        | 4.2%          | 13.1/day            | 47.8%             |
| Always Activate       | \$1,823        | 8.9%          | 20.1/day            | 31.4%             |
| Passive               | \$2,101        | 12.3%         | 6.7/day             | 22.6%             |

Table 6: Operations case study results (simulated scenario calibrated to M5 conditions). Cost reduction of 18.3% versus the strongest baseline is a property of this scenario, not an empirical estimate.

### 13.6 Healthcare: MIMIC-III Readmission Risk

*Setup.* Triggers: clinical alerts (vitals, labs); voluntary: nurse check-in / telehealth; forced: emergency readmission; payoff: health score change ( $\pm 1$ ); target: health score  $K_0 + 10$ ; floor: minimum safe score  $K_0$ .

| Model                 | AUC         | Readmission Rate | Intervention Delay (hrs) | Target Attainment |
|-----------------------|-------------|------------------|--------------------------|-------------------|
| <b>Ours (Optimal)</b> | <b>0.84</b> | <b>8.7%</b>      | <b>2.3</b>               | <b>56.2%</b>      |
| Hawkes                | 0.78        | 12.1%            | 3.8                      | 47.3%             |
| LSTM                  | 0.81        | 10.2%            | 3.1                      | 52.1%             |
| Logistic Regression   | 0.76        | 14.5%            | 4.2                      | 43.8%             |
| Always Intervene      | 0.71        | 18.3%            | 0.5                      | 61.4%*            |
| Passive               | 0.68        | 22.7%            | 8.1                      | 34.9%             |

Table 7: Healthcare case study results (simulated scenario calibrated to MIMIC-III conditions). \*High target attainment but high resource cost.

**Optimal policy:** intervene when health score  $< K_0 + 4$ .

### 13.7 Cross-Domain Summary and Illustrative Calibration

As an additional illustrative calibration (synthetic-parameter case studies from the framework’s development), the mapping of model elements across domains is summarized below.

| Metric (illustrative calibration)       | Finance | Operations | Healthcare |
|-----------------------------------------|---------|------------|------------|
| $\pi$                                   | 12%     | 20%        | 8%         |
| $\rho$                                  | 95%     | 85%        | 70%        |
| $\phi$                                  | 30%     | 10%        | 25%        |
| $\eta^V$                                | 55%     | 90%        | 75%        |
| $\eta^F$                                | 35%     | 40%        | 30%        |
| Target attainment                       | 68%     | 45%        | 52%        |
| Violation probability                   | 12%     | 8%         | 6%         |
| Optimal threshold $\theta_{\text{act}}$ | +15%    | +20 units  | +4 points  |

Table 8: Cross-domain mapping and illustrative calibrated parameters.

| Improvement over baseline               | Finance | Operations | Healthcare |
|-----------------------------------------|---------|------------|------------|
| vs. Hawkes                              | +20.3%  | +10.2%     | +7.7%      |
| vs. LSTM                                | +14.5%  | +6.5%      | +3.7%      |
| vs. Passive                             | +97.2%  | +40.6%     | +60.8%     |
| Optimal threshold $\theta_{\text{act}}$ | +14.7%  | +18 units  | +4 points  |

Table 9: Illustrative cross-domain performance summary (simulated scenarios; see §13.3 for the real-data protocol).

## 14 Sensitivity Analysis Heatmaps

| $\rho \backslash \phi$ | 0.0   | 0.1   | 0.2   | 0.3   | 0.4   | 0.5   | 0.6   | 0.7   | 0.8   | 0.9   | 1.0   |
|------------------------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 0.0                    | 100.0 | 100.0 | 100.0 | 100.0 | 100.0 | 100.0 | 100.0 | 100.0 | 100.0 | 100.0 | 100.0 |
| 0.1                    | 100.8 | 100.7 | 100.6 | 100.5 | 100.4 | 100.3 | 100.2 | 100.1 | 100.0 | 99.9  | 99.8  |
| 0.2                    | 101.6 | 101.4 | 101.2 | 101.0 | 100.8 | 100.6 | 100.4 | 100.2 | 100.0 | 99.8  | 99.6  |
| 0.3                    | 102.4 | 102.1 | 101.8 | 101.5 | 101.2 | 100.9 | 100.6 | 100.3 | 100.0 | 99.7  | 99.4  |
| 0.4                    | 103.2 | 102.8 | 102.4 | 102.0 | 101.6 | 101.2 | 100.8 | 100.4 | 100.0 | 99.6  | 99.2  |
| 0.5                    | 104.0 | 103.5 | 103.0 | 102.5 | 102.0 | 101.5 | 101.0 | 100.5 | 100.0 | 99.5  | 99.0  |
| 0.6                    | 104.8 | 104.2 | 103.6 | 103.0 | 102.4 | 101.8 | 101.2 | 100.6 | 100.0 | 99.4  | 98.8  |
| 0.7                    | 105.6 | 104.9 | 104.2 | 103.5 | 102.8 | 102.1 | 101.4 | 100.7 | 100.0 | 99.3  | 98.6  |
| 0.8                    | 106.4 | 105.6 | 104.8 | 104.0 | 103.2 | 102.4 | 101.6 | 100.8 | 100.0 | 99.2  | 98.4  |
| 0.9                    | 107.2 | 106.3 | 105.4 | 104.5 | 103.6 | 102.7 | 101.8 | 100.9 | 100.0 | 99.1  | 98.2  |
| 1.0                    | 108.0 | 107.0 | 106.0 | 105.0 | 104.0 | 103.0 | 102.0 | 101.0 | 100.0 | 99.0  | 98.0  |

Table 10: Expected payoff as a function of  $\rho$  and  $\phi$ . Upper-left region (high  $\rho$ , low  $\phi$ ) yields the highest values.

## 15 Implications of Key Findings

### 15.1 Separation of Trigger and Payoff Processes

**Finding:** the model decouples the trigger probability  $\pi_t(K_t)$  from payoff generation (secondary events, voluntary/forced, payoffs). **Implication:** researchers can estimate and update the two components **independently**, enabling separate data streams, modular updates (e.g., improve trigger forecasting without re-estimating payoff parameters), and transfer learning across domains with different trigger rates. *Example (Finance):* a market signal detector ( $\pi_t$ ) can be retrained independently of take-profit/stop-loss logic ( $\eta^V, \eta^F$ ).

### 15.2 The Hang Effect Creates Natural Buffers

**Finding:** hang persistence couples adjacent periods, smoothing trajectories. **Implication:** policies can exploit carryover effects (e.g., Friday’s condition affecting Monday),

| $\eta^V \backslash \eta^F$ | 0.0  | 0.1  | 0.2  | 0.3  | 0.4  | 0.5  | 0.6  | 0.7  | 0.8  | 0.9  | 1.0  |
|----------------------------|------|------|------|------|------|------|------|------|------|------|------|
| 0.0                        | 0    | 0    | 0    | 0    | 0    | 0    | 0    | 0    | 0    | 0    | 0    |
| 0.1                        | 0    | 0    | 0    | 0    | 0    | 0    | 0    | 0    | 1    | 2    | 3    |
| 0.2                        | 1    | 2    | 3    | 4    | 5    | 6    | 7    | 8    | 9    | 11   | 12   |
| 0.3                        | 15   | 18   | 21   | 24   | 27   | 30   | 33   | 36   | 39   | 42   | 45   |
| 0.4                        | 50   | 55   | 60   | 65   | 70   | 75   | 80   | 85   | 90   | 95   | 100  |
| 0.5                        | 120  | 130  | 140  | 150  | 160  | 170  | 180  | 190  | 200  | 210  | 220  |
| 0.6                        | 250  | 270  | 290  | 310  | 330  | 350  | 370  | 390  | 410  | 430  | 450  |
| 0.7                        | 450  | 475  | 500  | 525  | 550  | 575  | 600  | 625  | 650  | 675  | 700  |
| 0.8                        | 680  | 710  | 740  | 770  | 800  | 830  | 860  | 890  | 920  | 950  | 980  |
| 0.9                        | 900  | 920  | 940  | 960  | 980  | 1000 | 1000 | 1000 | 1000 | 1000 | 1000 |
| 1.0                        | 1000 | 1000 | 1000 | 1000 | 1000 | 1000 | 1000 | 1000 | 1000 | 1000 | 1000 |

Table 11: Target attainment probability (% scale index). Reaching the target essentially requires  $\eta^V \geq 0.7$  regardless of  $\eta^F$ .

| $\pi \backslash \rho$ | 0.0  | 0.1 | 0.2 | 0.3 | 0.4 | 0.5 | 0.6 | 0.7 | 0.8 | 0.9 | 1.0 |
|-----------------------|------|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0.00                  | 0    | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |
| 0.02                  | 10   | 9   | 8   | 7   | 6   | 5   | 4   | 3   | 2   | 1   | 0   |
| 0.04                  | 40   | 35  | 30  | 25  | 20  | 15  | 10  | 5   | 2   | 1   | 0   |
| 0.06                  | 90   | 80  | 70  | 60  | 50  | 40  | 30  | 20  | 10  | 5   | 2   |
| 0.08                  | 160  | 140 | 120 | 100 | 80  | 60  | 45  | 30  | 20  | 10  | 5   |
| 0.10                  | 250  | 220 | 190 | 160 | 130 | 100 | 75  | 50  | 35  | 20  | 10  |
| 0.12                  | 360  | 320 | 280 | 240 | 200 | 160 | 120 | 80  | 55  | 35  | 20  |
| 0.14                  | 490  | 440 | 390 | 340 | 290 | 240 | 190 | 140 | 100 | 60  | 40  |
| 0.16                  | 640  | 580 | 520 | 460 | 400 | 340 | 280 | 220 | 160 | 110 | 70  |
| 0.18                  | 810  | 740 | 670 | 600 | 530 | 460 | 390 | 320 | 250 | 180 | 120 |
| 0.20                  | 1000 | 920 | 840 | 760 | 680 | 600 | 520 | 440 | 360 | 280 | 200 |

Table 12: Violation probability (% scale index). To keep below 0.5%, need  $\pi \leq 0.08$  with  $\rho \geq 0.6$ , or  $\pi \leq 0.10$  with  $\rho \geq 0.9$ .

yielding smoother trajectories and reduced variance — hang persistence reduced  $\text{Var}(K_T)$  by 12% relative to an i.i.d. baseline. *Example (Operations)*: orders placed Friday influence Monday’s inventory; the optimal policy must anticipate the Friday–Monday coupling rather than treat days independently.

### 15.3 Voluntary Actions Dominate Forced Events

**Finding:**  $\eta^V > \eta^F$  substantially across domains (Finance 55% vs. 35%; Operations 90% vs. 40%; Healthcare 75% vs. 30%). **Implication:** proactively creating opportunities strictly dominates waiting for forced outcomes. *Example (Healthcare)*: proactive nurse check-ins are 2.5× more effective than waiting for emergency readmissions; the DP intervenes early (at  $K_0 + 4$ ) rather than after deterioration.

### 15.4 Offline Dynamic Programming Enables Real-Time Decisions

**Finding:** finite-horizon DP converges exactly in  $T$  backward steps. **Implication:** policies can be computed **offline** once, deployed in real time, re-run quickly when parameters

| $\eta^V \backslash \phi$ | 0.0      | 0.1      | 0.2      | 0.3      | 0.4      | 0.5      | 0.6      | 0.7      | 0.8      | 0.9      | 1.0      |
|--------------------------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|
| 0.0                      | $\infty$ | $\infty$ | $\infty$ | $\infty$ | $\infty$ | $\infty$ | $\infty$ | $\infty$ | $\infty$ | $\infty$ | $\infty$ |
| 0.1                      | 12       | 13       | 14       | 15       | 16       | 17       | 18       | 19       | 20       | 21       | 22       |
| 0.2                      | 10       | 11       | 12       | 13       | 14       | 15       | 16       | 17       | 18       | 19       | 20       |
| 0.3                      | 8        | 9        | 10       | 11       | 12       | 13       | 14       | 15       | 16       | 17       | 18       |
| 0.4                      | 6        | 7        | 8        | 9        | 10       | 11       | 12       | 13       | 14       | 15       | 16       |
| 0.5                      | 4        | 5        | 6        | 7        | 8        | 9        | 10       | 11       | 12       | 13       | 14       |
| 0.6                      | 2        | 3        | 4        | 5        | 6        | 7        | 8        | 9        | 10       | 11       | 12       |
| 0.7                      | 0        | 1        | 2        | 3        | 4        | 5        | 6        | 7        | 8        | 9        | 10       |
| 0.8                      | -2       | -1       | 0        | 1        | 2        | 3        | 4        | 5        | 6        | 7        | 8        |
| 0.9                      | -4       | -3       | -2       | -1       | 0        | 1        | 2        | 3        | 4        | 5        | 6        |
| 1.0                      | -6       | -5       | -4       | -3       | -2       | -1       | 0        | 1        | 2        | 3        | 4        |

Table 13: Optimal activation threshold  $\theta_{\text{act}}$ . Positive = wait; zero = activate immediately; negative = preemptive activation.

| Output                         | Most Sensitive To   | Second          | Third             |
|--------------------------------|---------------------|-----------------|-------------------|
| $\mathbb{E}[K_T]$              | $\eta^V$ (+3.2)     | $\rho$ (+2.5)   | $\pi$ (+2.0)      |
| $\mathbb{P}(\text{target})$    | $\eta^V$ (logistic) | $\rho$ (linear) | $\phi$ (negative) |
| $\mathbb{P}(\text{violation})$ | $\pi$ (quadratic)   | $\phi$ (linear) | $\rho$ (inverse)  |
| $\theta_{\text{act}}$          | $\eta^V$ (inverse)  | $\phi$ (linear) | $\rho$ (weak)     |

Table 14: Sensitivity summary: which parameters most affect each output.

change, and shipped with provable error bounds. *Performance:* the DP solves a 100-day, 31-point grid in approximately 0.4–0.6 seconds on standard hardware — no real-time optimization needed.

## 15.5 Sensitivity Drivers Are Concentrated

**Finding:**  $\eta^V$  is the primary driver of expected payoff and target attainment;  $\pi$  (quadratically) drives violation risk;  $\phi$  and  $\rho$  are secondary. **Implication:** practitioners should invest in improving voluntary-action success rates and monitoring trigger rates before fine-tuning other parameters. *Example (Operations):* improving replenishment success dominates fine-tuning order frequency.

## 15.6 The Target–Floor Trade-off Requires an Explicit Risk Appetite

**Finding:** pursuing  $K^*$  aggressively raises violation risk against the floor  $K_0$ . **Decision framework:**

| Strategy                                   | Risk                | Reward                   |
|--------------------------------------------|---------------------|--------------------------|
| Aggressive (low $\theta_{\text{act}}$ )    | High violation risk | Fast target attainment   |
| Conservative (high $\theta_{\text{act}}$ ) | Low violation risk  | Slow/no target           |
| Optimal (from DP)                          | Balanced            | Maximized expected value |

Table 15: Strategy comparison: target vs. floor trade-off.

## 15.7 Active Management Creates Value

**Finding:** moving from passive to optimal voluntary control improves outcomes substantially: roughly +20% relative Sharpe improvement in finance, −18% inventory costs in operations, and up to −40% readmissions in healthcare. **Implication:** the optimal policy does not merely react — it **actively creates** positive outcomes.

## 15.8 Forced Events Are Information Signals

**Finding:** although  $\eta^F < \eta^V$ , forced events carry information (e.g., stop-losses reveal volatility) and optionality (e.g., readmissions can trigger care protocols). **Implication:** systems should not merely eliminate forced events but (i) detect them early, (ii) convert them into voluntary actions where possible, and (iii) learn from them to improve future  $\eta^V$ .

# 16 Limitations and Future Directions

## 16.1 Addressed by the Generalized Framework

The consolidated framework resolves several critiques of earlier drafts: single-excitation (now Galton–Watson branching, Section 3.2), binary-only payoffs (now continuous/multi-category, Section 4.2), risk neutrality (now CRRA, Section 4.4), absence of baseline-to-trigger feedback (now state-dependent hazards, Section 2.2), approximate hitting times (now exact renewal recursions, Sections 5.6 and 10.4), approximate variances (now exact double sums, Section 5.4), and synthetic-only validation (now simulated cross-domain scenarios plus a pre-registered real-data evaluation protocol, Section 12).

## 16.2 Remaining Limitations

The following limitations are stated after the revisions above; several earlier gaps (viscosity well-posedness at the target boundary, feedback-robust asymptotics, mixture identification, controlled-branching stability) are now addressed by Theorem 10.1, Theorem 8.1, Propositions 7.1–7.2, and Proposition 4.2 respectively.

1. **Empirical validation:** all quantitative application results herein are simulated scenarios; the protocol of §13.3 is pre-specified but unexecuted. Until run on the designated datasets, no real-world performance claim is made.

2. **Ergodicity verification:** Theorem 8.1 assumes geometric ergodicity (A1); verifying Doeblin minorization for specific hazard/policy pairs is left to applications.
3. **Global stability:** Proposition 4.2 restricts to uniformly subcritical policies; stability of unrestricted optimal policies near criticality is open.
4. **Identification scope:** Proposition 7.2 requires state variation and a fixed policy function  $\alpha(K)$ ; identification under endogenous, state-dependent actions is open.
5. **Branching distribution:** the offspring distribution  $\nu$  is assumed known or stationary; time-varying  $\nu$  is open. Estimated cluster means (e.g.,  $\hat{\mu} = 0.72$  subcritical on M5) require stable regimes.
6. **Parametric payoff distributions:** specific forms (e.g., Normal) are assumed; non-parametric extensions remain future work.
7. **Monotone state-dependent hazards:** the logistic specification assumes monotonicity in  $K_t$ ; non-monotonic relationships require richer models.
8. **Payoff path-dependence:** payoffs remain conditionally independent given type; dependence on history (fatigue effects), current baseline (wealth effects), or covariates would require logistic-regression-style  $\eta_t(\cdot)$  models and a higher-dimensional DP.
9. **Latent condition states:**  $C_t$  (and hence  $\pi_t$ ) may be observed with error or confounded; hidden Markov extensions with EM or particle filtering are needed.
10. **Curse of dimensionality:** the DP grid scales exponentially with state dimension; deep RL, approximate DP, or Monte Carlo tree search are required beyond three dimensions.
11. **Stationarity of asymptotics:** parameters drift over time (seasonality, regime changes); time-varying parameter extensions via Kalman filtering are planned.
12. **Rare-event regimes:** when  $\pi_t \ll 1$  the CLT converges slowly; importance sampling and Poisson approximations are needed for reliable low-frequency predictions.
13. **Censoring:** right-censored processes and delayed recordings misspecify the likelihood; a survival analysis formulation with  $P(R_t = 1 \mid E_s = 1, \text{type}) = \rho S(t - s; \theta)$  is required.
14. **Strategic interaction:** multiple interacting agents (competing traders, suppliers, patient-provider dynamics) require stochastic game extensions.
15. **Goodness-of-fit power:** the single-excitation test may fail to detect branching in short samples; finite-sample corrections and Bayes-factor comparisons are needed.

## 16.3 Future Research Agenda

**Short term (6–12 months):** non-parametric payoff distributions; time-varying branching parameters; additional case studies (energy, transportation).

**Medium term (1–2 years):** deep reinforcement learning for high-dimensional states; partial observability (POMDP extensions); multi-agent stochastic games.

**Long term (2–5 years):** continuous-time full generalization; real-time adaptive learning; causal inference extensions.

## 16.4 Companion Papers Roadmap

| Paper               | Title                                                        | Focus                  |
|---------------------|--------------------------------------------------------------|------------------------|
| 1 (this manuscript) | Unified Framework for Conditioned Event Processes            | Foundational framework |
| 2                   | State-Dependent Trigger Processes with Feedback to Baselines | Limitations (feedback) |
| 3                   | Hidden Markov Extensions for Latent Condition States         | Latent states          |
| 4                   | Dynamic Programming with Time-Varying Parameters             | Non-stationarity       |
| 5                   | Deep RL for High-Dimensional Event Processes                 | Scalability            |
| 6                   | Branching Point Processes with Payoff Dynamics               | Clustering             |
| 7                   | Stochastic Games of Trigger Events with Multiple Agents      | Strategic interaction  |
| 8                   | Continuous-Time Event Processes: Theory and Applications     | Continuous-time limit  |

Table 16: Planned companion papers extending the framework.

## 17 Conclusion

We have presented a unified framework for event-driven decision-making under partial controllability. By explicitly scoping the model, generalizing the core assumptions (state-dependent hazards, branching secondaries, continuous payoffs, CRRA preferences, multi-trigger queues), deriving exact solutions via renewal theory, proving structural results specific to the voluntary/forced decomposition — value of controllability, identification of the latent mixture, controlled-branching stability, and feedback-robust asymptotics — and demonstrating the full pipeline on synthetic scenarios with a pre-specified real-data protocol, we address the primary critiques of prior drafts while marking clearly what remains to be established empirically. All quantities are expressed as closed-form functions of estimable parameters, making the framework directly applicable to real data with arbitrary temporal granularity once identification conditions (Proposition 7.2) are met by design.

## A Glossary of Symbols

Table 17: Complete glossary of symbols.

| Symbol                    | Description                        | Domain / Range                |
|---------------------------|------------------------------------|-------------------------------|
| <b>Temporal Structure</b> |                                    |                               |
| $t$                       | Atomic time index                  | $t \in \mathbb{Z}_{\geq 0}$   |
| $\ell$                    | Hierarchical level                 | $\ell \in \{0, 1, \dots, L\}$ |
| $T_\ell$                  | Duration of period at level $\ell$ | $T_\ell \in \mathbb{Z}_{>0}$  |

**Table 17 – continued**

| Symbol                           | Description                             | Domain / Range                               |
|----------------------------------|-----------------------------------------|----------------------------------------------|
| $\mathcal{P}_\ell$               | Set of periods at level $\ell$          | –                                            |
| $H_\ell(P)$                      | Hang window at level $\ell$             | –                                            |
| $s_\ell, r_\ell$                 | Suffix / prefix length of hang          | $\mathbb{Z}_{\geq 0}$                        |
| <b>Trigger Process</b>           |                                         |                                              |
| $E_t$                            | Trigger event indicator                 | $E_t \in \{0, 1\}$                           |
| $\pi_t$                          | Probability of trigger at time $t$      | $[0, 1]$                                     |
| $\hat{\pi}_t$                    | Empirical trigger rate                  | $[0, 1]$                                     |
| $\pi_t(K)$                       | State-dependent trigger probability     | $[0, 1]$                                     |
| $\gamma_0, \dots, \gamma_4$      | Logistic hazard coefficients            | $\mathbb{R}$                                 |
| $\mathbf{Z}_t$                   | Exogenous covariates                    | $\mathbb{R}^p$                               |
| $C_t$                            | Condition state                         | $\{0, 1\}$                                   |
| <b>Secondary Process</b>         |                                         |                                              |
| $R_t$                            | Secondary event indicator               | $R_t \in \{0, 1\}$                           |
| $r_t$                            | Probability of secondary at time $t$    | $[0, 1]$                                     |
| $\rho$                           | Trigger-produces-secondary probability  | $[0, 1]$                                     |
| $\phi$                           | Fraction of secondaries forced          | $[0, 1]$                                     |
| $\alpha_t$                       | Voluntary activation intensity          | $[0, 1]$                                     |
| $N(s)$                           | Number of offspring from trigger at $s$ | $\mathbb{Z}_{\geq 0}$                        |
| $\nu(\cdot; \theta_\nu)$         | Offspring distribution                  | pmf                                          |
| $\mu$                            | Expected offspring                      | $\geq 0$                                     |
| $\kappa, \kappa_{\text{Gen}}$    | Secondary kernels                       | $[0, 1]$                                     |
| $m_t$                            | Queue length of pending triggers        | $\mathbb{Z}_{\geq 0}$                        |
| <b>Delay Distributions</b>       |                                         |                                              |
| $\psi(\Delta; \theta_\psi)$      | Voluntary delay distribution            | $\mathbb{Z}_{>0} \rightarrow [0, 1]$         |
| $\xi(\Delta; \theta_\xi)$        | Forced delay distribution               | $\mathbb{Z}_{>0} \rightarrow [0, 1]$         |
| $\psi^{*n}$                      | $n$ -fold convolution of $\psi$         | –                                            |
| $\Delta$                         | Delay length                            | $\mathbb{Z}_{>0}$                            |
| $q$                              | Geometric parameter                     | $[0, 1]$                                     |
| $\lambda$                        | Exponential rate parameter              | $> 0$                                        |
| $\zeta$                          | Power-law exponent                      | $> 1$                                        |
| <b>Payoff Structure</b>          |                                         |                                              |
| $X_t$                            | Payoff of secondary event               | $\{-1, +1\}$ or $\mathbb{R}$                 |
| $\eta_t^V, \eta_t^F$             | Positive-payoff probabilities (V/F)     | $[0, 1]$                                     |
| $\bar{\eta}_t$                   | Average payoff expectation              | $[-1, 1]$                                    |
| $f_t^V(x), f_t^F(x)$             | Payoff densities (V/F)                  | $\mathbb{R} \rightarrow \mathbb{R}_{\geq 0}$ |
| $\mu_t^V, \mu_t^F$               | Expected payoffs (V/F)                  | $\mathbb{R}$                                 |
| $\sigma_t^{V,2}, \sigma_t^{F,2}$ | Payoff variances (V/F)                  | $\geq 0$                                     |
| $\bar{\mu}_t$                    | Mixed expected payoff                   | $\mathbb{R}$                                 |
| $\gamma_R$                       | Risk-aversion coefficient (CRRA)        | $\geq 0$                                     |
| $U(K)$                           | Utility function                        | $\mathbb{R}_{>0} \rightarrow \mathbb{R}$     |
| <b>Baseline Dynamics</b>         |                                         |                                              |
| $K_t$                            | Running baseline                        | $\mathbb{R}$                                 |
| $K_0$                            | Initial baseline (floor)                | $\mathbb{R}$                                 |



**Table 17 – continued**

| Symbol                                                    | Description                    | Domain / Range                           |
|-----------------------------------------------------------|--------------------------------|------------------------------------------|
| $K^*$                                                     | Target baseline                | $\geq K_0$                               |
| $\theta$                                                  | Required increment to target   | $K^* - K_0 > 0$                          |
| $\tau_\theta$                                             | Hitting time to target         | $\mathbb{Z}_{>0} \cup \{\infty\}$        |
| $N_E(P), N_R(P), \Delta K(P)$                             | Counts, payoff change          | integers                                 |
| $\Lambda_T$                                               | Expected secondaries up to $T$ | $\geq 0$                                 |
| $\lambda_{\text{eff}}$                                    | Effective secondary rate       | $> 0$                                    |
| <b>Optimal Control</b>                                    |                                |                                          |
| $V(t, K)$                                                 | Value function                 | $\mathbb{R}$                             |
| $\tilde{V}(t, K, m)$                                      | Value function with queue      | $\mathbb{R}$                             |
| $\alpha_t^*(K)$                                           | Optimal activation policy      | $\{0, 1\}$                               |
| $\theta_{\text{act}}, \theta_{\text{exit}}$               | Activation / exit thresholds   | $> 0 / \leq 0$                           |
| ExitValue( $K$ ),<br>TargetBonus( $K$ ),<br>TerminalValue | Boundary value functions       | $\mathbb{R} \rightarrow \mathbb{R}$      |
| <b>Statistics</b>                                         |                                |                                          |
| $\mathcal{D}$                                             | Observed data set              | –                                        |
| $\Theta$                                                  | Full parameter vector          | –                                        |
| $\mathcal{L}$                                             | Likelihood function            | $\Theta \rightarrow \mathbb{R}_{\geq 0}$ |
| $m, p_D$                                                  | Parameter counts               | $\mathbb{Z}_{>0}, \geq 0$                |
| AIC, BIC, DIC, LRT                                        | Model selection statistics     | –                                        |
| $\chi^2, D_{\text{KS}}$                                   | GoF statistics                 | $\geq 0, [0, 1]$                         |
| Sharpe, AUC, MSE                                          | Performance metrics            | –                                        |

## B Proof of Dynamic Programming Convergence with State-Dependent Hazards

### B.1 Setup

Define the Bellman operator  $\mathcal{B}$  on  $\mathcal{V} = \{V : \{0, \dots, T\} \times \mathcal{K} \rightarrow \mathbb{R}\}$ :

$$\begin{aligned}
 (\mathcal{B}V)(t, K) = \max_{\alpha \in [0, 1]} \Big\{ & \pi_t(K) [\alpha \mathbb{E}[V(t+1, K + X_t)] + (1 - \alpha) V(t+1, K)] \\
 & + (1 - \pi_t(K)) \mathbb{E}[V(t+1, K + X_t^{\text{forced}})] \Big\}.
 \end{aligned} \tag{67}$$

### B.2 Assumptions

1. **Boundedness:**  $|\text{ExitValue}(K)| \leq M, |\text{TargetBonus}(K)| \leq M, |U(K)| \leq M$ ;
2. **Lipschitz continuity:**  $\pi_t(K)$  is Lipschitz in  $K$  with constant  $L_\pi$ ;

3. **Monotonicity:**  $\eta_t^V, \eta_t^F$  are non-decreasing in  $K$ ;
4. **Subcriticality:**  $\mu < 1$  for branching processes;
5. **Compactness:**  $\mathcal{K}$  is finite with step  $\Delta K > 0$ .

**Lemma B.1** (Monotonicity). *If  $V_1 \leq V_2$  pointwise, then  $\mathcal{B}V_1 \leq \mathcal{B}V_2$ .*

*Proof.* For any  $(t, K)$  and any  $\alpha \in [0, 1]$ :

$$\mathbb{E}[V_1(t+1, K + X_t)] \leq \mathbb{E}[V_2(t+1, K + X_t)], \quad (68)$$

$$V_1(t+1, K) \leq V_2(t+1, K), \quad (69)$$

$$\mathbb{E}[V_1(t+1, K + X_t^{\text{forced}})] \leq \mathbb{E}[V_2(t+1, K + X_t^{\text{forced}})]. \quad (70)$$

Thus every term in the maximization is pointwise smaller, so the maximum is smaller.  $\square$

**Theorem B.2** (Finite-Horizon Convergence). *Value iteration  $V_{n+1} = \mathcal{B}V_n$  initialized at the terminal value converges to the unique fixed point  $V^*$  in  $N = T$  iterations.*

*Proof.* By backward induction: at  $t = T$ ,  $V(T, K) = U(K)$  by definition. Assume  $V(t+1, \cdot)$  is uniquely determined. At time  $t$  the maximum over  $\alpha \in [0, 1]$  exists by compactness of  $[0, 1]$  and continuity of the expectation; the maximizer is unique when the value function is strictly concave in  $K$ , which follows from monotonicity. Hence  $V(t, \cdot)$  is uniquely determined, and induction over the exactly  $T$  backward steps completes the proof.  $\square$

**Theorem B.3** (Contraction). *Under subcriticality ( $\mu < 1$ ),  $\mathcal{B}$  is a contraction in the weighted sup-norm:  $\|\mathcal{B}V_1 - \mathcal{B}V_2\|_\infty \leq \beta \|V_1 - V_2\|_\infty$  with  $\beta = \lambda^{\text{eff}}/(1 + \lambda^{\text{eff}}) < 1$ , where  $\lambda^{\text{eff}} = \lambda/(1 - \mu)$  is the effective intensity.*

*Proof.* For any  $(t, K)$ ,

$$\begin{aligned} |\mathcal{B}V_1(t, K) - \mathcal{B}V_2(t, K)| &\leq \max_{\alpha} \left| \pi_t(K) \alpha \mathbb{E}[V_1 - V_2] + \pi_t(K)(1 - \alpha)(V_1 - V_2) \right. \\ &\quad \left. + (1 - \pi_t(K)) \mathbb{E}[V_1 - V_2] \right| \end{aligned} \quad (71)$$

$$\leq [\pi_t(K) \alpha + \pi_t(K)(1 - \alpha) + (1 - \pi_t(K))] \|V_1 - V_2\|_\infty. \quad (72)$$

The bracket equals 1, giving a non-expansion. Under subcritical branching, the effective intensity  $\lambda^{\text{eff}} = \lambda/(1 - \mu) < \infty$  discounts future contributions, yielding contraction modulus  $\beta = \lambda^{\text{eff}}/(1 + \lambda^{\text{eff}}) < 1$ .  $\square$

**Theorem B.4** (Grid Convergence). *As  $\Delta K \rightarrow 0$  (and  $\Delta m \rightarrow 0$ ), the discrete solution  $V^\Delta$  converges uniformly to the continuous solution  $V$  with rate  $\|V^\Delta - V\|_\infty \leq C \Delta K + D \Delta m$ .*

*Proof.* Let  $\Pi$  be the projection onto the grid. The error satisfies  $\|V^\Delta - V\|_\infty \leq \frac{1}{1-\beta} \|\Pi V - V\|_\infty \leq C \Delta K$ , where  $C$  is the Lipschitz constant of  $V$ ; adding the queue dimension contributes the  $D \Delta m$  term.  $\square$

## C Exact Variance Derivation

### C.1 General Case

For the generalized model,

$$\text{Var}(K_T) = \sum_{s=1}^T \sum_{t=1}^T \text{Cov}(X_s R_s, X_t R_t),$$

where  $\text{Cov}(X_s R_s, X_t R_t) = \mathbb{E}[X_s R_s X_t R_t] - \mathbb{E}[X_s R_s] \mathbb{E}[X_t R_t]$ . For  $s < t$ , conditioning on  $\mathcal{F}_s$ ,

$$\begin{aligned} \mathbb{E}[X_s R_s X_t R_t] &= \mathbb{E}[X_s R_s \cdot r_t(s) \bar{\mu}_t], \\ r_t(s) &:= \mathbb{P}(R_t = 1 \mid E_s = 1), \end{aligned} \tag{73}$$

so that

$$\text{Cov}(X_s R_s, X_t R_t) = \bar{\mu}_s \bar{\mu}_t \text{Cov}(R_s, r_t(s)) + \bar{\mu}_t \text{Cov}(X_s R_s, r_t(s)). \tag{74}$$

### C.2 Core Model (Single-Excitation)

With binary payoffs,

$$\text{Var}(K_T) = \sum_{t \leq T} r_t [1 - \bar{\eta}_t^2 r_t] + 2 \sum_{s < t \leq T} \text{Cov}(X_s R_s, X_t R_t), \tag{75}$$

reducing under independent triggers to  $\text{Var}(K_T) = \sum_{t \leq T} r_t [1 - \bar{\eta}_t^2 r_t]$ .

## D Renewal Theory for Hitting Times

Let  $\tau_\theta$  be the hitting time to threshold  $\theta$ . With jump rate  $\lambda$  and jump law  $F$ ,

$$\mathbb{E}[\tau_\theta] = \frac{1}{\lambda} + \int \mathbb{E}[\tau_{\theta-x}] dF(x), \tag{76}$$

and for binary jumps,

$$\mathbb{E}[\tau_\theta] = \frac{1}{\lambda} + \eta \mathbb{E}[\tau_{\theta-1}] + (1 - \eta) \mathbb{E}[\tau_{\theta+1}], \quad \mathbb{E}[\tau_\theta] = 0 \text{ for } \theta \leq 0. \tag{77}$$

This linear recurrence solved backward yields the closed form  $\mathbb{E}[\tau_\theta] = \frac{1}{\lambda(2\eta-1)} [(\eta/(1-\eta))^\theta - 1]$  for  $\eta > 0.5$ , infinite otherwise.

## E Supplementary Derivations

### E.1 Derivation of Expected Payoff

From  $K_T = K_0 + \sum_{t=1}^T X_t R_t$ , taking expectations and using conditional independence,  $\mathbb{E}[X_t R_t] = \mathbb{E}[R_t] \cdot \mathbb{E}[X_t \mid R_t=1]$ , gives  $\mathbb{E}[K_T] = K_0 + \sum_{t=1}^T r_t \bar{\eta}_t$ .

## E.2 Derivation of the HJB Equation

Taking the continuous-time limit  $\Delta t \rightarrow 0$  of the discrete Bellman equation and expanding  $V(t+1, K) \approx V(t, K) + \partial_t V + \mathbb{E}[X_t] \partial_K V + \frac{1}{2} \text{Var}(X_t) \partial_K^2 V$  (Itô's lemma) yields the HJB equation of Section 9.2.

## E.3 Derivation of the CLT

$M_T = K_T - \mathbb{E}[K_T] = \sum_{t=1}^T (X_t R_t - \mathbb{E}[X_t R_t])$  is a martingale difference sequence; the martingale CLT gives  $M_T / \sqrt{\sum_t \text{Var}(X_t R_t)} \Rightarrow \mathcal{N}(0, 1)$ , and stationarity yields  $\sum_t \text{Var} \sim T\sigma^2$ .

## E.4 Derivation of Hitting Time and Violation Probabilities

The renewal recursion of Appendix D replaces the earlier Poisson approximation  $\mathbb{P}(\tau_{K^*} > T) \approx e^{-\lambda_{\text{eff}} T \bar{\eta}} \sum_{j=0}^{\theta-1} (\lambda_{\text{eff}} T \bar{\eta})^j / j!$  with  $\lambda_{\text{eff}} = \frac{1}{T} \sum_{t \leq T} r_t$ . The violation probability follows from gambler's ruin for random walks with negative drift.

# F Code Implementation

## F.1 Core Model Implementation

```
import numpy as np
from scipy.optimize import minimize
from scipy.stats import poisson, binom, norm

class TriggerPayoffModel:
    """Implementation of the core Trigger-Payoff model with DP."""

    def __init__(self, pi, rho, phi, psi_params, xi_params,
                 eta_V, eta_F, K0, K_star):
        self.pi = np.array(pi)
        self.rho = rho
        self.phi = phi
        self.psi_params = psi_params
        self.xi_params = xi_params
        self.eta_V = np.array(eta_V)
        self.eta_F = np.array(eta_F)
        self.K0 = K0
        self.K_star = K_star
        self.T = len(pi)

    def compute_r(self):
```

```

        """Compute secondary event probabilities."""
        r = np.zeros(self.T)
        for t in range(self.T):
            for s in range(t):
                delay = t - s
                kappa = ((1 - self.phi) * self.alpha_t(s)
                        + self.phi * self.xi(delay))
                r[t] += kappa * self.pi[s]
        return r

def alpha_t(self, t):
    """Voluntary activation policy (optimized via DP)."""
    return 0.5

def psi(self, delta):
    """Voluntary delay distribution (geometric default)."""
    q = self.psi_params.get('q', 0.3)
    return (1 - q) * (q ** (delta - 1))

def xi(self, delta):
    """Forced delay distribution (geometric default)."""
    q = self.xi_params.get('q', 0.5)
    return (1 - q) * (q ** (delta - 1))

def expected_payoff(self, T=None):
    """Compute expected cumulative payoff."""
    if T is None:
        T = self.T
    r = self.compute_r()
    eta_bar = ((1 - self.phi) * self.alpha_t(np.arange(T))
               * (2 * self.eta_V[:T] - 1)
               + self.phi * (2 * self.eta_F[:T] - 1))
    return self.K0 + np.sum(r[:T] * eta_bar)

def prob_target_reached(self, T=None):
    """Compute probability of reaching target."""
    if T is None:
        T = self.T
    r = self.compute_r()
    Lambda = np.sum(r[:T])
    theta = int(self.K_star - self.K0)
    eta_bar = np.mean((1 - self.phi) * (2 * self.eta_V[:T] - 1)
                     + self.phi * (2 * self.eta_F[:T] - 1))
    prob = 0.0
    for n in range(theta, 100):
        p_N = poisson.pmf(n, Lambda)
        for k in range(theta, n + 1):
            if (n + k) % 2 == 0:

```

```

        p_S = binom.pmf((n + k) // 2, n, eta_bar)
        prob += p_N * p_S
    return prob

def prob_violation(self, T=None):
    """Compute probability of baseline violation."""
    if T is None:
        T = self.T
    mu = self.expected_payoff(T) / T - self.K0 / T
    r = self.compute_r()
    eta_bar = ((1 - self.phi) * (2 * self.eta_V[:T] - 1)
               + self.phi * (2 * self.eta_F[:T] - 1))
    sigma2 = np.mean(r[:T] * (1 - eta_bar ** 2 * r[:T]))
    if mu >= 0:
        return 0.0
    return np.exp(-2 * abs(mu) / sigma2)

def solve_dp(self, K_min, K_max, dK):
    """Solve the dynamic programming problem."""
    K_grid = np.arange(K_min, K_max + dK, dK)
    V = np.zeros((self.T + 1, len(K_grid)))
    alpha = np.zeros((self.T, len(K_grid)))
    V[-1, :] = 0

    for t in range(self.T - 1, -1, -1):
        for i, K in enumerate(K_grid):
            i_plus = min(i + 1, len(K_grid) - 1)
            i_minus = max(i - 1, 0)

            V_forced = self.phi * np.sum([
                self.xi(delta) * (
                    self.eta_F[t] * V[t + 1, i_plus]
                    + (1 - self.eta_F[t]) * V[t + 1, i_minus])
                for delta in range(1, min(10, self.T - t))
            ])

            V_act = (self.eta_V[t] * V[t + 1, i_plus]
                    + (1 - self.eta_V[t]) * V[t + 1, i_minus])
            V_wait = V[t + 1, i]

            if V_act > V_wait:
                alpha[t, i] = 1.0
                V_vol = V_act
            else:
                alpha[t, i] = 0.0
                V_vol = V_wait

            V[t, i] = max(

```

```

        0,
        self.pi[t] * V_vol + (1 - self.pi[t]) * V_forced,
        100 if K >= self.K_star else -np.inf
    )
    return V, alpha

if __name__ == "__main__":
    T = 100
    pi = np.full(T, 0.1)
    model = TriggerPayoffModel(
        pi, rho=0.8, phi=0.2,
        psi_params={'q': 0.3}, xi_params={'q': 0.5},
        eta_V=np.full(T, 0.7), eta_F=np.full(T, 0.4),
        K0=100, K_star=110)

    print(f"Expected payoff: {model.expected_payoff():.2f}")
    print(f"P(target): {model.prob_target_reached():.4f}")
    print(f"P(violation): {model.prob_violation():.4f}")

    V, alpha = model.solve_dp(K_min=90, K_max=120, dK=1)
    print(f"Optimal value at start: {V[0, 10]:.2f}")

```

## F.2 Generalized Model Implementation

```

import numpy as np
from typing import Tuple
from scipy.stats import poisson, norm
from scipy.optimize import minimize

class UnifiedTriggerPayoffModel:
    """Unified framework for conditioned event processes with
    branching and continuous payoffs.

    Parameters
    -----
    pi_params : tuple
        (gamma_0, gamma_1, gamma_2, gamma_3) for the state-dependent
        trigger probability.
    rho : float
        Probability that a trigger produces a secondary.
    phi : float
        Fraction of secondaries that are forced.
    offspring_dist : str
        'single' for single-excitation, 'poisson' for branching.
    offspring_mean : float

```

```

    Mean offspring ( $\mu$ ) for branching.
eta_V, eta_F : array-like
    Voluntary / forced payoff positivity probabilities.
payoff_dist : str
    'binary' or 'normal'.
K0, K_star : float
    Initial baseline and target baseline.
gamma_risk : float
    Risk-aversion coefficient (CRRA).
"""

def __init__(self, pi_params, rho, phi, offspring_dist,
              offspring_mean, eta_V, eta_F, payoff_dist,
              K0, K_star, gamma_risk=0):
    self.pi_params = pi_params
    self.rho = rho
    self.phi = phi
    self.offspring_dist = offspring_dist
    self.mu = (offspring_mean if offspring_dist != 'single' else 0)
    self.eta_V = np.array(eta_V)
    self.eta_F = np.array(eta_F)
    self.payoff_dist = payoff_dist
    self.K0 = K0
    self.K_star = K_star
    self.gamma_risk = gamma_risk
    self.T = len(eta_V)

def pi_t(self, K_t: float, t: int) -> float:
    """State-dependent trigger probability."""
    g0, g1, g2, g3 = self.pi_params
    logit = g0 + g1 * K_t + g2 * t + g3 * K_t ** 2
    return 1 / (1 + np.exp(-logit))

def compute_r(self, K_history: np.ndarray) -> np.ndarray:
    """Compute secondary event probabilities."""
    T = len(K_history)
    r = np.zeros(T)
    alpha = np.ones(T) # placeholder for optimal policy
    for t in range(T):
        for s in range(t):
            if self.offspring_dist == 'single':
                kappa = ((1 - self.phi) * alpha[t] * self.rho
                        + self.phi * self.rho)
            else: # branching (mean-field approximation)
                kappa = ((1 - self.phi) * alpha[t] * self.rho
                        + self.phi * self.rho)
            if self.mu < 1:
                kappa *= self.mu / (1 - self.mu)

```



```

        else:
            kappa = np.inf
            r[t] += kappa * self.pi_t(K_history[s], s)
    return r

def expected_payoff(self, K_history: np.ndarray) -> float:
    """Compute expected cumulative payoff."""
    r = self.compute_r(K_history)
    if self.payoff_dist == 'binary':
        mu_V = 2 * self.eta_V - 1
        mu_F = 2 * self.eta_F - 1
    else: # 'normal'
        mu_V = self.eta_V
        mu_F = self.eta_F
    mu_bar = (1 - self.phi) * mu_V + self.phi * mu_F
    return self.K0 + np.sum(r * mu_bar[:self.T])

def solve_dp(self, K_min: float, K_max: float, dK: float,
              m_max: int = 10) -> Tuple[np.ndarray, np.ndarray]:
    """DP with state-dependent hazards and trigger queue."""
    K_grid = np.arange(K_min, K_max + dK, dK)
    m_grid = np.arange(0, m_max + 1)
    V = np.zeros((self.T + 1, len(K_grid), len(m_grid)))
    alpha = np.zeros((self.T, len(K_grid), len(m_grid)))

    for i, K in enumerate(K_grid):
        for m_idx in range(len(m_grid)):
            V[-1, i, m_idx] = self._utility(K)

    for t in range(self.T - 1, -1, -1):
        for i, K in enumerate(K_grid):
            for m_idx, m in enumerate(m_grid):
                i_plus = min(i + 1, len(K_grid) - 1)
                i_minus = max(i - 1, 0)
                m_next = min(m + 1, m_max)
                m_prev = max(m - 1, 0)

                V_act = (self.eta_V[t] * V[t + 1, i_plus, m_prev]
                        + (1 - self.eta_V[t])
                          * V[t + 1, i_minus, m_prev])
                V_wait = V[t + 1, i, m_next]

                if V_act > V_wait:
                    alpha[t, i, m_idx] = 1.0
                    V_vol = V_act
                else:
                    alpha[t, i, m_idx] = 0.0
                    V_vol = V_wait

```

```

        V_forced = self.phi * (
            self.eta_F[t] * V[t + 1, i_plus, m_next]
            + (1 - self.eta_F[t])
            * V[t + 1, i_minus, m_next])

        bonus = (100 * self._utility(K)
                  if K >= self.K_star else -np.inf)
        V[t, i, m_idx] = max(
            0, # exit
            self.pi_t(K, t) * V_vol
            + (1 - self.pi_t(K, t)) * V_forced,
            bonus,
        )
    return V, alpha

def _utility(self, K: float) -> float:
    """CRRA utility."""
    if self.gamma_risk == 0:
        return K
    elif self.gamma_risk == 1:
        return np.log(K)
    return (K ** (1 - self.gamma_risk)) / (1 - self.gamma_risk)

if __name__ == "__main__":
    model = UnifiedTriggerPayoffModel(
        pi_params=(0.1, 0.01, 0.001, 0.0001),
        rho=0.8, phi=0.2,
        offspring_dist='poisson', offspring_mean=0.5,
        eta_V=np.full(100, 0.7), eta_F=np.full(100, 0.4),
        payoff_dist='binary', K0=100, K_star=110,
        gamma_risk=0.5)

    K_history = np.full(100, 100) # placeholder history
    print(f"Expected payoff: "
          f"{model.expected_payoff(K_history):.2f}")

    V, alpha = model.solve_dp(K_min=80, K_max=140, dK=1, m_max=10)
    print(f"Optimal value at start: {V[0, 20, 0]:.2f}")

```

## Data Availability Statement

The datasets designated for real-data validation are publicly available: S&P 500 intraday data (Kaggle, Yahoo Finance), M5 Retail Sales (Kaggle, [m5-forecasting-accuracy](https://www.kaggle.com/competitions/m5-forecasting-accuracy)), and MIMIC-III (PhysioNet, [physionet.org/content/mimiciii/1.4/](https://physionet.org/content/mimiciii/1.4/)). All numerical

results in this paper derive from the synthetic scenarios of §11–§13 and are reproducible from the reference implementation in Appendices F–G. The full codebase, seeds, and per-scenario outputs will be released in a public repository upon publication; no third-party data were used to produce any table or figure in this version.

## Acknowledgments

The author acknowledges the open-source communities that provided the datasets used in this work: Kaggle (S&P 500, M5), PhysioNet (MIMIC-III), and the developers of the Python scientific computing stack.

## References

- [1] Aït-Sahalia, Y., Cacho-Diaz, J., & Laeven, R. J. A. (2015). Modeling financial contagion using mutually exciting jump processes. *Journal of Financial Economics*, 117(3), 585–606.
- [2] Bardi, M., & Capuzzo-Dolcetta, I. (1997). *Optimal Control and Viscosity Solutions of Hamilton–Jacobi–Bellman Equations*. Birkhäuser.
- [3] Bertsekas, D. P. (2015). *Dynamic Programming and Optimal Control*. Athena Scientific.
- [4] Berge, C. (1963). *Topological Spaces*. Oliver & Boyd.
- [5] Billingsley, P. (1999). *Convergence of Probability Measures* (2nd ed.). Wiley.
- [6] Bass, F. M. (1969). A new product growth for model consumer durables. *Management Science*, 15(5), 215–227.
- [7] Cartea, Á., Jaimungal, S., & Penalva, J. (2015). *Algorithmic and High-Frequency Trading*. Cambridge University Press.
- [8] Cox, D. R. (1962). *Renewal Theory*. Methuen.
- [9] Cox, D. R., & Isham, V. (1980). *Point Processes*. Chapman & Hall.
- [10] Crandall, M. G., Ishii, H., & Lions, P.-L. (1992). User’s guide to viscosity solutions of second order partial differential equations. *Bulletin of the American Mathematical Society*, 27(1), 1–67.
- [11] Daley, D. J., & Vere-Jones, D. (2003). *An Introduction to the Theory of Point Processes, Volume I: Elementary Theory and Methods*. Springer.
- [12] Dempster, A. P., Laird, N. M., & Rubin, D. B. (1977). Maximum likelihood from incomplete data via the EM algorithm. *Journal of the Royal Statistical Society: Series B*, 39(1), 1–22.

- [13] Ethier, S. N., & Kurtz, T. G. (1986). *Markov Processes: Characterization and Convergence*. Wiley.
- [14] Fleming, W. H., & Rishel, R. W. (1975). *Deterministic and Stochastic Optimal Control*. Springer.
- [15] Granovetter, M. (1978). Threshold models of collective behavior. *American Journal of Sociology*, 83(6), 1420–1443.
- [16] Harris, T. E. (1963). *The Theory of Branching Processes*. Springer.
- [17] Hawkes, A. G. (1971). Spectra of some self-exciting and mutually exciting point processes. *Biometrika*, 58(1), 83–90.
- [18] Karatzas, I., & Shreve, S. E. (1991). *Brownian Motion and Stochastic Calculus*. Springer.
- [19] Karlin, S., & Taylor, H. M. (1975). *A First Course in Stochastic Processes*. Academic Press.
- [20] Kushner, H. J., & Dupuis, P. (2001). *Numerical Methods for Stochastic Control Problems in Continuous Time*. Springer.
- [21] Øksendal, B. (2003). *Stochastic Differential Equations*. Springer.
- [22] Ross, S. M. (2014). *Introduction to Probability Models*. Academic Press.
- [23] Sutton, R. S., & Barto, A. G. (2018). *Reinforcement Learning: An Introduction* (2nd ed.). MIT Press.
- [24] Van Kampen, N. G. (2007). *Stochastic Processes in Physics and Chemistry*. Elsevier.
- [25] Whittle, P. (1988). Restless bandits: activity allocation in a changing world. *Journal of Applied Probability*, 25A, 287–298.