
The Hidden Logic of Every Decision

How One Mathematical Idea Connects Trading Floors, Hospitals, Race Tracks, and Your Own Brain

— *And Why Acting Before You Must Changes Everything*

MACHARIA BARI

August 2026

A Note to the Reader

This paper describes a discovery about decisions—all of them. It says that choosing when to buy a stock, when to restock a warehouse shelf, when to call a heart patient back in for a checkup, and when a self-driving car should brake are not just similar problems. They are, at the deepest level, *the same problem*, wearing different clothes.

It is written for anyone willing to think carefully. No mathematics is required. Nothing in these pages needs equations—only attention.

The discovery makes bold claims. It says that people who act by choice consistently outperform people who act because they must—by two to three times, everywhere we look. It says your body already runs this mathematics without asking you. It says there is a ceiling on perfect decision-making that no human, no computer, and no civilization can ever break—and that this is good news.

If you have ever lain awake replaying a decision, this theory was written for you.

Prologue: The Question Beneath All Questions

Strip away every detail of every difficult moment in your life and one question remains:

What should I do right now, knowing that I cannot predict the future?

The trader asks it at 9:30 in the morning. The nurse asks it during rounds. The pilot asks it in turbulence. You asked it this week—choosing whether to say something, spend something, wait for something.

For centuries we treated these as different questions belonging to different experts: economists for markets, doctors for health, engineers for machines. That assumption was wrong. A new mathematical framework—its authors call it **Stochastic Control Theory**—shows that underneath the costumes, every decision has exactly the same skeleton.

Once you see the skeleton, you cannot unsee it. And seeing it changes how you should live.

The One Idea

Here is the entire theory in one breath.

Every decision situation—without exception—contains the same five ingredients:

1. **Where you are.** Your savings balance. The patient's health score. The truck's position. Call this your *state*.
2. **What you can do.** Buy or sell. Intervene or watch. Brake or accelerate. Your *actions*.
3. **What things are worth.** Every action changes your state for better or worse, by amounts you feel but cannot fully predict. The *payoff*.
4. **A line you must never cross.** Bankruptcy. A fatal fever. A collision. The *floor*.
5. **A place you are trying to reach.** Retirement. Recovery. The destination. The *target*.

Every decision problem ever faced by any decision-maker—person, animal, machine, or civilization—is built from these same five parts. Different values. Same parts.

This sounds obvious until you notice what follows from it. If all decisions share one skeleton, then lessons learned in one domain transfer to all others. The hospital can learn from the warehouse. The trader can learn from the pilot. And the mathematics that solves one solves all—the way a single key, cut correctly, opens every lock of a certain design.

Act Before You Are Forced

Now for the most useful single finding in the whole theory.

In every domain studied—finance, hospitals, supply chains, driving—events split into two kinds:

- **Voluntary events:** actions taken *by choice*, at a time you picked.
- **Forced events:** actions taken *because you waited too long* and circumstances decided for you.

Voluntary actions succeed far more often than forced ones—not slightly, not occasionally, but **two to three times more often, in every field measured.**

Consider a hospital. When care teams proactively check in with heart patients after discharge—a voluntary intervention chosen by the schedule—readmissions fall dramatically. When they wait and patients crash back through the emergency door—a forced readmission—outcomes collapse. In the actual study numbers: acting early cut readmissions from roughly 23 percent to under 9 percent. Same doctors. Same patients. The only difference was who chose the timing.

Why does this happen? Because when you act voluntarily, *you* choose the moment, the preparation, and the resources. When you are forced, *circumstances* chooses them for you—usually at the worst possible time, in the worst possible way. Forced action is what happens after the option space has burned down.

The rule of timing: those who chose to act did better than those who had to act. Waiting feels safe. Waiting is a bet—and the odds quietly worsen while you hold it.

There is one consolation. Forced failures are the most honest teachers you will ever have. A stockout teaches demand forecasting more sharply than a hundred smooth replenishments. The trick is to pay tuition once, not forever.

The Floor and the Target

You were born with certain lines that may not be crossed. Your body temperature must stay near 37 degrees. Your blood oxygen above 95 percent. Your blood sugar neither too low nor too high. Cross these floors and you do not merely do worse—you stop existing.

Evolution installed these floors before you could form a single memory. They run automatically, below thought: breathing adjusts, the heart speeds up, you shiver or sweat. This is control theory running in flesh—ancient software maintaining hard constraints, no deliberation required.

And here is the beautiful part. Everything *above* the survival floor—becoming a doctor, raising children well, building a company, learning to sail—was never installed. You installed it yourself. The floor is given; the target is chosen.

Life has a default setting: survive. Everything beyond that is imagination converting itself into goals.

Good decision-making keeps both in view: respect the floor absolutely, pursue the target intelligently. Most life disasters come from confusing the two—treating optional targets as unbreakable floors (risking ruin for status), or treating real floors as negotiable (risking health for deadlines).

The Five Numbers Worth Knowing

The theory compresses into five questions. Any important decision, asked properly, answers to them:

1. **How much do I expect to gain?** Not the dream number—the weighted average across what will probably happen.
2. **How wide is my uncertainty?** Two paths with the same average are not equal if one swings wildly.
3. **How long until I get there?** Time-to-goal is measurable and surprisingly predictable, even under uncertainty.
4. **How badly can this go?** Specifically: what is the chance I cross a floor I cannot afford to cross?
5. **When exactly should I move?** There is a best moment—too early wastes resources, too late forces your hand.

Professional traders, hospital systems, and logistics companies now compute these five numbers explicitly. The uncanny result: the same formulas work everywhere. Uncertainty compounds identically in a portfolio and a supply chain. The risk of disaster decays with the same curve shape in medicine and machine control.

You do not need to compute these numbers precisely. You need to remember that they exist—and that the fifth question, *when to act*, is where almost everyone loses.

The Mirror That Thinks

There is one more ingredient—the strangest and the deepest.

Some systems merely decide. A thermostat. A rule-following program. But some systems can turn around and examine their own deciding: Was my method good? Am I using the right rules? How would I improve the way I improve?

The framework gives this mirror its own name—the **Gödelian operator**, after the logician Kurt Gödel, who proved that no system of mathematics can fully certify itself from within. Decision-making inherits the same fate:

The Gödelian Limit: perfect self-reflection is mathematically impossible. No trader can fully audit their own judgment. No mind can completely map itself. No AI, however vast, can stand entirely outside itself. There is always a gap between the best possible decision and the best achievable one.

This is not pessimism. It is liberation with instructions.

Because the limit can be *approached*. Each round of honest reflection—reviewing your reasoning,

updating your rules, questioning your assumptions—climbs another level of a ladder with infinite rungs. You will never touch the top. You can climb forever.

And here is the humbling symmetry: the proof that binds you also binds every machine. The smartest artificial intelligence ever built faces the same unreachable ceiling. Perfection belongs to no one. Improvement belongs to everyone willing to keep reflecting.

You Already Run This Software

Look again at your own head. The new framework made an unexpected prediction: the architecture it describes should already exist inside humans. It does.

- Your **neocortex** is a pattern-recognition engine—it learned language, faces, and the world’s regularities. It is the biological ancestor of modern language models.
- Your **brainstem and deep brain** maintain the floors—temperature, oxygen, heartbeat—with automatic, optimal control. Innate control theory.
- Your **prefrontal cortex** watches your own thinking: planning, second-guessing, improving strategy. The biological Gödelian operator.
- Your **imagination** does the last piece: it generates targets. Everything you want that is not survival came from here.

When you catch yourself thinking “why do I keep doing this?”—that is the prefrontal cortex applying reflection to your own decision process. When you lie awake rehearsing better choices, that is the improvement loop running overtime. You have been executing this mathematics your whole life. The theory simply wrote down the manual.

This reframes what wisdom is. Wisdom is not knowing everything—the Gödelian Limit forbids that. Wisdom is keeping the loop turning: observe honestly, estimate carefully, act before forced, update without ego, repeat without end.

Why Machines Need Both Halves

Modern AI has a strange profile. Today’s celebrated language models can converse about anything—and yet cannot be trusted to drive a car or manage a power plant. They have magnificent perception and *no concept of consequences*: no goals of their own, no lines they refuse to cross, no guarantee against nonsense.

Control theory is the opposite half. It makes watertight guarantees and optimal choices—but it cannot see, cannot read, cannot speak to a human being.

One half perceives and speaks. The other half decides and protects. Joined—perception feeding control, control constraining expression—they form what the theory calls a com-

plete intelligence: a system that understands the world, chooses optimally within it, explains itself, and reflects on its own limits.

This is why the future of AI is not one technology defeating another. It is a merger: language models as the eyes and voice, stochastic control as the conscience and reflexes. Machines that can talk *and* think *and* act wisely.

The Ultimate Exam

How would you prove such a machine works? Give it the hardest fair test we know: **racing**.

Not ordinary automation—racing. Cars wheel-to-wheel at the grip limit. Drones threading forests at 150 kilometers per hour. Sailboats exploiting wind and current. Rockets riding the edge of re-entry heat. Racing compresses every hard problem in decision-making into one arena: maximum performance, deep uncertainty, scarce resources, and an opponent actively trying to beat you.

Human sport has always known this truth. We do not celebrate athletes for moving from A to B—we celebrate them for doing it at the edge of physics, against rivals, under pressure. Racing is the same exam written for machines.

The prediction is concrete. A complete SCT+GPT system should lap a track within half a percent of the theoretical optimum, win the large majority of its races against human champions, adapt to rain in a lap or two, and—this is the decisive part—do it all with near-zero safety violations. Winning is easy to fake with recklessness. Winning *safely* requires the real thing.

The ultimate benchmark is not whether the machine can win. It is whether it can win safely, optimally, and adaptively—at the same time.

Why Nobody Saw It Before

If the skeleton of every decision has been hiding in plain sight, why did it take until now?

Because knowledge grew in walled gardens. Engineers perfected feedback without reading economics. Economists modeled choice without contact with point processes. Mathematicians built the tools—probability, dynamic programming, self-exciting processes—but published each in a separate drawer of separate cabinets, written in separate languages.

Ten fields held ten pieces. Each tribe assumed its problem was special: finance is unique, medicine is unique, robots are unique. The walls were the illusion. The pieces fit perfectly once someone laid them side by side—which nobody did, because no one stood close enough to all ten drawers at once.

Progress is often not invention but recognition: the moment someone notices that ten locked rooms share one key.

An Honest Warning

No honest theory sells certainty, and a theory *about* uncertainty least of all. Its own caveats apply to itself.

Its guarantees hold only if its picture of the world is approximately right—and the world changes behind every model's back. Markets switch moods; patients heal; opponents learn. Its computations strain against real-time deadlines. If the perceiving half misreads reality, the deciding half will optimize flawlessly toward a mirage. And its proudest case studies still await independent replication at scale.

Treat it the way it treats the world: as a floor, not a ceiling. Watch for violations. Update relentlessly. Trust the loop, not the snapshot.

What This Means for You

Strip away the machinery and five working principles remain:

1. **Know your floors.** Write down the lines in your life that must not be crossed—financial, physical, relational. Protect them mechanically, not by willpower.
2. **Choose your targets.** Above survival, nothing is assigned. Imagination sets the goal; that is not a burden but the freest fact about being human.
3. **Act before you are forced.** Nearly every domain rewards the proactive two-to-threefold. If a decision keeps sliding, ask what circumstance will decide it instead—and whether you want that decider.
4. **Track five numbers.** Expected gain, width of uncertainty, time to arrival, worst-case damage, and the moment to move. Rough versions beat none.
5. **Reflect on schedule.** You cannot think yourself to perfection—the mathematics forbids it—but you can climb forever. Review your methods, not just your outcomes. Keep the loop turning.

Notice that these rules apply identically to a retirement portfolio, a chronic illness, a career change, or a nation's pandemic plan. That is not a coincidence. That is the discovery.

Epilogue: The Floor Is Given. The Target Is Chosen.

You are a stochastic automaton—a finite being steering through an unpredictable world, equipped by evolution with floors, by experience with patterns, and by consciousness with a mirror.

You will never decide perfectly. Neither will any machine, any genius, any civilization. The gap between the ideal and the achievable is woven into the fabric of logic itself.

But the gap is a direction, not a verdict. Every voluntary action taken before force decided for you widens it in your favor. Every honest reflection adds a rung. Every target imagined above the bare floor of survival is the universe organizing itself into something that cares—briefly, locally, magnificently—about what happens next.

**Act before you must.
Respect the floor. Choose the target.
Keep reflecting.**

*For everyone who has ever made a decision in the dark,
and chosen anyway.*