

Grounded Agents: A Founder’s Compass

Application Paper IV of the *Grounded Agents* series
The Advice Essay, Formalized as a Companion to the Framework

Macharia Bari

Abstract

The advice essay *Advice to a Founder Entering AI* was written to be acted on, not proved. This paper is its formal companion: it restates each counsel of the essay as a claim with a stated antecedent, and ties each claim to the results of Paper 0, *The Inference-First Picture of Mind, Logic, and Mathematics* and of the preceding application papers. The central thread is the distinction between a map and a territory: the framework predicts where value will accrue; a business is a specific customer, a specific pain, and a specific price. We prove that intelligence is a commodity and grounded experience is not; that the moat is the logged experience of a closed loop, not the model; that the cheapest on-ramp is a literal instruction rather than a metaphor; that a founder who logs everything and understands the logs outcompetes a founder with a better model and no data; that the platform position is earned, not declared; that staying quiet is an information asymmetry play; that the inversion will be visible only retrospectively; and that a loop-directed bet is the structurally aligned bet. We close with the Compass—the essay’s summary rendered as a single conditional. All results are conditional on the axioms of Paper 0, *The Inference-First Picture of Mind, Logic, and Mathematics*.

Remark 0.1 (Reading). Theorem numbers in this paper are local. The framework, definitions, and proofs of the results recalled in Section 2 live in Paper 0, *The Inference-First Picture of Mind, Logic, and Mathematics*; the industrial, entry, transfer, and stress-test results invoked below are from Application Paper I, Application Paper II, and Application Paper III. The counsels formalized here originate in the advice essay; this paper gives each a stated antecedent.

Contents

1	Introduction	2
2	Recalled Framework	2
3	The Map and the Territory	3
4	The Moat	4
4.1	What a Founder Is Actually Selling	4
4.2	Why the Moat Is the Compass	4
5	Domain Selection	4
6	The Cheapest Loop	5
7	Logging	6

8	Resisting the Sale of the Loop	6
9	Staying Quiet	7
10	The Shape of the Transition	7
11	The Asymmetry	7
12	The Unglamorous Truth	8
13	The Compass	8
14	Open Problems	9
15	Conclusion	9
A	Summary of Definitions	10
B	Summary of Theorems	10
C	A Note on Method	10

1 Introduction

The framework of Paper 0, *The Inference-First Picture of Mind, Logic, and Mathematics* explains where value accrues in the grounded-agent economy. The application papers translate that explanation into industrial structure (Application Paper I), entry points and transfer (Application Paper II), and a stress test (Application Paper III). This paper is addressed, explicitly and exclusively, to the person who read all of that and asked: *what do I do on Monday?*

The answer is a set of cautions and instructions collected in the advice essay. This paper performs one task on that essay: it says what each counsel would have to be true for it to hold, and which framework result it rests on. The result is a set of conditionals. Some are theorems—their antecedents are framework axioms. Others are strategic heuristics—their antecedents are the empirical shape of the market, and they hold conditional on that shape.

The structure mirrors the essay. Section 2 recalls the framework. Section 3 formalizes the map–territory distinction. Section 4 formalizes the moat. Section 5 isolates the domain-selection criteria. Section 6 is the cheapest loop. Section 7 is the logging discipline. Section 8 is the refusal to sell the loop. Section 9 is the value of silence. Section 10 is the shape of the transition. Section 11 is the structural alignment of bets. Section 12 is the unglamorous truth. Section 13 collects the section-level counsels as a single compass. We close with open problems.

One methodological note, stated once. A framework that tells a founder nothing actionable is not a theory; it is a description. A founder who acts on a framework without knowing which counsels are theorems and which are heuristics is building on unknown ground. This paper exists to separate the two.

2 Recalled Framework

We recall the framework results used below, without proof. They are proved in Paper 0, *The Inference-First Picture of Mind, Logic, and Mathematics* and in the application papers.

Definition 2.1 (Agent and Grounding). An *agent* is a tuple $\mathcal{A} = (\mathcal{I}, \mathcal{S}, \mu)$ with inference system $\mathcal{I}$, sensory states $\mathcal{S}$, and grounding map μ . The *grounding coefficient* $G(\mathcal{A})$ measures the fraction of belief contents grounded in $\mathcal{A}$'s own sensory history.

Definition 2.2 (Loop Closure and Grounded Data). Agent $\mathcal{A}$ *closes the loop* in environment $\mathcal{E}$ if it has a self-sourced observation operator with respect to $\mathcal{E}$ and $G_{\mathcal{E}}(\mathcal{A}) > 0$. Its *grounded data* is the set of belief contents in $\mathcal{E}$ not derivable from public corpora.

Theorem 2.3 (The Value Layer Shifts). *The value layer of the AI industry shifts structurally from the model layer to the loop layer. (Application Paper I, Application Paper II.)*

Theorem 2.4 (Grounding Cannot Be Purchased). *No purchase of components yields grounding without the agent's own interaction with the environment. (Paper 0, The Inference-First Picture of Mind, Logic, and Mathematics, Application Paper I.)*

Theorem 2.5 (The Integrator Moat). *An integrator who closes the loop in $\mathcal{E}$ acquires grounded data not obtainable by component purchase. (Application Paper I.)*

Theorem 2.6 (Cheapest Entry Point). *Among general entry points, the minimum loop cost is achieved by a tabletop manipulation loop with a low-cost arm. (Application Paper II.)*

Theorem 2.7 (Domain Value Ranking). *Among equal-cost entry points, domain value $DV = V \cdot P \cdot T / c_{\text{loop}}$ ranks pharma compounding and lab automation first, followed by electronics, food, micro-fulfillment, and medical laboratories. (Application Paper II.)*

Theorem 2.8 (Transfer Determines Industrial Shape). *The industrial landscape is a function of the transfer coefficient T : high T everywhere gives one or a few general winners; $T = 1$ gives a federation of domain integrators. (Application Paper II.)*

Remark 2.9 (This Paper's Place in the Series). The preceding papers answer *what* will happen. This paper answers what a founder should *do*. The other papers' results are imported here as constraints and justifications; the contribution of this paper is the export—a set of actionable conditionals with their antecedents stated.

3 The Map and the Territory

Definition 3.1 (Map and Territory). The *map* is the framework's account of where value will accrue. The *territory* is the specific customer with a specific pain who pays a specific price for a specific solution.

Proposition 3.2 (Map and Territory Are Distinct Questions). *Where value accrues (the map) and what to build (the territory) are different questions. Conflating them is the most common failure mode for founders who have read the theory.*

Proof. The map is a claim about the value layer and its occupants' characteristics (Theorems 2.3, 2.5). It answers: which positions exist. It does not answer: which customer, which pain, which price. A map with no territory is a philosophy; a territory with no map is a guess. Neither condition alone constitutes a business. $\square$ $\square$

Prediction 3.3 (The Map Says Close the Loop; the Territory Says for Whom). The map tells the founder to close a loop in a high-value domain. The territory tells the founder which specific loop, for which specific customer, at what price. Both are required. Neither follows from the other.

Corollary 3.4 (Founders Need Both). *A founder who has read the theory but skipped the territory has no business. A founder who knows a customer but not the structure is building on unknown ground. The framework’s answer to the founder is not either/or but both/and.*

4 The Moat

4.1 What a Founder Is Actually Selling

Theorem 4.1 (Intelligence Is a Commodity; Grounded Experience Is Not). *Under the framework, intelligence—the generative component of the agent—is a commodity, getting cheaper over time. Grounded experience, the data generated by closing a loop in a specific environment, is not.*

Proof. By the model-layer-commoditization result of Paper 0, *The Inference-First Picture of Mind, Logic, and Mathematics* (restated in Application Paper I), substituting one generative component for another leaves G invariant. Any model available today will be beaten by a cheaper one within the relevant horizon. By Theorem 2.4, grounded experience cannot be purchased; it can only be lived. Hence the two assets have opposite supply curves: the model’s price falls, the grounded data’s availability does not. $\square$

Corollary 4.2 (The Moat Is Not the Model). *A founder’s moat is not the model. It is the logged experience of the closed loop—grounded data, trial logs, falsification history. This data is not on the internet, cannot be scraped, cannot be bought, and can only be lived.*

Proof. By Theorem 4.1 and Theorem 2.5. $\square$

Remark 4.3 (The Single Most Important Counsel). Corollary 4.2 is the essay’s load-bearing counsel. If a founder internalizes nothing else, this is the claim to internalize: your moat is not your model; it is your logged experience.

4.2 Why the Moat Is the Compass

Proposition 4.4 (The Moat Is Both the Defense and the Direction). *The grounded-data moat simultaneously defends the founder’s position and points at the next position.*

Proof. As a defense: by Theorem 2.5, the grounded data is not obtainable by component purchase, so it resists competitive copying. As a direction: by the transfer result of Application Paper II, the grounded data determines the transfer coefficient T to adjacent environments, and the direction of expansion is the direction of maximum T . The single asset thus answers both strategic questions: where the position is defensible, and where to expand. $\square$

Corollary 4.5 (The Framework Gives the Founder One Thing). *Everything the framework gives a founder reduces to one instruction: close a loop, log it, and own what the loop produces. The moat is the compass; the compass is the moat.*

5 Domain Selection

The framework ranks domains by value; it does not assign the founder a domain. The essay’s counsel is to pick a domain the founder can win. We formalize the test.

Definition 5.1 (Domain Selection Test). A domain $\mathcal{E}$ is winnable for founder Φ if all four of the following hold:

- (D1) **Cheap loop.** The loop is cheap to close: bounded environment, resettable, low safety risk.
- (D2) **Money and pain.** The customer has money and pain: high labor cost replaced, high error cost, or both.
- (D3) **Proprietary data by construction.** The grounded data is proprietary by construction: no one else holds the trial logs.
- (D4) **Transfer.** The capability transfers: what is learned here applies next door.

Proposition 5.2 (The Domain Test Implements the Domain Value). *The domain selection test of Definition 5.1 instantiates the domain value $DV = V \cdot P \cdot T / c_{loop}$ of Application Paper II: each criterion corresponds to a factor of DV.*

Proof. (1) corresponds to c_{loop} (loop cost); (2) corresponds to V (commercial value); (3) corresponds to P (data propriety); (4) corresponds to T (transfer coefficient). A domain passes the test exactly when its DV factors are favorable. $\square$

Corollary 5.3 (Do Not Start Broad). *Do not start with humanoids, general-purpose robots, or a platform. Start with one task, one customer, one loop.*

Proof. By Theorem 2.6, the cheapest loop is a narrow one. By Theorem 2.8, a general-purpose starting point requires $T > 1$ across all environments, which by Application Paper II is the strong regime rather than the default one. A broad start therefore violates (D1) of Definition 5.1 and buys the costliest antecedent. $\square$

Prediction 5.4 (Narrow Outcompetes Broad at the Margin). The founders who attempt to be general from day one burn capital on breadth and end with nothing; the founders who win start narrow and expand along the transfer gradient.

Justification. By Application Paper II’s prediction on optimal entry strategy, given the transfer regime. $\square$

6 The Cheapest Loop

Proposition 6.1 (The Entry Point Theorem Is a Literal Instruction). *The claim that a tabletop manipulation loop with a sub-\$2,000 arm is the cheapest on-ramp (Theorem 2.6) is not a metaphor. It is a literal instruction about hardware and environment choice.*

Proof. Theorem 2.6 minimizes $c_{hw} + c_{env} + c_{int}$ over entry points; the optimum is a tabletop arm in a bounded environment. The numbers are quantities, not parables. $\square$

Remark 6.2 (The Logic of the Cheap Prototype). The early prototype should cost less than a laptop. Not because the founder is cheap, but because the speed of iteration matters more than the quality of the hardware: thousands of trials, most of them failures, require the cost of failure to be near zero. The founders who spend \$100,000 on a robot before they have a working loop run out of money before they have a product. The reason is structural, not aesthetic: the grounded-data flywheel of Application Paper I is built by accumulating trials, and trial frequency is inversely proportional to trial cost.

Prediction 6.3 (Iteration Rests on a Function of Trial Cost). The rate at which a founder accumulates grounded data is decreasing in the cost of a failed trial. Cheap hardware is therefore an epistemic instrument, not a cost-optimization habit.

7 Logging

Definition 7.1 (The Logging Discipline). The *logging discipline* consists of four operations on the grounded data $\mathcal{D}_A(\mathcal{E})$:

- (L1) **Instrument.** Record every sensor reading, model output, and comparison result. If it is not logged, it did not happen.
- (L2) **Structure.** Store raw data with metadata—time, position, object, outcome, error type—so that it is queryable.
- (L3) **Protect.** Encrypt the logs, do not share them with vendors, do not let them leak through a demo. The logs are the moat.
- (L4) **Mine.** The logs are not only for training; they are for understanding. What fails, why, and what patterns emerge determine what to build next.

Theorem 7.2 (The Logging Theorem). *A founder who logs everything and understands the logs outcompetes a founder with a better model and no data.*

Proof. Let A be a founder who logs everything, and B a founder with a better model and no data. By Theorem 4.1, B 's asset—the better model—is a commodity: its value decays as cheaper models appear. By Theorem 2.5, A 's asset—the grounded data—is not purchasable and not derivable from public corpora. By operations (L1)–(L4), the grounded data is captured completely, made queryable, protected, and converted into decisions. B 's model confers no grounded data (Theorem 2.4); hence B has no asset that A cannot eventually match (by purchase), while A has an asset B cannot acquire at any price. The asymmetry is structural. $\square$

Corollary 7.3 (The Logs Are the Moat). *Corollary 4.2 operationalized: instrument everything, structure the logs, protect the logs, mine the logs. The founder who treats the logs as an asset compiles the moat; the founder who treats the model as the asset compiles a commodity.*

8 Resisting the Sale of the Loop

Definition 8.1 (The Platform Temptation). The *platform temptation* is the temptation to sell the loop itself—*loop closure as a service*—before owning a domain.

Proposition 8.2 (The Platform Position Is Earned, Not Declared). *The Microsoft-type position of the loop layer—the standard interface everyone writes to—is earned by owning a domain first, and declared at one's peril before that.*

Proof. By Application Paper III (the cast of the standardized era), the platform gate is defined relative to a standard, and standardization requires a market. A loop sold with no reference implementation is a framework with no demonstration. By the moat logic of Theorem 2.5, the only way to prove the loop is to own a domain in which it closes. Hence: prove the loop in one domain first; then, if the transfer coefficients hold, consider expanding. $\square$

Prediction 8.3 (Selling Too Early Is the Standard Failure). A founder who sells loop closure before owning a domain will be treated as selling vaporware; a founder who owns a domain and then sells the loop will be treated as a platform. The order is the content of the counsel.

Remark 8.4 (The Genesis of the Position). The Microsoft position of the loop layer will be occupied by whoever closes the loop profitably first, in a domain, and then generalizes the interface. This is a theorem about structure (Proposition 8.2) and a prediction about sequencing. The declared platform is a hypothesis; the proved loop is a position.

9 Staying Quiet

Proposition 9.1 (The Quiet Advantage Is Information Asymmetry). *The current AI industry is loud: model labs announce capability, robotics companies announce funding, all claim proximity to general-purpose. A founder closing a narrow loop who stays quiet runs the asymmetry that the incumbents are distracted by layer-directed strategy and will not notice the loop until it is too late.*

Proof. By Application Paper II, layer-directed strategy—spending on the model layer to make it permanent—does not move the value layer, and the inversion is not visible from the present: observers attend to the currently dominant layer. By (Theorem 2.3 in Application Paper I), the loop layer’s occupants are positioned where the incumbents are not looking. Announcement would supply the incumbents with the one fact they lack—where the value layer is actually being built. Silence withholds the map. $\square$

Prediction 9.2 (Staying Quiet Is Playing the Aligned Bet). Stay quiet. Build the loop. Log the data. Own the domain. The incumbents’ attention is a scarce resource that the founder should not voluntarily consume.

10 The Shape of the Transition

Proposition 10.1 (The Inversion Is Not Sudden to Participants). *The structural inversion of Theorem 2.3 will not happen suddenly. It will happen slowly, then all at once, and then everyone will claim they saw it coming.*

Proof. By Application Paper I and Application Paper II, the inversion is a structural consequence of grounding, which accumulates incrementally inside closed loops: each trial is an increment. The accumulation is continuous; the visibility is not. By the invisibility result (Application Paper I), the shift becomes apparent only retrospectively, at which point participants reconstruct it as foreseeable. $\square$

Corollary 10.2 (Position Before Visibility). *Build for the loop layer now, even when the market does not understand it, even when investors push back, even when the press ignores it. By the time the inversion is visible, the positions will be taken.*

Proof. By Proposition 10.1, visible arrival of the inversion is retrospective. Occupying a position before its visibility is the only way to hold a value-layer position when the inversion completes. $\square$

11 The Asymmetry

Definition 11.1 (Alignment). A bet is *structurally aligned* if its success requires only that the framework’s structural predictions hold. A bet is *structurally misaligned* if its success requires that the framework’s structural predictions fail.

Proposition 11.2 (Alignment of Loop-Directed and Layer-Directed Bets). *A loop-directed bet is aligned; a layer-directed bet is not.*

Proof. By Theorem 2.3 and the strategy results of Application Paper II, the value layer shifts to the loop layer, and no layer-directed strategy moves it. Building a closed loop in a high-value domain succeeds if the loop layer becomes the value layer. Building a better model succeeds only if the model layer remains the value layer—i.e., only if the framework’s structural prediction fails. $\square$ $\square$

Prediction 11.3 (The Asymmetry Changes Odds, Not Certainty). Venture is asymmetric: large upside, low probability. The framework does not change this distribution. It changes which bets are aligned. The founder who builds a loop in a high-value domain, logs the data, and expands along maximum transfer holds a structurally defensible position; the founder who builds a better model does not. Both may succeed; only one is playing the game the framework says is winnable.

Corollary 11.4 (Accept the Asymmetry; Play the Aligned Bet). *Accept the asymmetry. Play the aligned bet. Increase the odds. Do not expect certainty.*

12 The Unglamorous Truth

Proposition 12.1 (The Winners Are Not on Stage). *The founders who will win the grounded-agent era will not look like the founders of the current AI moment. They will not be on stage, will not be quoted in the press, will not have huge followings.*

Proof. By Application Paper I and Application Paper II, the value layer accrues to loop-closers, and the loop closes in a lab, a workshop, or a warehouse. The instruction set is not visibility; it is trial accumulation, failure logging, and loop closure. Visibility is a model-layer currency, and the model layer is not the value layer (Theorem 2.3). $\square$ $\square$

Prediction 12.2 (The Reward Is Structural, Not Glamorous). The work is unglamorous; the reward is structural. The founders who understand this outperform the founders who do not, other things equal.

Justification. By Theorem 7.2 and the alignment result (Proposition 11.2). $\square$ $\square$

13 The Compass

The essay’s closing counsel, rendered as a single conditional:

Theorem 13.1 (The Compass). *If a founder closes the loop in a high-value domain, logs the data, owns the domain, expands along maximum transfer, stays quiet, and plays the aligned bet, then the founder’s position is structurally defensible when the inversion arrives.*

Proof. Closure of a loop yields the moat (Theorem 2.5); logging the data compounds it (Theorem 7.2); owning the domain secures the value-layer position (Theorem 2.3); expansion along T extends it (Theorem 2.8); silence preserves the asymmetry (Proposition 9.1); the aligned bet is aligned (Proposition 11.2). Hence the conjunction is sufficient for a structurally defensible position at inversion time. $\square$ $\square$

Corollary 13.2 (The Map Is the Compass; the Territory Is the Business). *The map tells the founder where value will accrue; the compass tells the founder which direction to go. Both are the same object read twice. The founder has the map. The task is to find the territory.*

14 Open Problems

Open Problem 14.1 (Actionable Signal for the Transfer Coefficient). The direction of expansion is the direction of maximum T , but T is measured only by closing the next loop (Application Paper II). What negligible-cost internal signal approximates T well enough to steer expansion?

Open Problem 14.2 (Valuing a Grounded-Data Asset). If grounded data is the moat, how should it be valued—as an asset, an option on adjacent domains, or as insurance against model commoditization? The framework says it is the moat (Theorem 2.5); it does not say how to price it, in a funding context or an exit context.

Open Problem 14.3 (The Founders-in-Canvas Problem). The framework names positions, not occupants (Application Paper I, Application Paper II). For a founder this is the central epistemic gap: the position is visible, the occupant is not. Can the axioms be extended to place constraints on *conditions* under which the occupant is likely to appear (for example, the tinkering threshold, the trial rate, or the availability of cheap loops)?

Remark 14.4 (The Status of the Heuristics). Several counsels of this paper are heuristics conditional on the empirical shape of the market: the \$2,000 arm, the laptop-priced prototype, the sequence narrow-then-expand. They are consequences of the framework’s structural claims given present cost curves, and they should be re-derived as those curves change. The theorems they rest on do not change.

15 Conclusion

This paper has formalized the advice essay as a set of conditionals, each attached to a framework result. The line of argument is now complete:

1. The map (where value accrues) and the territory (which customer, which pain, which price) are distinct questions; a founder needs both (Proposition 3.2).
2. Intelligence is a commodity; grounded experience is not (Theorem 4.1). The moat is the logged experience of the closed loop, not the model (Corollary 4.2).
3. The moat is also the compass: it defends the position and points to the next one (Proposition 4.4).
4. Domain selection is the domain value test read as a research program (Proposition 5.2); the founder starts narrow (Corollary 5.3).
5. The cheapest loop is a literal instruction, not a metaphor (Proposition 6.1); cheap prototypes are an epistemic instrument (Remark 6.2).
6. Log everything, and mean everything: instrument, structure, protect, mine (Definition 7.1); a founder who logs outcompetes a founder who models (Theorem 7.2).
7. The platform position is earned, not declared (Proposition 8.2).
8. Quiet is an information asymmetry play, aligned with the inversion (Proposition 9.1).
9. The inversion is slow-then-sudden-then-claimed; the founder must be in position before visibility (Proposition 10.1, Corollary 10.2).

10. Bets are either aligned or misaligned with the structure; play the aligned bet (Proposition 11.2).
11. The winners are not on stage; the reward is structural (Proposition 12.1).
12. The Compass: close the loop, log the data, own the domain, expand along maximum transfer, stay quiet, play the aligned bet (Theorem 13.1).

The advice essay ends with a line that bears restating here, because this paper’s whole purpose is to make it precise: *the map is the compass. The territory is the business.*

The structure is visible to anyone who reasons from the axioms. The occupants are not. The founder can be one of them, if willing to do the work.

A Summary of Definitions

Symbol	Name	Reference
Map	Where value will accrue	Definition 3.1
Territory	Customer, pain, price	Definition 3.1
(D1)–(D4)	Domain selection test	Definition 5.1
(L1)–(L4)	Logging discipline	Definition 7.1
Platform temptation	Selling the loop before owning a domain	Definition 8.1
Alignment	Success requires the structure to hold	Definition 11.1

B Summary of Theorems

Theorem	Statement
4.1	Intelligence is a commodity; grounded experience is not
4.2	The moat is the logged experience, not the model
4.4	The moat is both defense and direction
5.2	The domain test implements DV
5.3	Do not start broad
6.1	The entry point theorem is literal
7.2	Logging outcompetes modeling
8.2	The platform position is earned, not declared
9.1	Quiet is information asymmetry
10.1	Slow, then sudden, then claimed
10.2	Position before visibility
11.2	Loop-directed is aligned; layer-directed is not
12.1	The winners are not on stage
13.1	The Compass
13.2	Map is compass; territory is business

C A Note on Method

This paper formalizes normative counsel. Its method is to attach each counsel to the framework result that justifies it, and to state explicitly which counsels are theorems (conditional only on the

axioms) and which are heuristics (conditional on present cost curves). The distinction matters for a founder for the same reason it matters to the framework itself: a conditional theory should not be confused with a prophecy, and a heuristic should not be confused with a law. Both are useful. Both are different.