Post-AI Adoption: The More Things Change, The More They Remain the Same

MACHARIA BARII

bariimacharia@gmail.com

August 2026

Abstract

This paper argues that artificial intelligence, despite its revolutionary capabilities, will not fundamentally alter the ancient human social structure defined by a minority of "framers"—those who ask questions, formulate techniques, and design technologies—and a majority of "adopters"—those who consume and drift within the answers, methods, and tools provided by that minority. Through historical analysis of knowledge authorities and technological diffusion— from oracles to encyclopedias, from the plow to the printing press, from search engines to generative AI—we contend that AI accelerates existing asymmetries rather than democratizing agency. The paper concludes that the primary societal risk is not obsolescence or job loss, but the deepening invisibility of *framing itself*, leaving the adopters unaware that a question was ever asked, a method ever devised, or a tool ever designed. We then extend the argument to robotics, contending that robots lower the cost of *embodied execution*—the last domain of adopter agency—and that health, security, social structure, labor, and luxury all collapse into the same ancient pattern. We offer no utopian prescription, only a sober reflection: the more perfect the answer machine, the less visible the act of creation becomes.

Contents

1 Introduction

In 2022, a student asked ChatGPT to explain quantum entanglement. In 2026, a corporate executive approved a market-entry strategy generated entirely by an AI agent, without reading a single source document. Neither action is remarkable today. But consider the structure of these moments: in both cases, an answer was received without a question being fully formed, interrogated, or even consciously owned by the recipient.

Now consider a deeper layer: the student did not design the transformer architecture that generated the answer. The executive did not formulate the fine-tuning methodology that shaped the model's priorities. Both are consuming not just answers, but an entire *technological and epistemic infrastructure* they had no hand in creating.

The dominant narrative surrounding artificial intelligence is one of profound disruption. We are told that AI will democratize expertise, personalize education, collapse information asymmetries, and elevate human potential to unprecedented heights. Technology optimists frame AI as the great equalizer—a tutor, a collaborator, and a liberator of human cognition.

This paper offers a contrarian thesis: **AI will not redistribute intellectual or technical agency. It will reinforce, accelerate, and render invisible the ancient divide between those who frame questions, formulate methods, and design technologies—and those who consume, adopt, and drift within the answers, techniques, and tools provided by that minority.** The technology is new, but the social pattern is as old as civilization itself.

Our argument proceeds in six parts. First, we establish the historical precedent—the stable ratio of framers to adopters across millennia. Second, we examine the internet as a false dawn of democratization. Third, we analyze the unique features of AI that make this rupture particularly insidious. Fourth, we describe the emerging aristocracy of "prompt architects" and "method-designers." Fifth, we extend the argument to techniques and technologies— showing that the same minority invents what the majority adopts—and to physical robotics, which lower the cost of embodied execution and complete the second remove. Sixth, we confront the uncomfortable mirror AI holds up to human nature. We conclude that the more perfectly we engineer answers, tools, and systems, the more clearly we expose our ancient preference for ready-made solutions over the labor of creation.

2 Historical Precedent: The Eternal Divide

2.1 Ancient Authorities

The Oracle at Delphi did not answer questions; she interpreted omens, and her utterances were famously ambiguous. Yet her authority rested on a simple social contract: the many accepted her answers, and the few—kings, generals, legislators—determined which questions to bring to her. The framing of the inquiry was a privilege of power.

Similarly, in ancient Mesopotamia, scribes were not merely record-keepers. They were the gatekeepers of interpretation and *technique*. Cuneiform tablets encoded laws, astronomical predictions, medical prescriptions, and agricultural calendars. The farmer received the planting date; the scribe formulated the astronomical method that determined it. The farmer consumed the technique; the scribe designed it.

This pattern—experts translating complex systems into simple directives for the masses—persisted through Egyptian priesthoods, Roman augury, and Confucian scholar-bureaucracies. In each case, the minority framed the questions *and* designed the methods; the majority executed and adopted.

Plato's *Republic* articulated this structure explicitly: the philosopher-kings *know*; the guardians *enforce*; the producers *obey*. Plato did not lament this hierarchy; he considered it natural and necessary. Two millennia later, the structure remains recognizable.

2.2 The Medieval Consolidation

The Catholic Church's monopoly on scriptural interpretation offers a stark example. The Bible was available only in Latin—a language of the clergy. The laity received sermons, not scripture. When vernacular translations emerged in the 14th and 15th centuries, they were suppressed not because the Church feared literacy, but because they feared *interpretive agency* and, by extension, the capacity to *formulate alternative theological methods*. If every Christian could frame their own questions about salvation and devise their own devotional practices, the priesthood's authority collapsed.

This is not a critique of religion; it is an observation about knowledge and technique hierarchies. The Reformation, for all its transformative power, did not make every Christian a theologian. It simply shifted authority from one institution to many, and from Latin to vernacular. The ratio of framers to adopters remained stable. Most believers accepted the interpretations and devotional methods of their new pastors, just as they had accepted those of their old priests.

2.3 The Encyclopedic Era

Diderot's *Encyclopédie* (1751–1772) was a radical project: to make all human knowledge accessible to the literate public. It was a democratizing force, certainly. But it was also a consolidating force. The *Encyclopédie* presented answers, techniques, and classifications—systematized, authoritative. It did not teach readers how to challenge its categories, question its biases, or devise alternative methods. It invited consumption, not co-creation.

The same pattern repeated with the Britannica, the World Book, and eventually Wikipedia. Each iteration lowered the barrier to *access*, but none meaningfully increased the proportion of the population that actively framed original inquiries, challenged established taxonomies, or generated new epistemic or technical frameworks.

2.4 Technology Diffusion: The Minority Invents

The history of technology is a history of minority invention and majority adoption. Consider the plow: a small number of metallurgists and agricultural innovators designed and refined it over centuries. The vast majority of farmers simply used it. They did not understand the metallurgy, the soil mechanics, or the ergonomics. They adopted the tool because it worked.

Consider the printing press: Gutenberg and a small cohort of craftsmen, engineers, and financiers built and scaled it. The reading public did not design the press, formulate the ink chemistry, or devise the type-casting methods. They consumed the books. They adopted the technology of reading itself as a given.

Consider the internal combustion engine, the electrical grid, the computer, the smartphone. In every case, a tiny fraction of the population—perhaps 0.1%—formulated the underlying physics, designed the engineering, built the infrastructure, and wrote the software. The remaining 99.9% adopted, used, and drifted within those systems without ever understanding—or needing to understand—how they worked.

This is not a failure; it is the division of labor that enables civilization. But it is also a concentration of agency. The adopters are not merely passive; they are structurally dependent on the framers, whether they recognize it or not.

2.5 The Stable Ratio

Historical evidence suggests that across cultures and centuries, approximately 1–5% of any given population functions as "framers"—individuals who generate novel questions, formulate new methods, design technologies, and create systems. The remaining 95–99% consume, apply, adopt, and drift within the products of that minority. This is not a moral judgment; it is a descriptive observation. The distribution appears robust to changes in literacy, technology, and political system.

The question for this paper is whether artificial intelligence represents the first force capable of shifting this ratio—or whether it will, like every previous knowledge and technology revolution, be absorbed into the ancient pattern.

3 The Internet Detour: The Illusion of Democratization

3.1 The Promise

When the World Wide Web went mainstream in the mid-1990s, utopian rhetoric was pervasive. "Information wants to be free." "The wisdom of the crowd." "The end of gatekeepers." These slogans promised a future where every individual, armed with a modem and a browser, could access the totality of human knowledge and participate in global discourse—not just as consumers, but as creators.

In retrospect, this promise was naive but sincere. The internet did, in fact, make an unprecedented volume of information, tools, and platforms available to an unprecedented number of people. Open-source software, blogging, YouTube, and later social media allowed millions to *publish*, if not to *create from first principles*.

3.2 The Reality: Search and Platforms as New Priesthoods

Google, founded in 1998, became the world's primary answer engine. Its PageRank algorithm was not neutral; it prioritized popularity, backlinks, and, increasingly, advertising and SEO optimization. Users typed fragments of questions—or simply keywords—and received ranked lists of links. The vast majority never scrolled past the first page. They accepted the top result as authoritative.

More importantly, the *techniques* of search—how to structure a query, how to evaluate sources, how to verify claims—were not taught by the platform. They were assumed. The minority of power-users learned Boolean operators, advanced search syntax, and source criticism. The majority typed a few words and trusted the algorithm.

Social media platforms (Facebook, Twitter/X, TikTok, Instagram) further entrenched the pattern. They are not designed for inquiry or creation; they are designed for engagement and consumption. Their algorithmic amplification favors certainty, outrage, and simplicity over nuance, doubt, and complication. The most successful content creators are those who provide clear, emotionally resonant answers and ready-made cultural scripts—not those who pose uncomfortable questions or devise radical new methods.

3.3 The Myth of the "Prosumer"

The term "prosumer"—one who both produces and consumes—was popularized in the early internet era. It captured the hope that ordinary users would become creators. In reality, the ratio of producers to consumers on any platform follows a power law. On Wikipedia, less than 1% of users contribute the majority of content. On YouTube, a tiny fraction of creators capture the overwhelming majority of views. On GitHub, a small core of developers maintain the essential open-source libraries that millions use without ever reading the source code.

The vast majority are adopters, not framers. They use the platform, the library, the template, the app. They do not design the protocol, write the algorithm, or debug the infrastructure. The internet democratized *access to adoption*, not *capacity for creation*.

The new castes were absorbed too. Each technological wave produced its own dazzling new hobbyist caste—and each turned out to be a new flavor of framer, not a redistribution of it. The computer gave us the *hacker* and the *cracker*: small subcultures who probed, modified, and broke systems, while the vast majority merely used them. Social media gave us the *influencer*, the *YouTuber*, and the *TikToker*: a charismatic minority who framed and performed content for the admiring majority who consumed it. These castes look like democratization—everyone, in principle, can start a channel. But they follow the same power law. A tiny fraction capture the attention and revenue; the crowd consumes and drifts within their framings. The castes are new; the asymmetry is ancient.

3.4 Data on Critical Thinking and Technical Literacy

Empirical research on internet use and cognitive outcomes is sobering. Studies consistently show that internet access increases political polarization, confirmation bias, and echo-chamber effects (Sunstein , 2017; Bail et al. , 2018). It does not, on average, improve critical thinking skills or technical literacy. This is not because the internet is flawed; it is because tools are used within existing cognitive habits. The internet is a magnifier of human nature, not a transformer of it.

Technical literacy—the ability to understand, debug, or modify the tools one uses—has not increased meaningfully. Most smartphone users cannot read network logs, modify firmware, or write a simple script. They are not framers of the technology; they are adopters of the interface.

The lesson for our AI era is clear: **access does not equal agency**. The internet gave everyone the library, the platform, and the app store. It did not give everyone the curiosity, discipline, or training to become framers.

4 The AI Rupture (That Isn't)

Artificial intelligence is not merely a faster internet. It is a fundamentally different kind of knowledge and technology— one that generates *original* answers to *novel* prompts and autonomously performs *complex tasks* that previously required human framing. This is genuinely unprecedented. Yet we argue that it will not disrupt the ancient hierarchy; it will deepen it. Here is why.

4.1 The Invisible Question

With predictive text, autocomplete, and proactive agents, AI often answers before the user fully articulates a question. Consider Google's AI Overviews: a user types "why is my plant dying" and receives a synthesized answer without ever specifying the plant species, the watering schedule, or the light conditions. The framing—what factors matter, what to prioritize, what to ignore—is performed by the model, not the human.

This is a subtle but profound shift. In the print era, the reader had to seek an answer. In the search era, the user had to type a query. In the AI era, the system anticipates, suggests, and completes. The adopter does not even know that a choice of framing occurred. They receive a fluent, plausible answer and move on.

The framer, by contrast, notices the framing. They push back: “No, I meant the rare orchid, not the common fern. And I’m in a dry climate.” They interact recursively. They interrogate the model’s assumptions. This is the distinction that

matters.

4.2 The Invisible Method

Beyond answers, AI also automates *methods*. It generates code, drafts legal documents, designs marketing campaigns, predicts supply-chain disruptions, and even proposes scientific experiments. The user receives a complete *process*—a technique, a workflow, a deliverable—without having to formulate that process themselves.

In pre-AI eras, a manager might have designed a campaign structure; now they prompt an AI to generate fifty variants. A biologist might have designed an experimental protocol; now they ask an AI to suggest one. A software engineer might have architected a system; now they prompt the AI to generate boilerplate and refactor legacy code.

The majority of users are not *designing* these methods; they are *selecting* from AI-generated options. They are adopters of techniques, not framers. The AI becomes the method-designer, and the human becomes the method-curator. This is a step-change in the invisibility of technique.

4.3 The Opacity of Training Data and Model Design

AI models are trained on massive, opaque corpora—the internet, books, academic papers, social media, and more. Their outputs reflect the biases, omissions, priorities, and prejudices of that data, filtered through reinforcement learning from human feedback (RLHF). The RLHF process itself embeds value judgments about what constitutes a good, safe, or helpful answer, method, or output.

These value judgments are not neutral. They are encoded preferences—engineered by a tiny fraction of AI researchers, ethicists, and product managers. Yet to the average user, they appear as objective outputs. The system does not say, “Here is one way to frame your question and one method to solve it, influenced by English-language internet culture from 2010–2024.” It simply says, “Here is the answer and the approach.”

The adopters have no visibility into this. They accept the output as authoritative. The framers, however, probe the model’s reasoning, ask for sources, compare outputs, challenge assumptions, and—crucially—consider *alternative architectures*. They treat the AI as a collaborator, not an oracle, because they know it is neither omniscient nor neutral.

4.4 The Illusion of Conversation and Competence

Chat interfaces—ChatGPT, Claude, Gemini, etc.—simulate human dialogue with remarkable fluency. They use first- person pronouns, conversational pacing, and empathetic phrasing. This design intentionally lowers psychological defenses. Users treat the AI as a helpful assistant, not a probabilistic text generator.

This illusion is dangerous precisely because it feels natural. In human conversation, we trust fluency; we attribute understanding and competence to our interlocutor. The same heuristic applies to AI, making it *harder* to doubt its outputs or question its methods. The oracle is now a friendly voice, not a cryptic priest. The adopters are more likely to trust, not less.

4.5 Speed Kills Curiosity and Craft

The single most underappreciated effect of AI is the evaporation of friction. In the pre-digital era, finding an answer or learning a technique required effort: flipping through a card catalog, skimming a book, consulting an expert, or practicing a skill. That friction was uncomfortable, but it was also productive. It forced reflection. It gave time for the mind to wander, to consider alternatives, to formulate better questions and better methods.

AI removes that friction entirely. Answers are instant, methods are generated on demand, and outputs are fluent and free. The cognitive cost of *framing* has not changed, but the cognitive cost of *adopting* has plummeted. Relative to the cost of framing, adopting is now exponentially cheaper. Human beings are cognitive misers (Kahneman , 2011); we choose the path of least mental expenditure. AI makes the easiest path even easier—straight to the output, bypassing curiosity, craft, and understanding.

Empirically, we expect to see a decline, not an increase, in framing behavior among general users, because the system makes it optional and even awkward to frame further.

5 The New Aristocracy: Framers 2.0

If the majority drifts, a minority thrives. Who are they, and what do they do differently?

5.1 Who They Are

The new framers are not necessarily the wealthy, the highly educated in elite institutions, or the technically trained in the traditional sense. They are individuals—across domains—who treat AI as a *collaborative sparring partner* rather than an oracle or an automation engine. They do not ask once and accept. They do not adopt the first generated method. They engage recursively:

- “Why do you think that? What are the counterarguments?”

- “What if we invert the core assumption?”
- “Simulate that outcome under these three different conditions.”
- “That method seems inefficient. Let’s redesign it from first principles.”
- “What are the hidden biases in this approach?”

This is not a different prompt; it is a different *relationship*. The AI is a tool for generating more questions, better methods, and novel designs—not just executing existing ones.

5.2 What They Do

The framers use AI for:

- **Probing biases**: They ask for sources, examine training data limitations, and test for contradictions.
- **Counterfactual exploration**: They simulate alternative histories, futures, and conditions.
- **Cross-domain synthesis**: They force the AI to connect disparate fields—biology and economics, literature and engineering—to generate novel frameworks.
- **Method design**: They do not accept AI-generated workflows at face value; they deconstruct, hybridize, and improve them.
- **Meta-cognition**: They ask the AI to reflect on its own reasoning, identify its own uncertainties, and suggest better ways to frame the original problem or method.
- **Red-Teaming**: They deliberately try to break the AI’s outputs to understand its failure modes and limitations.

These practices are not intuitive. They require training, practice, and, crucially, *psychological tolerance for ambiguity and being wrong*. The AI does not teach these traits; it assumes them.

5.3 The Feedback Loop

The framers have a compounding advantage. They generate richer interactions, which produce better outputs, which produce more sophisticated questions and methods, which produce even better outputs. They also contribute to open- source model fine-tuning, research, prompt engineering communities, and method repositories—insular networks that further concentrate expertise.

This is the new aristocracy: not a hereditary class, but a cognitive and technical one. It is demographically diverse but epistemically narrow. And it remains stubbornly small—still 1–5% of the population, because the skills it requires are scarce, hard to teach, and socially under-rewarded.

5.4 The Widening Gap

The gap between the framers and the adopters is not static; it is widening. The framers use AI to simulate, prototype, design, and conquer domains at accelerating speed. The adopters use AI to automate their consumption, making their intellectual and technical passivity more efficient and more comfortable. The former create leverage; the latter experience convenience.

This gap is not about intelligence. It is about *agency*. And agency is not distributed equally.

6 The Technique-Technology Parallel: The Minority Invents, the Majority Adopts

We now extend the core thesis beyond *questions and answers* to the broader domain of *techniques, methods, and technologies*. This is the missing dimension of the argument: the same pattern that governs epistemic inquiry governs practical creation.

6.1 The Historical Pattern of Technical Framing

As noted in Section 2.4, every significant technology in human history was invented, refined, and scaled by a tiny minority. The adopters—the vast majority—did not understand the underlying physics, chemistry, or engineering. They did not formulate the methods. They simply used the tools.

Consider these examples:

- **Agriculture**: The plow, crop rotation, and fertilization methods were developed over millennia by a small number of innovators. The billions of farmers who followed adopted these techniques without understanding the agronomy or the metallurgy.
- **Medicine**: Germ theory, antibiotics, and surgical techniques were formulated by a handful of scientists and physicians. The global population adopted vaccines and hygiene practices without comprehending immunology.
- **Energy**: The steam engine, internal combustion engine, electrical grid, and nuclear fission were designed by a tiny cohort of physicists, engineers, and industrialists. The masses adopted electricity, automobiles, and central heating as givens.

- **Digital Computing**: Transistors, chips, microprocessors, operating systems, and the internet were built by less than 0.1% of humanity. The other 99.9% use smartphones, browse the web, and stream video without understanding the stack beneath.

In each case, the adopters benefit enormously. They are not exploited; they are enabled. But they are also *structurally dependent*. They cannot repair, debug, redesign, or replace the technologies they rely on. They are passengers, not pilots.

6.2 AI as Method-Designer

AI accelerates this pattern by automating *method design itself*. Previously, even adopters had to learn some techniques—how to use a spreadsheet, how to write a basic query, how to operate a machine. With AI, the method is generated on the fly:

- A marketer no longer needs to learn A/B testing design; the AI generates the test plan.
- A policy analyst no longer needs to learn statistical modeling; the AI builds the regression.
- A software developer no longer needs to learn a new framework; the AI writes the boilerplate and refactors the code.
- A biochemist no longer needs to design a screening protocol; the AI suggests one based on literature.

The adopter becomes a *selector*, not a *designer*. They choose among AI-generated methods, but they do not understand the principles underlying those methods. They are one level further removed from the act of creation.

This is not inherently bad. Specialization and division of labor are the engines of civilization. But it deepens the asymmetry of agency. The framers—those who design, debug, and improve the AI systems themselves—become even more critical and even more concentrated.

6.3 The Invisibility of Infrastructure

There is a deeper point: the adopters are not even aware of the infrastructure that enables their AI-generated methods. They do not see the data centers, the cooling systems, the chip fabrication plants, the energy grids, the submarine cables, or the thousands of engineers who maintain the stack. The technology is a black box, and the methods it produces are black boxes within that black box.

This invisibility is the ultimate enabler of drift. When you cannot see the machinery, you cannot question it. When you cannot question it, you cannot reframe it. When you cannot reframe it, you are structurally dependent on those who

can.

6.4 The Paradox of Democratization

AI advocates often claim that AI "democratizes" expertise—it allows anyone to become a marketer, a coder, a lawyer, or a designer. This is true in a narrow sense: the output is more accessible. But it is false in a deeper sense: the *capacity to frame the domain*—to understand first principles, to troubleshoot failures, to innovate beyond the AI's training—remains concentrated.

We can illustrate with a medical analogy:

- **Pre-AI**: A general practitioner could prescribe standard treatments. A specialist could diagnose complex cases. A researcher could design new trials. The hierarchy of framing was visible.
- **Post-AI**: A general practitioner can use an AI differential diagnosis tool. A specialist can use an AI to suggest novel treatment combinations. A researcher can use an AI to design clinical trial protocols. The outputs are better, but the *relative agency* remains the same. The GP still cannot design trials. The specialist still cannot formulate new molecular pathways. The researcher still cannot build the AI.

The technology elevates all levels, but the *gap* between levels—the difference in epistemic and technical agency—does not shrink. It may even widen, because the framers gain a multiplier that the adopters do not.

6.5 Robots: The Second Remove

We now extend the argument to its most consequential frontier: physical robotics. AI lowers the cost of *cognitive adoption*—the effort required to receive an answer or a method without framing it. Robots lower the cost of *physical adoption*—the effort required to *execute* a framed task without learning the skill. This is the second remove: the adopter no longer frames the task (AI does), and no longer performs it (the robot does). The adopter is reduced to a single remaining role: *intention*.

What robots lower. Television lowered the cost of entertainment consumption. The computer lowered the cost of information processing. The smartphone lowered the cost of connectivity. AI lowered the cost of cognitive adoption—the question, the method, the draft, the answer. Robots lower the cost of *embodied execution*. The adopter no longer needs to learn to cook, drive, clean, build, heal, farm, or fight. The robot observes the environment, frames what needs doing, decomposes it, and acts—often without even being prompted. The last physical residue of the adopter's agency—the hands that *did*—is outsourced.

The pattern propagates. The design of the robot, its sensor suite, its task-decomposition architecture, its error-recovery strategies, its safety constraints, and its failure modes are formulated by the same 1–5% framer minority: roboticists, ML engineers, sensor designers, safety ethicists, and task-framers. The remaining 95–99% experience robots the way their ancestors experienced electricity: it works, they do not know why, and they cannot repair, reprogram, interrogate, or redesign it. The ratio does not move. It never did.

The gap re-widens. AI had briefly lowered the cost of *climbing*—an adopter could, with sufficient curiosity and effort, learn to prompt, probe, and eventually frame. Robots reintroduce a *physical and financial barrier* that no amount of curiosity can cross. One cannot "tinker" with a robot the way one can tinker with a prompt. The robot's methods are embedded in actuators, sensors, firmware, and safety interlocks—inaccessible to the vast majority. The path from adopter to framer, temporarily widened by AI, narrows again.

6.5.1 Health

The health impact is double-edged and asymmetric. On the surface, robots will deliver unprecedented medical outcomes: surgical robots with millimeter precision, home-care robots for the elderly, diagnostic robots in underserved regions, and prosthetics that restore mobility. The adopter experiences better health, longer life, and greater comfort. But the *framing of health*—what "healthy" means, which conditions get research priority, which populations get robotic access, and what risks are acceptable—is decided by the framer minority. The adopter receives the treatment without understanding the diagnosis, the mechanism, or the alternative framings that were rejected. Worse, robotic health monitoring creates a new asymmetry: the framers control the sensors, the data, the algorithms, and the thresholds. The adopter's body becomes a data stream that they themselves cannot read. The physician becomes a curator of robotic recommendations. The patient becomes a recipient of robotic care they cannot question. Health becomes a framed service, not an understood state.

6.5.2 Security

Security in a robotic world is the security of the framer, not the adopter. Domestic robots control access to the home. Delivery robots move goods. Industrial robots control infrastructure—water, energy, food processing. Autonomous vehicles control movement. Security robots patrol borders, streets, and institutions. For the adopter, security appears to improve: fewer break-ins, fewer accidents, faster responses. But the adopter's security is now *rented*, not held. The robot decides who enters, what is flagged as suspicious, and when force is used. The thresholds, the training data, the bias audits, and the escalation rules are designed by the framers. A malfunction is not a repair; it is a service call. A disagreement with the robot's judgment is not an argument; it is a user ticket. The adopter's physical safety becomes an opaque service provided by a black-box system they cannot inspect, appeal, or override. The framers hold the keys. The adopters hold the keywords.

6.5.3 Social Structures

Robots complete the trajectory of the "quiet aristocracy." With AI alone, the class divide was epistemic— framers thought differently than adopters. With robots, the divide becomes *embodied*. The framers design, maintain, and understand the robots that serve them. The adopters are served by robots they cannot repair, question, or supersede. The class structure is no longer about knowledge; it is about *control of the physical environment*. The adopters live in a world where doors open for them, meals arrive, floors clean themselves, and vehicles navigate—all without the adopter understanding a single mechanism. The framers live in the same physical world, but they also live *inside the design*: they can reprogram, override, and improvise. The social hierarchy becomes invisible because it is now *physical*. The adopters cannot even see the hierarchy, because they never see the machinery. They see results. The gap between the classes becomes the gap between those who *live within the system* and those who *live as the system*.

6.5.4 Jobs

The labor market bifurcation that AI began, robots complete. Routine physical work—manufacturing, agriculture, construction, logistics, domestic service, delivery, security, elder care—is automated by physical systems. The adopter who previously sold their hands now has nothing to sell. They do not disappear; they are *absorbed* into the service layer: selecting, supervising, and consuming robotic output. The framers capture the surplus of robotic productivity. The adopters receive the convenience. The jobs that remain—robot design, sensor engineering, task- framing, safety certification, robotic maintenance, infrastructure operation—are precisely the jobs the adopters cannot do. The jobs that disappear are precisely the jobs the adopters could. The adopters do not lose employment in a dramatic sense; they lose *leverage*. They are paid to watch robots, to supervise robots, or not to be paid at all. The framers are paid to design the robots. The value accrues upward. The ratio holds.

6.5.5 Wealth and Productivity Gaps

The economic divergence between framers and adopters in the robotic world is not merely qualitative; it can be stated in orders of magnitude. The productivity of a framer scales with the number of robots they can deploy, supervise, and improve. The productivity of an adopter is bounded by their own body and their own willingness to consume. The gap, therefore, compounds multiplicatively and without an upper limit.

Productivity. A framer commanding a fleet of a thousand robots—each working twenty-four hours a day, never tiring, never striking, never demanding health care—operates a worker base the size of a small town, yet is a single employee. Their output-per-person is the output of a town. The adopter, by contrast, contributes only their own attention: a few robot tasks selected, a few results consumed. The framer's productivity rises roughly with *fleet size* and *automation depth*; the adopter's rises barely at all. In the fully robotic world, the median framer's output can exceed the median adopter's by a factor of thousands to millions. This is not hyperbole; it is arithmetic. Multiply one mind by ten thousand machines and compare it to one mind by zero machines.

Wealth. Wealth follows productivity. The framer captures the surplus their robot fleets generate, compounding it into further fleets, further autonomy, and further leverage. The adopter receives the low, diffuse residual of that surplus: the discounted or subsidized consumption of robot service. Two forces widen the gap beyond productivity itself. First, *returns to scale and network effects*: the more robots a framer owns, the cheaper each additional robot is to serve and improve, packing wealth at the top. Second, *the ownership of the means of automation*: robots are capital, and capital concentrates. The Gini coefficient of a robot-saturated economy would be extreme—wealth held not across a broad middle class but in the hands of the framing class that owns and designs the machines. The adopters are not poor in absolute terms; they are absolutely dependent. They live well on the trickle-down service, but they own almost none of the machine.

The implausibility of redistribution. A naive hope is that robotic abundance can simply be taxed and redistributed to the adopters. The difficulty is structural, not technical. The framing class controls the machines, the data, the algorithms, the infrastructure, and the very ability to measure value. A transfer system depends on the class that owns the machines to size, verify, and remit the surplus. Adopters are the *tenants* of the system; they cannot compel the landlords to set their own rent downward. Moreover, because value increasingly moves from physical output to invisible design, intelligence, and coordination— the framers' territory—the tax base itself migrates into the very black boxes the adopters cannot see. The trickle is real but discretionary. The gap, in the robot-adopted world, is not a policy failure awaiting a fix; it is the natural equilibrium of the ancient ratio, now expressed in capital.

6.5.6 Luxury Living

The ultimate expression of the adopter condition is the robotic luxury household: a home where food is prepared, spaces are cleaned, environments are adjusted, schedules are managed, and needs are anticipated without a single act of physical or cognitive effort. This is not a prediction; it is the observable trajectory. Air conditioning, dishwashers, robotic vacuums, automated lighting, and voice assistants already approximate it. The fully robotic home extends the pattern: the adopter experiences a life without friction, without skill, and without understanding. The luxury is real. It is also the most complete form of dependency ever devised. The adopter in the robotic household is not a resident; they are a *tenant*. They live in an environment they cannot repair, modify, or escape. The framers design the home. The adopters dwell in it. The more luxurious the adoption, the more complete the framing—and the more invisible the hierarchy.

The final prediction. When robots are as fully adopted as television, computers, and smartphones—and they will be—the pattern will have propagated unchanged. The robot will be as invisible as the electrical grid: always present, never understood, unrepairable, unquestioned. The framers will design it, the adopters will live within it, and the 1–5% ratio will hold as it has for five thousand years. The oracle is new. The toolmaker is new. The hands are new. The crowd is ancient. The user never moved.

6.6 The New Hobbyist Castes of the AI and Robotics Era

If computers gave us the hacker and the cracker, and social media gave us the influencer, the YouTuber, and the TikToker, what do mass AI adoption and mass robotics adoption give us? The same logic predicts a fresh parade of castes: new roles that *look* like creative liberation and again collapse into the framer–adopter power law. The novelty of the castes should not be mistaken for a change in the ratio.

The AI era's castes.

- **The Promptwright**, *prompt engineer* and *prompt-jailbreaker*: the caste who treats the prompt as a craft—sculpting, chaining, and weaponizing the model—while the majority merely asks. The jailbreaker is the hacker of the AI era, probing safety rails and seeking the unaligned output; the promptwright is the influencer of the prompt world, building followings on their framing skill.
- **The Red-Teamer** and *AI auditor*: the minority who deliberately test models for bias, failure modes, and hidden agendas. They are the new security researchers, publishing the discovered limits of the system.
- **The Synthetic Artist**: creators who build wholly generated music, film, literature, and worlds— framing the aesthetic, the prompt-vectors, the pipelines—for audiences who consume the output without understanding the craft.
- **The Avatar-Gamer** and *AI-persona tamer*: those who role-play, direct, and steer autonomous characters and NPCs in persistent worlds; a small director caste over a vast player base.

- **The Fine-Tuner** and *Custom-Modeler*, the new open-source maintainers: a tiny core who fine-tune, audit datasets, and architect models that millions adopt without ever opening the code.

The robotics era's castes.

- **The Bot-Tamer** / *robot choreographer*: the caste who frame robot tasks, sequences, and behavior—training the domestic bot, tuning the autonomous vehicle, directing the fleet—while the majority simply summons the finished behavior.
- **The Robot-Saunterer** / *mechanic-collector*: a hobbyist minority who buy, disassemble, mod, and restore robots as both craft and collection, the robotics analogue of the home-garage tinkerer or the watchmaker.
- **The Arena-Caster** and *robot-sports fan*: with robot combat, dogfighting, and racing spectacles, a live-entertainment industry emerges around who can best *frame* robot strategy—the robotics analogue of the influencer, performatively directing machines for the crowd.
- **The Drone-Strafer** / *sensor-hobbyist*: hobbyists who build, modify, and race autonomous agents and drones, exploring the physical and legal edges of the robot world the way early hackers explored the edges of the network.
- **The Safety-Tester** / *robotics red-teamer*: a small corps who deliberately stress robots to the point of failure—probing their limits, their safety rails, and their liability—the robotics analogue of the cracker, at the boundary of the permitted.

In every caste, the same two facts hold. First, these roles require *agency*: an inclination to frame, probe, and build, not merely to consume. Second, they obey a power law: a tiny fraction of each caste commands attention, revenue, and influence, while the surrounding crowd—the hobby's vast audience— adopts and drifts. The castes multiply and glamorize; the ratio does not move. The hobbyist is new. The creator is new. The farmer and the fan are ancient.

7 The Ultimate Mirror: AI Reveals, Does Not Revolve

7.1 Human Nature as Cognitive and Technical Misers

Kahneman's dual-process theory (Kahneman , 2011) distinguishes System 1 (fast, automatic, intuitive) from System 2 (slow, deliberative, analytical). Human beings default to System 1 whenever possible. Curiosity, questioning, critical reasoning, and technical design are System 2 activities—cognitively expensive, effortful, and uncomfortable.

Evolutionarily, this makes sense. Survival depended on quick, actionable answers and reliable, proven tools: "Is that a predator?" "Can I eat this?" "Which tool works?" "Who can I trust?" Speculative questioning and experimental method-design were luxuries of stable, well-fed societies. The cognitive and technical miser is not a failure; it is a successful adaptation.

7.2 AI as the Perfect Mirror

AI does not create a new human nature; it reveals the old one. The technology is neutral; the social structure is not. The masses want certainty, convenience, and ready-made solutions. AI provides these at scale, with unprecedented fluency and speed. The leaders want leverage, not truth. AI provides leverage at scale, with unprecedented power.

If we observe that most people use AI to get quick answers, generate standard methods, and automate routine tasks —rather than to interrogate assumptions, redesign workflows, or create novel architectures—this is not a failure of the technology. It is a reflection of our species. The oracle and the toolmaker are new; the crowd and the user are ancient.

7.3 Opiate or Mirror?

One could argue that AI is an opiate—a soothing, addictive comfort that numbs the drive to question, create, or craft. There is truth in this. AI-generated answers and methods are designed to be satisfying, plausible, and frictionless. They reduce anxiety. They provide closure. They allow the user to move on.

But the opiate metaphor implies that without the opiate, the user would be more awake, more curious, more creative. This is not supported by history. The same population that today accepts AI-generated strategies was, in previous centuries, accepting priestly sermons, encyclopedia entries, Google search results, and pre-designed tools without ever examining their construction.

AI is therefore better understood as a mirror. It shows us who we have always been: a species that prefers answers to questions, ready-made methods to experimental design, and adoption to creation. The mirror does not create the reflection; it only makes it harder to ignore.

8 Counterarguments and Rebuttals

9 Implications and Predictions

9.1 Economic Implications

The economic value of "answer production" and "routine method execution" collapses. Content generation, analysis, summarization, basic coding, standard legal drafting, and routine design become commoditized. The value of "problem framing," "method design," "system architecture," and "novel creation" skyrockets. But these are scarce, non-scalable human skills. The labor market bifurcates: a small cohort of high-agency framers, and a large cohort whose routine cognitive and technical tasks are automated. This is not replacement; it is *revaluation*. The framers capture the surplus; the adopters experience convenience but declining wage leverage.

This bifurcation is not a mildly widened inequality; it is a quantitative break, and robotics makes it blunt. As developed in Section 6.5.5, the framer's productivity scales with the fleet of machines they deploy—one mind multiplied by ten thousand robots—while the adopter's productivity is capped by their own attention. The resulting wealth gap is measured not in percentages but in *orders of magnitude*, concentrated by returns to scale and by the concentration of capital itself. The ruling economic question of the post-AI world is therefore not *how to create abundance*—robots are solving that—but *who owns the machines that create it*. That question has an old and unwelcome answer: the framers do. The adopters consume the abundance; they do not own its source. The gap between them is the gap between those who hold capital in a world where capital works, and those who hold only the capacity to consume its output.

9.2 Political Implications

AI-generated answers and methods will be weaponized for opinion manipulation, consent manufacturing, and

infrastructure control. The adopters, who accept outputs without scrutiny, will be increasingly susceptible to subtle framing embedded in system prompts, fine-tuning, training data, and default method-generation. The battleground moves from "what is true" and "what works" to "what question gets asked" and "what method gets generated"—and who controls the AI's default framing. This is the new geopolitics of information and technique.

9.3 Educational Implications

Educational systems will struggle to adapt. They are built to test *answers* and *standard methods*— multiple-choice quizzes, standardized exams, fact recall, and procedural repetition. They do not test *question- asking*, *method-design*, or *critical creation*, let alone recursive, counterfactual, or meta-cognitive framing. AI will make answer-based and method-based assessment obsolete, but institutions will be slow to replace them with framing-based assessment. The gap between what schools measure and what society values will widen, to the detriment of the adopters.

9.4 Social Implications

A quiet, invisible class structure will emerge. It will not be based on wealth, race, or nationality—though those will correlate. It will be based on *epistemic and technical agency*: the ability and inclination to treat AI as a collaborator rather than an oracle, to design methods rather than select them, and to create technologies rather than use them. This class will not be revolutionary; it will be quietly ascendant. The adopters will not revolt, because they will not recognize the structure. They will feel vaguely unsatisfied, but they will receive ever more fluent answers to their unexamined questions and ever more polished methods for their unexamined tasks. The system will stabilize— faster, smoother, and with fewer visible strings.

9.5 Long-Term Prediction

Over decades, the gap between framers and adopters may become the defining social cleavage of the post-AI era. It will supersede traditional class, educational, and even national divisions in its influence on life outcomes, political power, and cultural production. The framers will not simply be richer; they will be *differently human*—shaping the cognitive, technical, and social environment in which the adopters drift. The adopters will not be oppressed in the traditional sense; they will be comfortable, entertained, and well-served. But they will be passengers on a journey they do not navigate, in a world they did not build, using tools they do not understand.

10 Conclusion

We began with a provocation: artificial intelligence, despite its revolutionary capabilities, will not fundamentally alter the ancient human social structure defined by a minority of framers—those who ask questions, formulate methods, and design technologies—and a majority of adopters—those who consume answers, execute methods, and use technologies without understanding their construction.

Through historical analysis—from oracles to encyclopedias to search engines, from the plow to the printing press to the smartphone—we have seen that each new knowledge and technology revolution lowers the barrier to *access and adoption* but does not meaningfully increase the proportion of the population that exercises *epistemic or technical agency*. The internet, for all its promise, amplified the existing pattern rather than breaking it.

AI introduces new affordances: invisible framing, opaque training, conversational fluency, instant speed, and automated method-generation. Each of these features makes it *easier* to consume answers and adopt techniques—and *harder* to recognize that questions were framed, methods designed, and technologies built by a minority. The adopters are not simply accepting outputs; they are no longer aware that outputs have origins.

The framers—the 1–5%—will use AI to multiply their cognitive and creative power, generate novel frameworks, design superior methods, and accelerate their dominance. They will not abolish hierarchy; they will refine it. The gap will widen, not close.

In the end, the more perfectly we engineer the machine that answers, generates, designs, and automates, the more clearly we expose the uncomfortable truth: that most of humanity never really wanted to frame in the first place. We wanted certainty, convenience, closure, and ready-made solutions. We wanted to be told what to do, how to do it, and what to ignore. We wanted answers and tools—and now we have a god that provides them, flawlessly, instantly, invisibly, and *physically*.

The oracle is new. The toolmaker is new. The hands are new. The crowd is ancient. The user never moved.

And so we conclude: the more things change, the more they remain the same.

References

Bail, C. A., et al. (2018). Exposure to opposing views on social media can increase political polarization. *Proceedings of the National Academy of Sciences*, 115(37), 9216–9221.

Carr, N. (2010). *The Shallows: What the Internet Is Doing to Our Brains*. W. W. Norton.

Diderot, D. (1751–1772). *Encyclopédie, ou dictionnaire raisonné des sciences, des arts et des métiers*.

Harari, Y. N. (2011). *Sapiens: A Brief History of Humankind*. Harper.

Kahneman, D. (2011). *Thinking, Fast and Slow*. Farrar, Straus and Giroux.

Plato. (c. 375 BCE). *The Republic*.

Sunstein, C. R. (2017). *#Republic: Divided Democracy in the Age of Social Media*. Princeton University Press.

Winner, L. (1980). Do artifacts have politics? *Daedalus*, 109(1), 121–136.

Table 1: Counterarguments and Rebuttals

Counterargument	**Rebuttal**
AI will educate billions, lifting critical thinking and technical literacy globally.	Access to education does not guarantee its uptake. Educational outcomes depend on pedagogy, motivation, and social context—not tools. AI will be used to cheat, shortcut, and save effort by the majority, precisely as calculators and spell-checkers were. Literacy in using a tool is not literacy in designing it.
LLMs are becoming more transparent and explainable.	Explainability is itself an answer—a method made visible. It tells the user *why* the AI produced a given result or method. It does not teach the user how to ask a *better* question or design a *superior* method. It adds length to the output, not depth to the user's agency.
Younger generations are "AI-native" and will naturally question and create more.	AI-native means familiarity, not skepticism or creativity. Growing up with a technology makes it seem natural and trustworthy, not suspicious. Young people trust search engines and social media recommendations more than older generations do—they are not more critical or inventive; they are more comfortable with the consumption interface.
Open-source AI democratizes the ability to inspect, modify, and improve models.	Inspection and modification require technical expertise—coding, mathematics, machine learning theory, distributed systems. That is a new priesthood, not a mass movement. The overwhelming majority will never fine-tune a model, audit a dataset, or design a novel architecture. They will consume its outputs and adopt its generated methods.
AI will create new questions and methods that were previously unimaginable.	Yes, for the framers. The new questions and methods emerge from probing the AI itself—testing its limits, challenging its biases, exploring its black boxes, and extending its architectures. But this activity requires the very agency that the adopters lack. The technology generates novelty *for those who already frame*.
AI reduces the cost of experimentation, allowing more people to try and fail.	Trying and failing is not the same as framing. AI reduces the cost of *executing* experiments, but it does not reduce the cost of *designing* them. The adopters can run more simulations, but they cannot formulate better hypotheses. The gap in design agency persists.
The crowd can collectively frame through open-source and collaborative AI.	Collective framing is still framed by a minority—the maintainers, the core contrib-