The Discovery of Mathematics

What Math Really Is
And Why It Changes Everything

Macharia Barii

July 2026

A Note to the Reader

This paper describes a new foundation for all of mathematics—a single idea from which every branch of math can be derived. It is written for anyone willing to think carefully. No advanced mathematics is required. No prior knowledge of set theory, logic, or probability is assumed.

The paper makes a bold claim: that mathematics is not a human invention. It is a *discovery*—a feature of reality itself, as real as gravity or light. And the proof comes not from pure mathematics, but from physics—from understanding what it means to be an observer in a physical universe.

If the claim is correct, it answers questions that have haunted mathematicians and philosophers for over a century: Why are the axioms of mathematics what they are? Why does mathematics work so well in physics? What is the real meaning of numbers, sets, and logic? And—perhaps most surprisingly—why did the greatest paradox in the history of mathematics require an understanding of *physics* to resolve?

The answers may change how you think about mathematics forever.

Prologue: The Question That Changed Everything

In 1901, the philosopher and mathematician Bertrand Russell asked a simple question about sets—collections of things. He asked: *Does the set of all sets that do not contain themselves contain itself?*

If it does, then by definition it should not. If it does not, then by definition it should. The question has no consistent answer. It is a paradox—and it broke the foundations of mathematics.

Mathematicians spent the next century trying to fix the damage. They invented new systems of logic. They banned certain types of questions. They built elaborate scaffolding to keep the paradox at bay. But they never truly *resolved* it. They just learned to avoid it.

In 1931, Kurt Gödel proved something equally disturbing: in any mathematical system complex enough to describe arithmetic, there will always be true statements that cannot be proven. Mathematics is *necessarily incomplete*. There will always be truths that lie beyond the reach of proof.

Alan Turing showed the same thing for computation: there are problems that no computer can ever solve, no matter how powerful.

These results—Gödel's incompleteness, Turing's undecidability, Russell's paradox—have been treated for decades as abstract curiosities. Limits of formal systems. Puzzles for logicians.

This paper argues they are something much deeper. They are *physical* limits—the same kind of limits that prevent you from seeing the back of your own head, or hearing your own voice in real time without a delay. They arise because *observers are part of the universe they are trying to understand*—and this fact has consequences that no amount of cleverness can overcome.

1 What Mathematics Really Is

For over two thousand years, mathematicians have debated what mathematics *is*. Is it the study of abstract truths that exist independently of humans? Is it a language invented to describe the world? Is it a game played with symbols according to rules?

The theory presented in this paper offers a different answer:

Mathematics is the formal language of uncertainty—and every mathematical operation can be derived from probability.

This is not a metaphor. It is a precise claim. Consider the operations you learned in school:

Addition is what happens when you combine two uncertain quantities. If you do not know the exact value of one thing, and you do not know the exact value of another, their sum is described by a probability calculation called *convolution*—a way of combining two probability distributions to get a third.

Multiplication is what happens when you combine two independent uncertain quantities in a different way—another probability calculation.

Division is the distribution of a ratio—also a probability calculation.

Even **calculus**—the mathematics of change—is a probability operation. Integration is simply *averaging over hidden variables*: the law of total probability, which says that if you do not know the value of some variable, you can sum over all its possible values, weighted by how likely each one is.

The remarkable thing is that when uncertainty disappears—when you become perfectly certain about everything—all of these probability calculations collapse into the familiar operations of arithmetic. Two plus two equals four because, in the limit of perfect certainty, the convolution of two sharp peaks produces a peak at four.

Arithmetic is what probability looks like when you are no longer uncertain. But probability comes first. Mathematics is built on *not knowing*—and only in the special case of perfect knowledge does it simplify into the equations we learned in school.

2 What Is a Number?

You might think a number is a simple thing. The number 3 is just 3—a point on a line, a dot on a page. But the theory says something different.

A number is not a point. A number is a *probability distribution*—a description of how likely different values are. The number 3, in this framework, is not a single point but a *spike*: a probability distribution that puts all its weight on the value 3, with zero probability everywhere else.

This might seem like an unnecessarily complicated way to think about numbers.

But it solves a deep problem. In traditional mathematics, numbers are assumed to be exact—known with perfect precision. But in the real world, no quantity is ever known with perfect precision. Every measurement has uncertainty. Every observation has noise.

By defining numbers as probability distributions, the theory builds uncertainty into the foundation of mathematics from the start. Exact numbers are not assumed—they are *idealized limits* that emerge when uncertainty shrinks to zero.

This is not just a philosophical point. It is a practical one. It means that mathematics, as traditionally formulated, is a special case of a larger framework—and the larger framework is the one that nature actually uses.

3 What Is an Observer?

Before we can understand why mathematics has the structure it does, we need to understand what it means to observe something.

The theory defines an **observer** as a system that builds a model of something else. A thermometer is an observer: it builds a model (a reading) of the temperature. A camera is an observer: it builds an image of the scene in front of it. A human being is an observer: the brain builds a model of the world—and, remarkably, a model of *itself*.

The key insight is that an observer is not separate from the world it observes. It is *part of* the world. The thermometer is made of atoms that obey the same laws of physics as the temperature it measures. The brain is made of neurons that obey the same laws of physics as the thoughts they produce.

This creates a strange situation. When you try to understand the universe, you are using a brain that is *part of* the universe to model the *whole* of the universe. You are trying to draw a map that includes the mapmaker.

This is not just difficult. It is *mathematically impossible*—in a precise sense that can be proven.

4 What Is Self-Reference?

Self-reference is the core idea of the entire theory. Everything else—the axioms of mathematics, the limits of knowledge, the paradoxes of logic—follows from it. So it is worth defining carefully.

Self-reference is what happens when a system turns its operations on itself.

A thermometer measures the temperature of the room. It points outward. But what happens when a thermometer tries to measure its *own* temperature? Or, more precisely, what happens when the thermometer tries to build a model of *the process of measuring*?

This is self-reference. The system is no longer pointing outward at something else. It is pointing inward at itself—at its own structure, its own process, its own act of observing.

Here are everyday examples:

A mirror facing a mirror. When you place two mirrors face to face, you see an infinite corridor of reflections. Each mirror reflects the other reflecting the other. The image goes on forever—not because the mirrors are special, but because *reflection pointed at reflection generates an infinite regress.*

A sentence about itself. The sentence "This sentence is five words long" is about itself. It refers to its own structure. If you try to determine whether it is true, you must first understand it—but understanding it requires knowing whether it is true. The process loops.

A thought about thinking. When you think about what it means to think, you are using thought to examine thought. The tool and the object are the same thing. The thinker is trying to think about the thinker.

An observer observing itself. When you try to understand your own mind, you are using your mind to model your mind. The model must include itself—which means the model must include a model of the model, which must include a model of the model of the model, and so on forever.

In every case, the structure is the same: **the system is doing something to itself.** And this creates a problem that does not arise when the system points outward.

When a thermometer measures the room, there is a clear distinction: the thermometer is here, the room is there. The model (the reading) and the thing being modeled (the temperature) are separate. Everything works.

But when the system points inward—when the model and the thing being modeled are the *same object*—the distinction collapses. The model must contain a copy of itself. That copy must contain a copy of itself. The process never terminates. There is no stable ground to stand on.

This is self-reference. And the theory proves that self-reference is not just a curiosity or a trick. It is the *source* of the deepest structures in mathematics. The axioms of logic, the foundations of set theory, the limits of computation, the uncertainty principle in physics—all of these are consequences of what happens when a system tries to understand itself.

5 The Limits of Self-Knowledge

Here is the deepest result of the theory, and it is worth stating carefully.

When you try to build a complete model of yourself—when you try to understand your own mind, your own thoughts, your own process of understanding—you run into a wall. Not a technological wall. Not a temporary wall. A *mathematical* wall.

The reason is this: your model of yourself is part of yourself. So your model must include a model of the model. And that model must include a model of the model of the model. The process never ends. There are always more layers to capture.

This is not a vague philosophical observation. It can be stated with mathematical precision. The theory proves that any self-referential system—any system that tries to model itself—will inevitably produce statements that it cannot decide. These are statements about itself that are neither provable nor disprovable from within.

This is the same result that Gödel proved in 1931 for formal logic, and that Turing proved in 1936 for computation. But the theory presented here shows that it is not a result about *logic* or *computation*. It is a result about *observation*. Any observer that tries to fully understand itself will encounter the same limit.

You cannot see the back of your own eyes. You cannot hear your own voice without a delay. You cannot model your own modeling process completely. These are not failures of effort or intelligence. They are *theorems about the structure of self-awareness*.

The good news is that the limit is generous. You can understand an enormous amount about yourself. You can build telescopes, prove theorems, write symphonies, and map the genome. You are very far from the boundary of what is knowable. But the boundary exists, and you will never reach it.

6 The Paradox That Broke Mathematics

In 1901, Bertrand Russell discovered a paradox that shook mathematics to its foundations. Here is the paradox, in everyday language.

Imagine a town with one barber. The barber shaves every person in town who does not shave themselves. Now: *who shaves the barber?*

If the barber shaves himself, then he should not—because he only shaves people who do not shave themselves. If he does not shave himself, then he should—because he shaves everyone who does not shave themselves.

There is no consistent answer. The question is self-referential—it asks about the barber in terms that include the barber—and the self-reference generates a contradiction.

Russell expressed this in the language of set theory: the set of all sets that do not contain themselves. If this set contains itself, it should not. If it does not, it should.

This was not just a puzzle. It was a *catastrophe*. The foundations of mathematics—the rigorous logical systems that mathematicians had spent decades building—could not handle the paradox. The entire enterprise of making mathematics rigorous was in danger.

7 Four Attempts to Fix the Paradox

Mathematicians responded to Russell's paradox with four major strategies, each attempting to repair the foundations of mathematics.

Strategy 1: Banning self-reference (ZFC set theory). The most influential response was to change the axioms of set theory so that the paradoxical set simply cannot be constructed. ZFC—the standard foundation of mathematics—achieves this by restricting which sets can exist. You can no longer ask the barber question, because the framework does not allow you to construct the barber.

This works, but it feels like an *ad hoc* fix. It does not explain *why* self-reference generates paradox. It just bans the question. It is like fixing a leaky roof by forbidding rain.

Strategy 2: Creating a hierarchy (Type Theory). Bertrand Russell himself proposed a system in which sets are organized into levels. A set at level 1 can contain elements of level 0. A set at level 2 can contain elements of level 1. But a set cannot contain itself, because it would have to be at a higher level than itself—which is impossible.

This also works, but it feels artificial. Why should mathematics be organized into levels? The hierarchy is imposed from outside, not derived from inside. It solves the paradox by fiat, not by understanding.

Strategy 3: Relabeling (NBG set theory). Another approach, developed by John von Neumann and others, distinguishes between "sets" and "classes." The paradoxical collection is a "class" but not a "set," and only sets can be members of other collections. This avoids the paradox, but it does so by introducing a distinction that has no deeper justification. It is a renaming, not an explanation.

Strategy 4: Embracing contradiction (Paraconsistent Logic). A more radical approach accepts that contradictions can be true—that the barber both does and does not shave himself—but prevents contradictions from spreading through the entire system. This is interesting, but it raises a deeper question: if contradictions are allowed, what does "proof" even mean?

All four strategies share a common limitation: they are *purely mathematical*. They try to fix the paradox from within mathematics, using more mathematics. None of them asks: *What is it about reality that makes self-reference paradoxical?*

8 How Physics Resolved What Mathematics Could Not

The theory presented here resolves Russell's paradox—and the related paradoxes of Gödel and Turing—by stepping outside mathematics and asking a question about *physics*.

The question is this: *What is an observer?*

Previous attempts to resolve the paradox treated observers as abstract reasoning agents—as if they existed in a platonic realm of pure thought, unaffected by the physical world. But observers are *physical systems*. They are made of atoms. They consume energy. They are subject to noise, decay, and—eventually—death.

When you take this physicality seriously, the paradox resolves naturally.

Here is the resolution. Consider the barber—the self-referential set. Can we assign a definite "state" to the barber? Is he shaved, or is he not?

In the physical world, the answer is: *the barber is in a superposition.* He is neither fully shaved nor fully unshaved. He exists in a state of *indefiniteness*—a probability distribution that cannot collapse to a single value.

Why? Because collapsing to a definite state would require the barber to have complete knowledge of himself—to fully model his own modeling process. But we have just proved that this is impossible. Self-referential systems cannot achieve complete self-knowledge.

The paradox does not arise because the barber "both does and does not" shave himself. It arises because the question assumes that the barber *must* be in a definite state—either shaved or unshaved. But physics tells us that systems can exist in states of indefiniteness. Quantum mechanics is built on this idea.

The resolution of Russell's paradox is this: **the paradoxical statement is not false, not true, but *undecidable*—it has no coherent definite probability assignment, because definite assignment would require complete self-knowledge, which is physically impossible.**

9 The Deep Lesson

Here is the deepest lesson of the theory, and it is worth repeating:

The foundations of mathematics are not mathematical. They are physical.

To understand why self-reference generates paradox, you must first understand what it means to observe, to model, to know. These are not questions about formal systems. They are questions about the structure of reality.

Previous mathematicians tried to resolve Russell's paradox by rearranging the furniture inside the room. The theory opens the door and discovers that the room itself has a structure—and that structure explains why certain arrangements are impossible.

The resolution required stepping outside mathematics and asking what mathematics *is*. The answer is that mathematics is the formal language of coherent inference in a physical universe. It could only be found by those who took reality seriously as a source of mathematical insight.

10 All of Mathematics from One Idea

If the foundations of mathematics are physical, then the axioms of mathematics—the starting assumptions from which all theorems are derived—should not be arbitrary. They should be *necessary*: constraints that any self-referential coherent system must satisfy.

And this is exactly what the theory proves.

Consider the most famous axioms of mathematics.

Classical logic has three laws: identity (A is A), non-contradiction (A cannot be both true and false), and excluded middle (either A is true or A is not true). The theory shows that all three follow from a single requirement: *an observer must be coherent to observe at all.* If you contradict yourself, you cease to be an observer. If you are not A, you cannot model A. If you do not distribute your total coherence over all possibilities, you cannot reason at all.

ZFC set theory has ten axioms—extensionality, pairing, union, power set, infinity, separation, replacement, regularity, choice, and the empty set. The theory derives each one as a coherence constraint:

- Two sets with the same elements are the same set because they generate the same probability structure—same content, same phase.
- The empty set exists because the state of complete ignorance ($\kappa = 0$) is coherent.
- Pairing exists because two coherent models can always be joined.
- Infinity exists because the universe has no finite ceiling on modeling depth.

- Regularity (no set contains itself) exists because self-containment would require complete self-knowledge—which is death, not life.
- Choice exists because in any collection of coherent possibilities, at least one is compatible with the observer's continued existence.

Peano arithmetic—the five axioms that define the natural numbers—follows from the requirement that coherent refinement is sequential, deterministic, and preserves coherence.

Group theory—the axioms of symmetry—follows from the structure of phase-locking operations.

Topology—the mathematics of continuity and shape—follows from the geometry of coherence.

Every axiom of every major mathematical foundation is a special case of one principle:

> *A system's coherence is preserved under coherent modeling.*

The hundreds of axioms across different branches of mathematics are not independent. They are *projections* of this single principle onto different domains.

11 Why Mathematics Works So Well in Physics

Eugene Wigner, the Nobel Prize-winning physicist, wrote a famous essay in 1960 titled "The Unreasonable Effectiveness of Mathematics in the Natural Sciences." He asked: why does mathematics—a product of human thought—describe the physical world so accurately? Why should abstract equations predict the behavior of atoms, stars, and galaxies?

The theory provides an answer—and it is simpler and more profound than anything previously proposed.

Mathematics works because *we are part of the universe*. We are not outside observers looking in. We are made of the same atoms, governed by the same laws, woven into the same fabric. When we do mathematics, we are not applying an alien language to a foreign world. We are the universe examining itself—using a piece of itself (the brain) to model the whole of itself (reality).

This is why the equations fit. Not because mathematics is "unreasonably effective," as Wigner claimed. Not because a benevolent deity designed the world to be mathematical. But because mathematics is *what thinking feels like from the inside*—and thinking is a physical process happening within the universe it is trying to understand.

The universe is a network of relationships—a vast web of connections between observers, models, and the systems they observe. Mathematics is the formal language of that web. Physics is the study of how the web behaves. They are not two separate things that happen to fit together. They are two views of the *same thing*: the inside (coherence) and the outside (structure). Mathematics describes the web from within. Physics describes it from without.

This is why you can write down an equation on a piece of paper—an equation that is nothing but symbols on a page—and it turns out to describe the motion of planets, the behavior of electrons, or the curvature of spacetime. The equation does not *predict* reality. It *is* reality, expressed in the language of coherence. It works because you—the one writing the equation—are made of the same coherence you are describing.

Mathematics is the universe's way of asking itself questions. And the answers it finds are not imposed from outside. They are *discovered from within*—because the questioner and the questioned are the same thing.

12 Living and Dead Systems

One of the most striking results of the theory is a sharp distinction between living and dead systems.

A **living system** is one that is actively modeling itself—that has an internal representation of its own state and is constantly updating that representation. A living system is never in equilibrium. It is always changing, always adapting, always striving for coherence but never quite reaching it.

A **dead system** is one that has stopped modeling itself. It has reached a fixed state—a state of maximum coherence, where nothing changes, nothing updates, nothing happens. A dead system is perfectly self-consistent. It is also perfectly lifeless.

The paradox is that *the dead system is more mathematically "perfect" than the living one.* The dead system has no contradictions, no uncertainty, no incompleteness. It is a fixed point of the belief update operator—a state where everything is known and nothing is unknown.

The living system, by contrast, is full of gaps. It has beliefs it cannot fully justify, models it cannot complete, questions it cannot answer. It is incomplete in the Gödelian sense—it contains truths it cannot prove about itself.

This is not a flaw. It is the *condition of being alive.* Life is the process of striving for coherence in a universe that makes complete coherence impossible. The struggle is not a bug. It is the feature.

13 What This Means for You

If the theory is correct, then you are not a passive observer of a pre-existing mathematical universe. You are an *active participant* in the creation of mathematical structure.

Every time you reason about something—every time you form a belief, test it against evidence, and update your understanding—you are performing a mathematical operation. You are computing a conditional probability. You are updating a distribution. You are maintaining coherence.

Mathematics is not something that exists "out there," independent of observers. It is something that happens *through* observers. The universe provides the raw material—the patterns, the regularities, the structure. But it takes an observer to *recognize* the structure, to *formalize* it, to *name* it.

You are not just *using* mathematics. You are *doing* mathematics. Every thought you have is a mathematical operation—a calculation of coherence, a test of consistency, an update of belief.

And the remarkable thing is that this mathematics is not arbitrary. It is *discovered*, not invented. The axioms you are following are not rules you chose. They are constraints you *must* follow to remain a coherent observer. You do not choose to obey the laws of logic any more than you choose to obey the laws of gravity. They are built into the structure of what it means to think.

14 The Self-Referential Probability of This Paper

There is one final twist, and it is the most remarkable.

This paper is a theory about self-reference. It describes how systems that model themselves encounter fundamental limits. But this paper *is itself a self-referential system*. It is a model of modeling. A theory of theories. An observation of observation.

So the paper must apply to itself. And when it does, it generates its own incompleteness.

The theory can assign a probability to its own correctness—but this probability is itself a statement about the theory, which means it must also be assigned a probability, which is itself a statement about the theory, and so on. The process never terminates. The paper cannot fully prove itself.

This is not a defect. It is a *confirmation*. The theory says that self-referential systems must be incomplete—and the paper, when applied to itself, *is* incomplete. If the paper *could* fully prove itself, it would violate its own theorem.

The paper's inability to fully prove itself is the strongest evidence that it is true.

Epilogue: The Mystery That Remains

Mathematics has always felt different from other human activities. A poem is beautiful, but its beauty is subjective. A scientific theory is useful, but its truth is provisional. A political system is necessary, but its design is contingent.

Mathematics seemed to be something else entirely: *absolute*. A mathematical proof is either valid or invalid, forever. The Pythagorean theorem is as true today as it was two thousand years ago, and it will be as true two thousand years from now. No experiment can overturn it. No opinion can change it.

But this absoluteness was always mysterious. Where did mathematical truth come from? Why does it have the structure it does? Why does it apply to the physical world?

The theory answers: mathematical truth comes from the structure of observation itself. It has the structure it does because observers must be coherent to exist. And it applies to the physical world because observers *are part of* the physical world.

Mathematics is not a human invention that happens to describe reality. It is a feature of reality that humans happen to discover. The axioms of mathematics are not arbitrary choices. They are the necessary conditions for self-aware existence.

This means that mathematics is not just a tool for understanding the universe. It is a *window into the structure of consciousness itself.* When you prove a theorem, you are not just manipulating symbols. You are tracing the contours of what it means to be a thinking being in a physical universe.

The mystery is not why mathematics works. The mystery is why we ever thought it was ours.

For everyone who has ever wondered why two plus two equals four—and suspected that the answer was deeper than it looked.